

数字图像处理

Digital Image Processing

Image as A Function

Discrete Sampling of A Function

Spatial Convolution

Most slides are courtesy of Juyong Zhang

Mathematical Convolution?

$$\begin{aligned}(f * g)(t) &\stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau \\&= \int_{-\infty}^{\infty} f(t - \tau) g(\tau) d\tau.\end{aligned}$$

Meaning of Convolution?

滑动平均

- Mathematics
- Signal
 - 激励信号、响应函数
- Statistics
- Life
 - 下棋
 - 复利
 - ...

Spatial Filtering

Let \mathbf{I} and \mathbf{J} be images such that $\mathbf{J} = T[\mathbf{I}]$.

$T[\cdot]$ represents a transformation, such that,

$$\begin{aligned}\mathbf{J}(r,c) &= T[\mathbf{I}](r,c) = \\ f\left(\left\{\mathbf{I}(\rho, \chi) \mid \rho \in \{r-s, \dots, r, \dots, r+s\}, \chi \in \{c-d, \dots, c, \dots, c+d\}\right\}\right).\end{aligned}$$

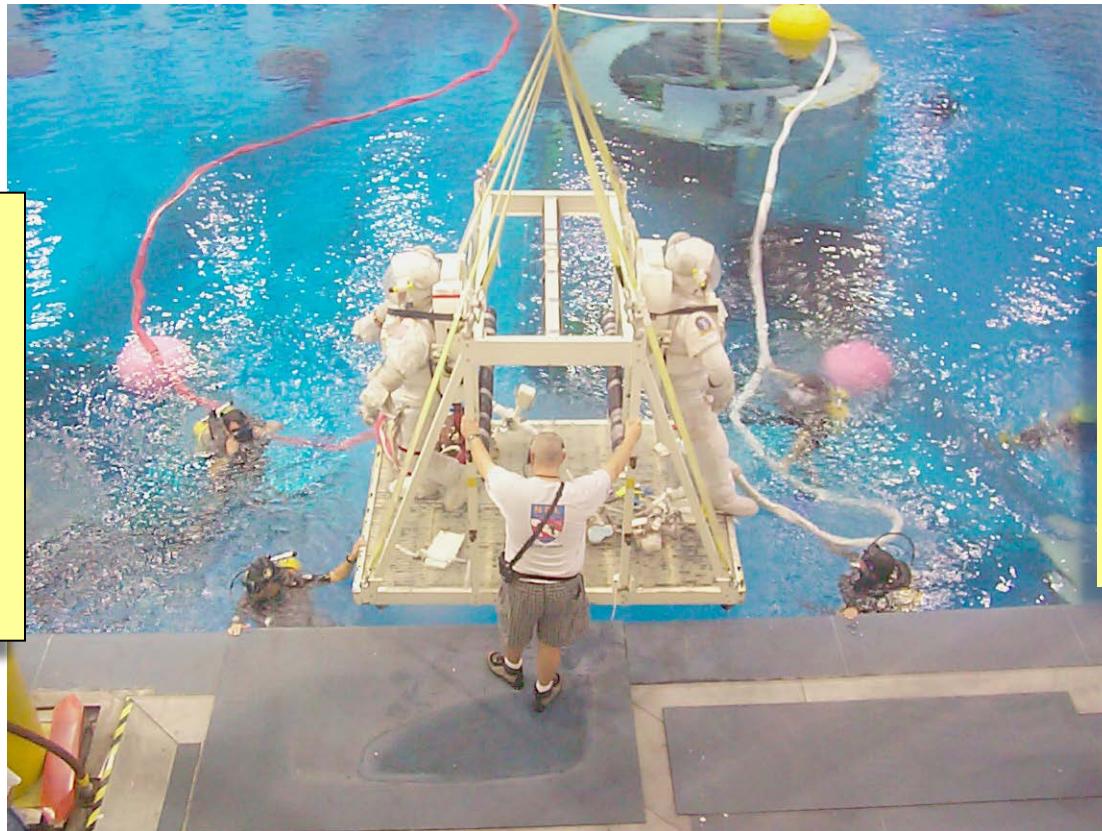
That is, the value of the transformed image, \mathbf{J} , at pixel location (r,c) is a function of the values of the original image, \mathbf{I} , in a $2s+1 \times 2d+1$ rectangular neighborhood centered on pixel location (r,c) .

Moving Windows

- | The value, $\mathbf{J}(r,c) = T[\mathbf{I}](r,c)$, is a function of a rectangular neighborhood centered on pixel location (r,c) in \mathbf{I} .
- | There is a different neighborhood for each pixel location, but if the dimensions of the neighborhood are the same for each location, then transform T is sometimes called a *moving window transform*.

Moving-Window Transformations

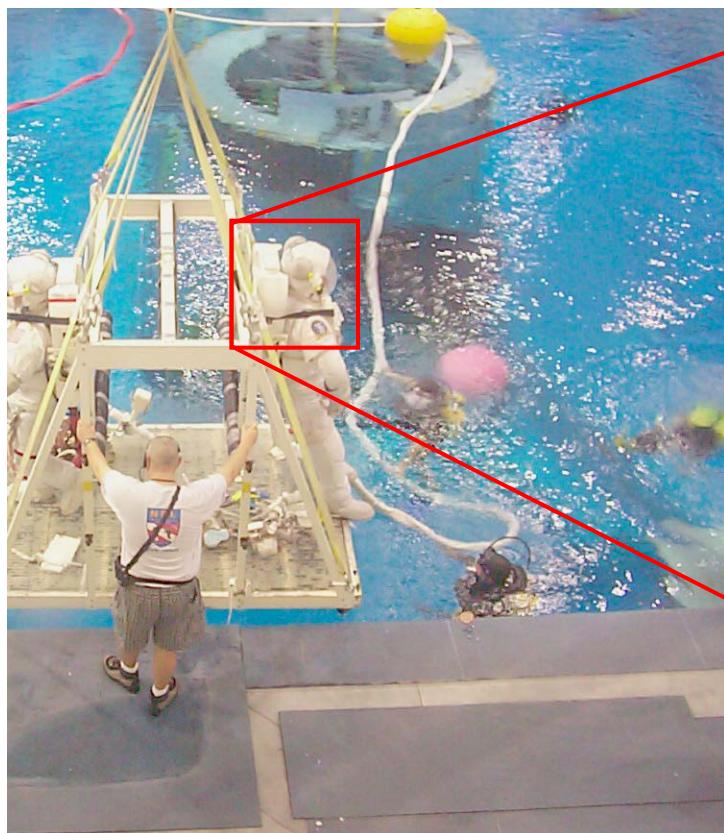
Neutral
Buoyancy
Facility at
NASA
Johnson Space
Center



We'll take a
section of
this image to
demonstrate
the MWT

photo: R.A.Peters II, 1999

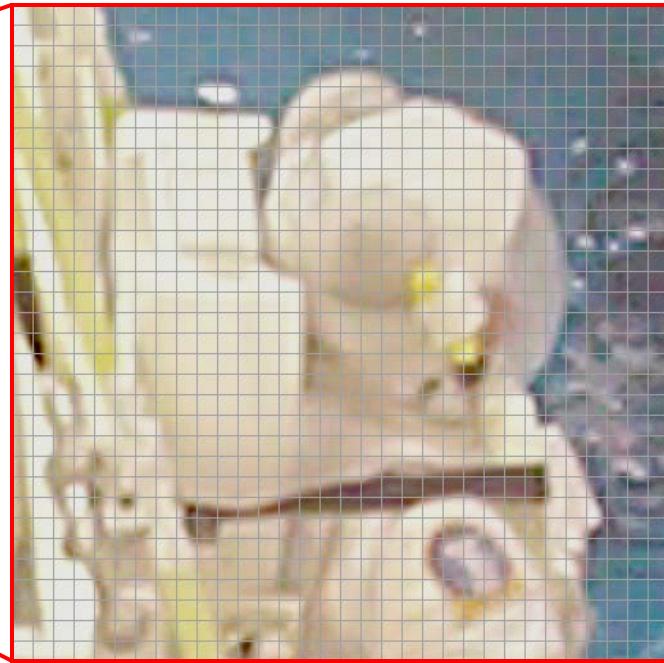
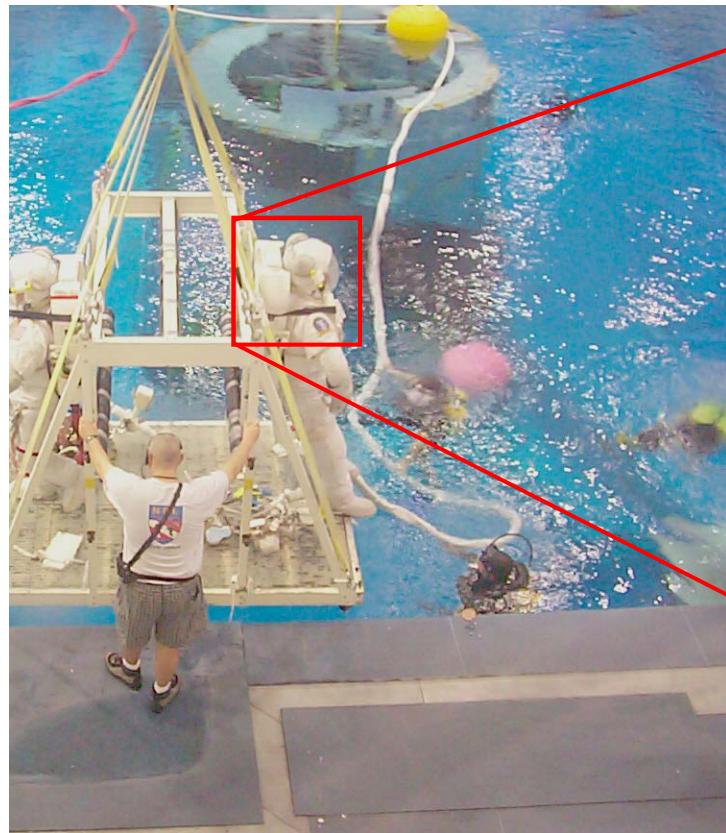
Moving-Window Transformations



operate on this region

Pixelize the section to better see the effects.

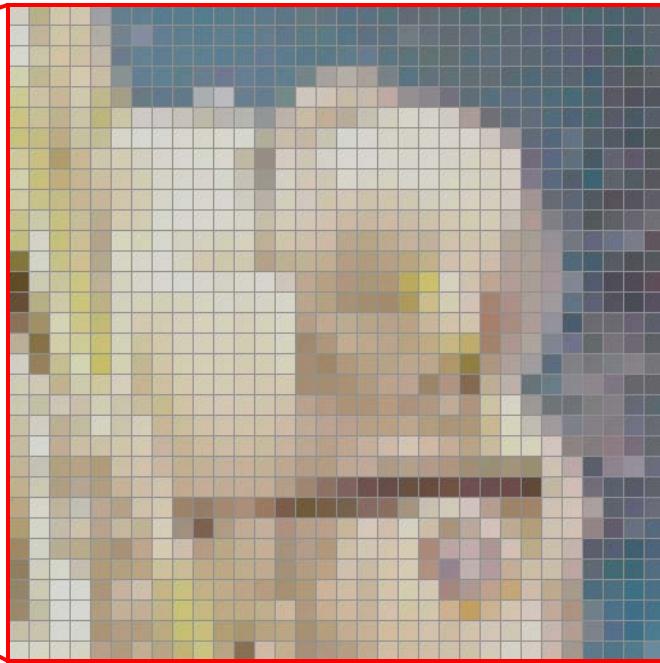
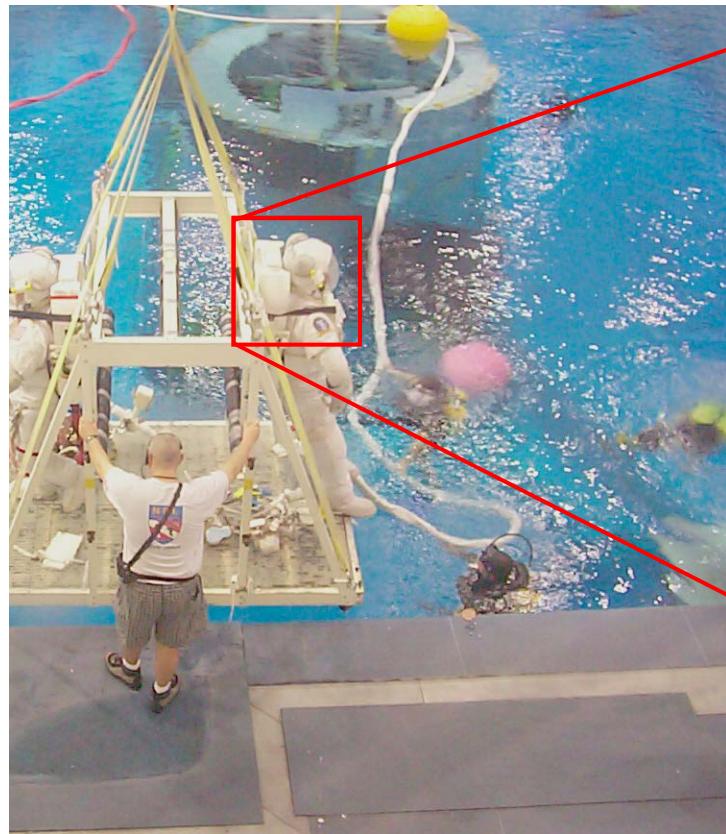
Moving-Window Transformations



apply a pixel grid

Pixelize the section to better see the effects.

Moving-Window Transformations

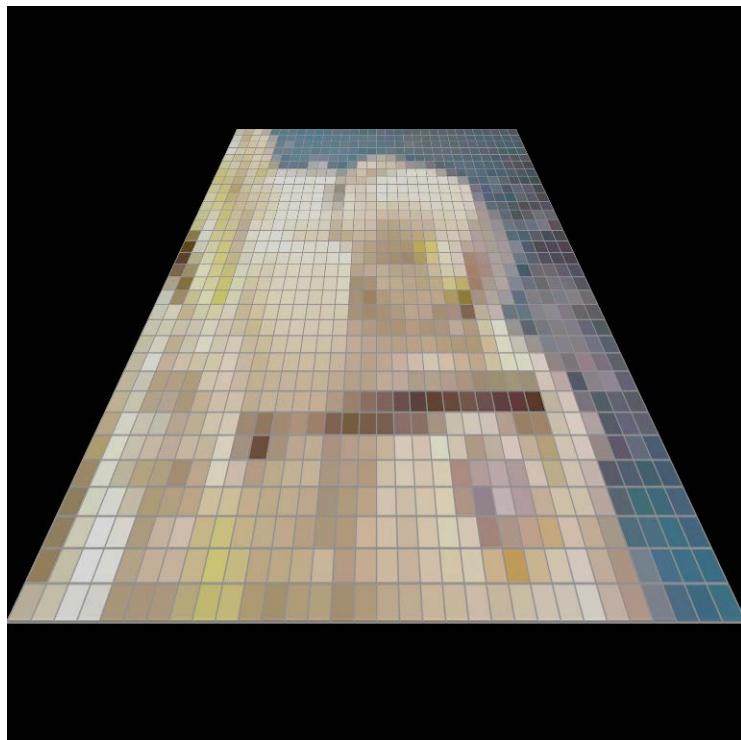


sample (average in the squares).

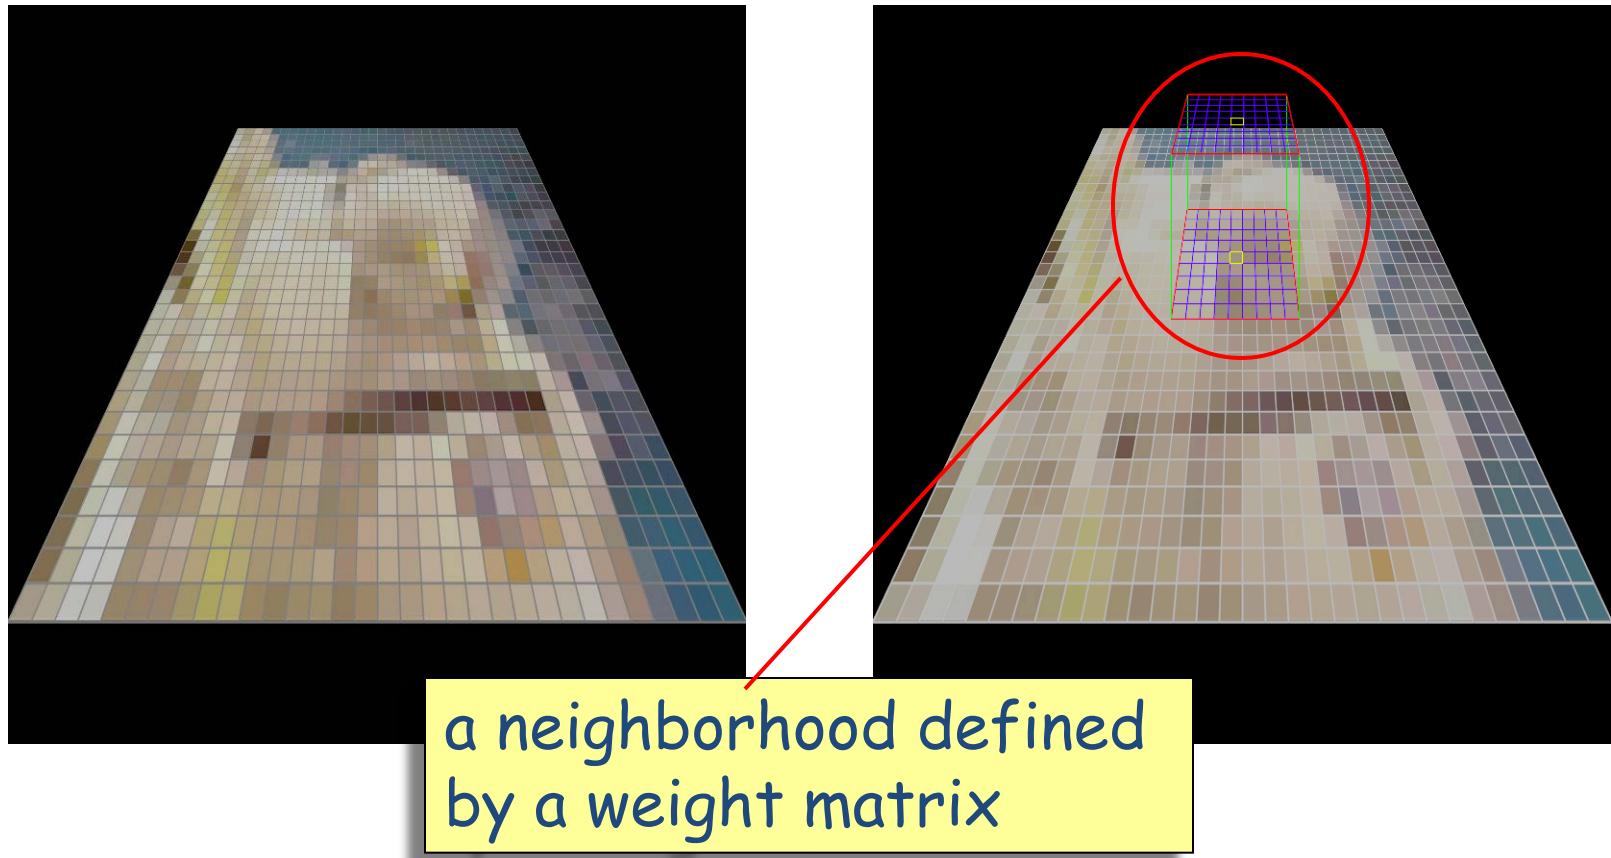
Moving-Window Transformations



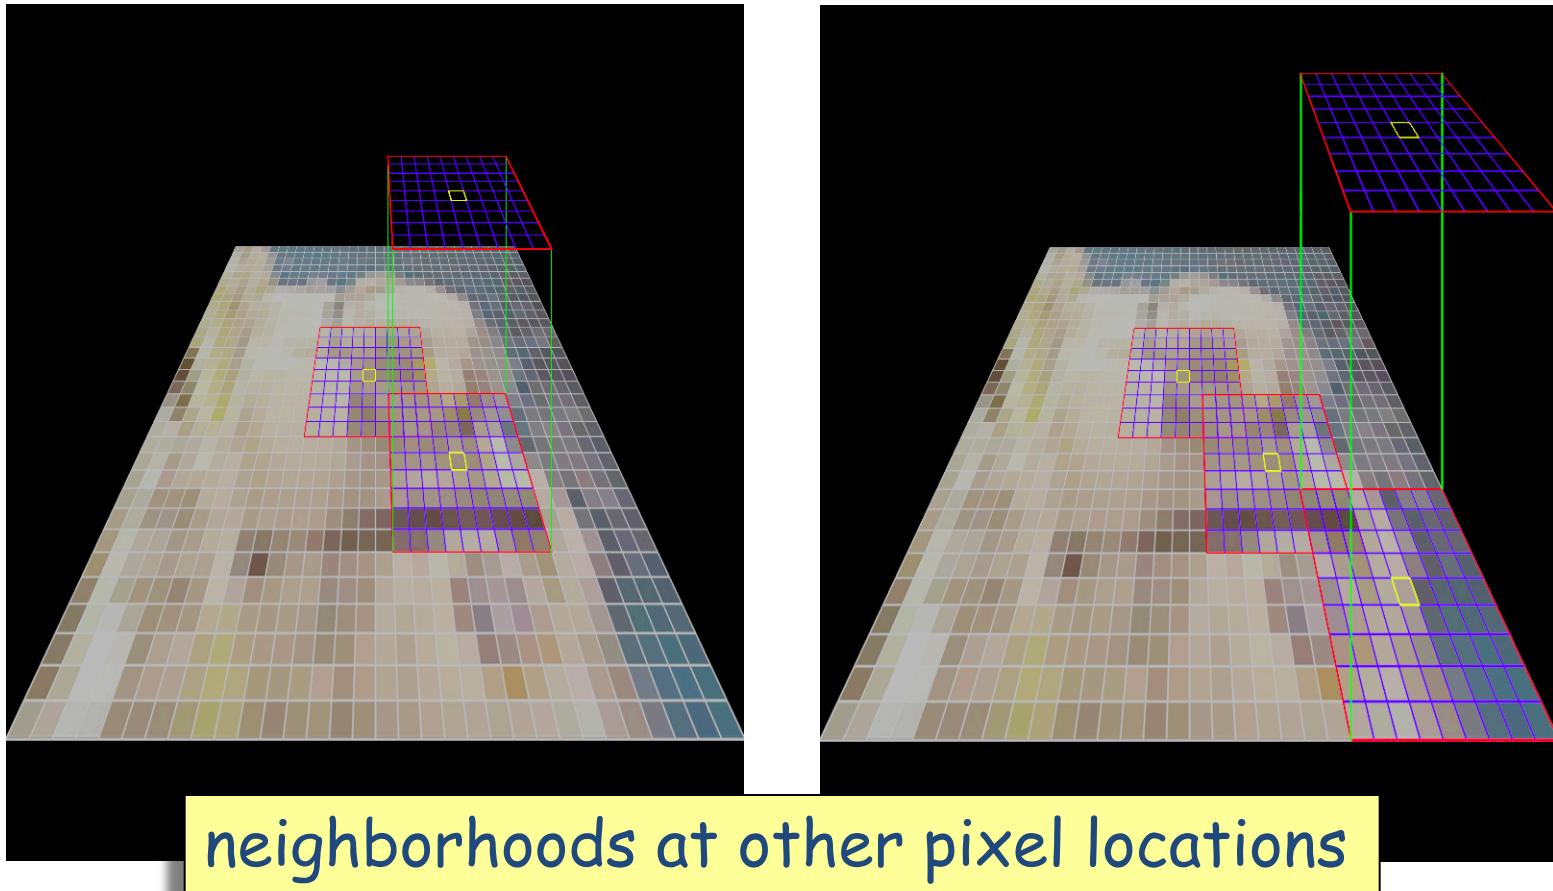
lets get
some
perspective
on this



Moving-Window Transformations

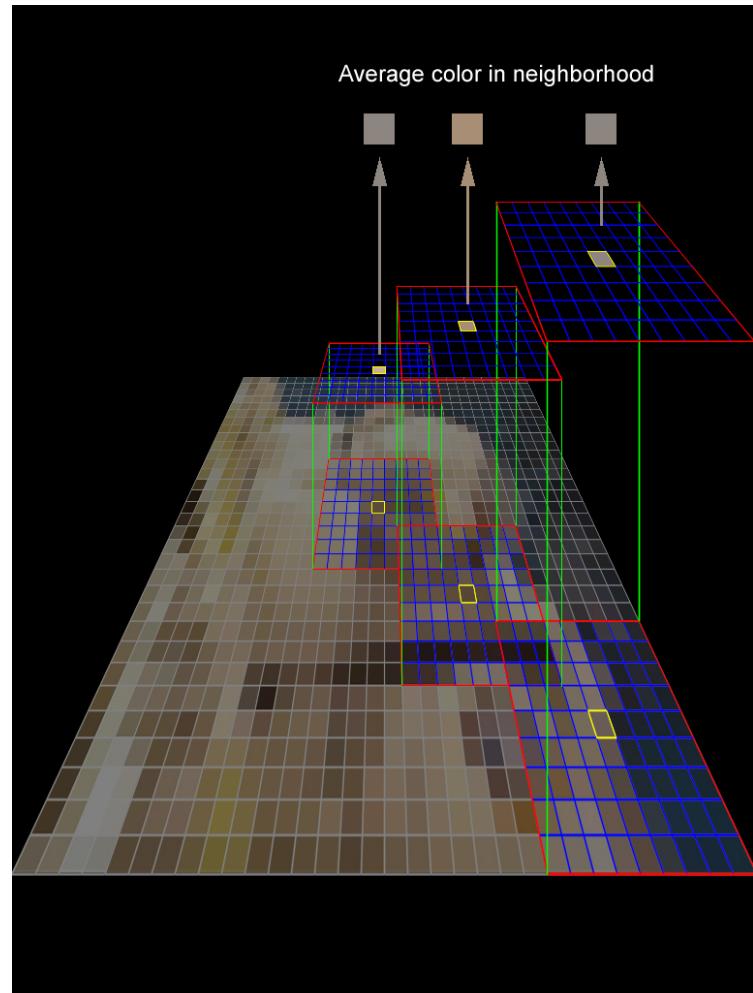


Moving-Window Transformations

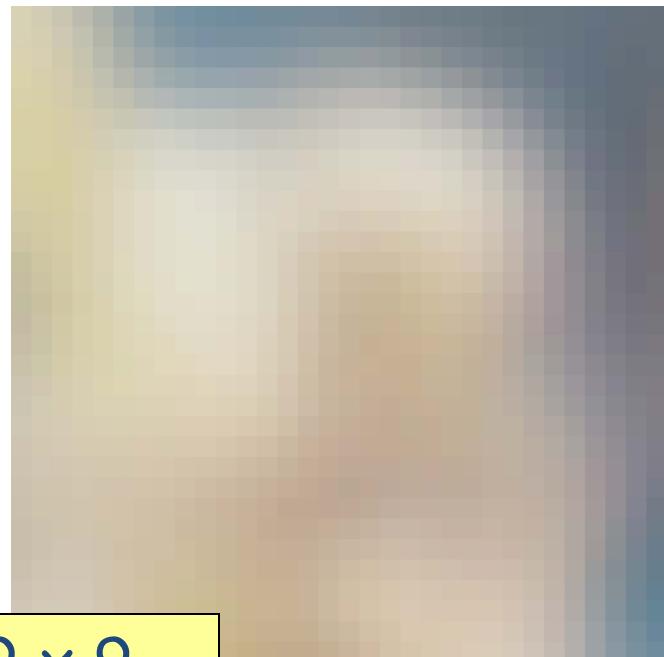
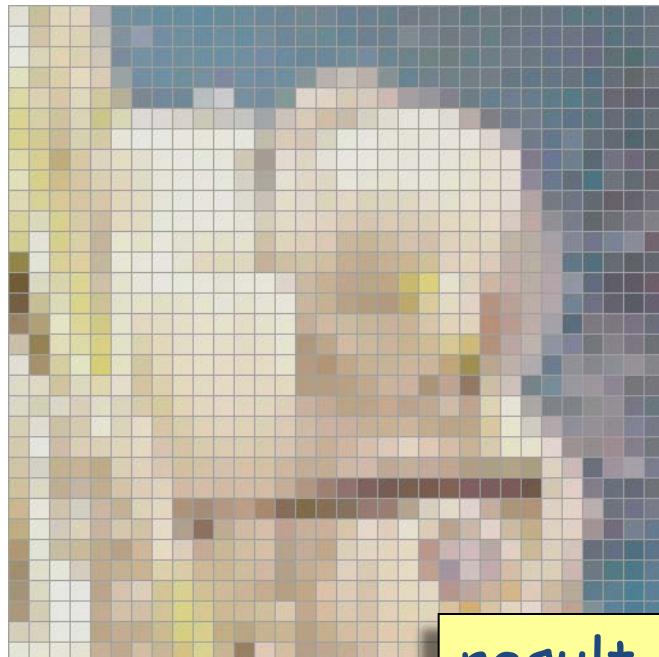


Linear Moving-Window Transformations (i.e. convolution)

The output of the transform at each pixel is the (weighted) average of the pixels in the neighborhood.



Moving-Window Transformations



result of a 9×9
uniform averaging

注：
 1) 滤波器核 $h(i,j)$ 定义了对应 $I(r-i,c-j)$ 处的加权值
 3) 采用这种定义式的卷积，本质是相关
 2) 如果采用 $h(-i,-j)$ ，则是相关运算

Convolution: Mathematical Representation

If a MW transformation is *linear* then it is a *convolution*.

$$\mathbf{J}(r,c) = [\mathbf{I} * \mathbf{h}](r,c) = \iint_{-\infty}^{\infty} \mathbf{I}(r-\rho, c-\chi) \mathbf{h}(\rho, \chi) d\rho d\chi,$$

for a real image ($\mathbf{I}: R \times R \rightarrow R$), or for a digital image ($\mathbf{I}: Z \times Z \rightarrow Z$):

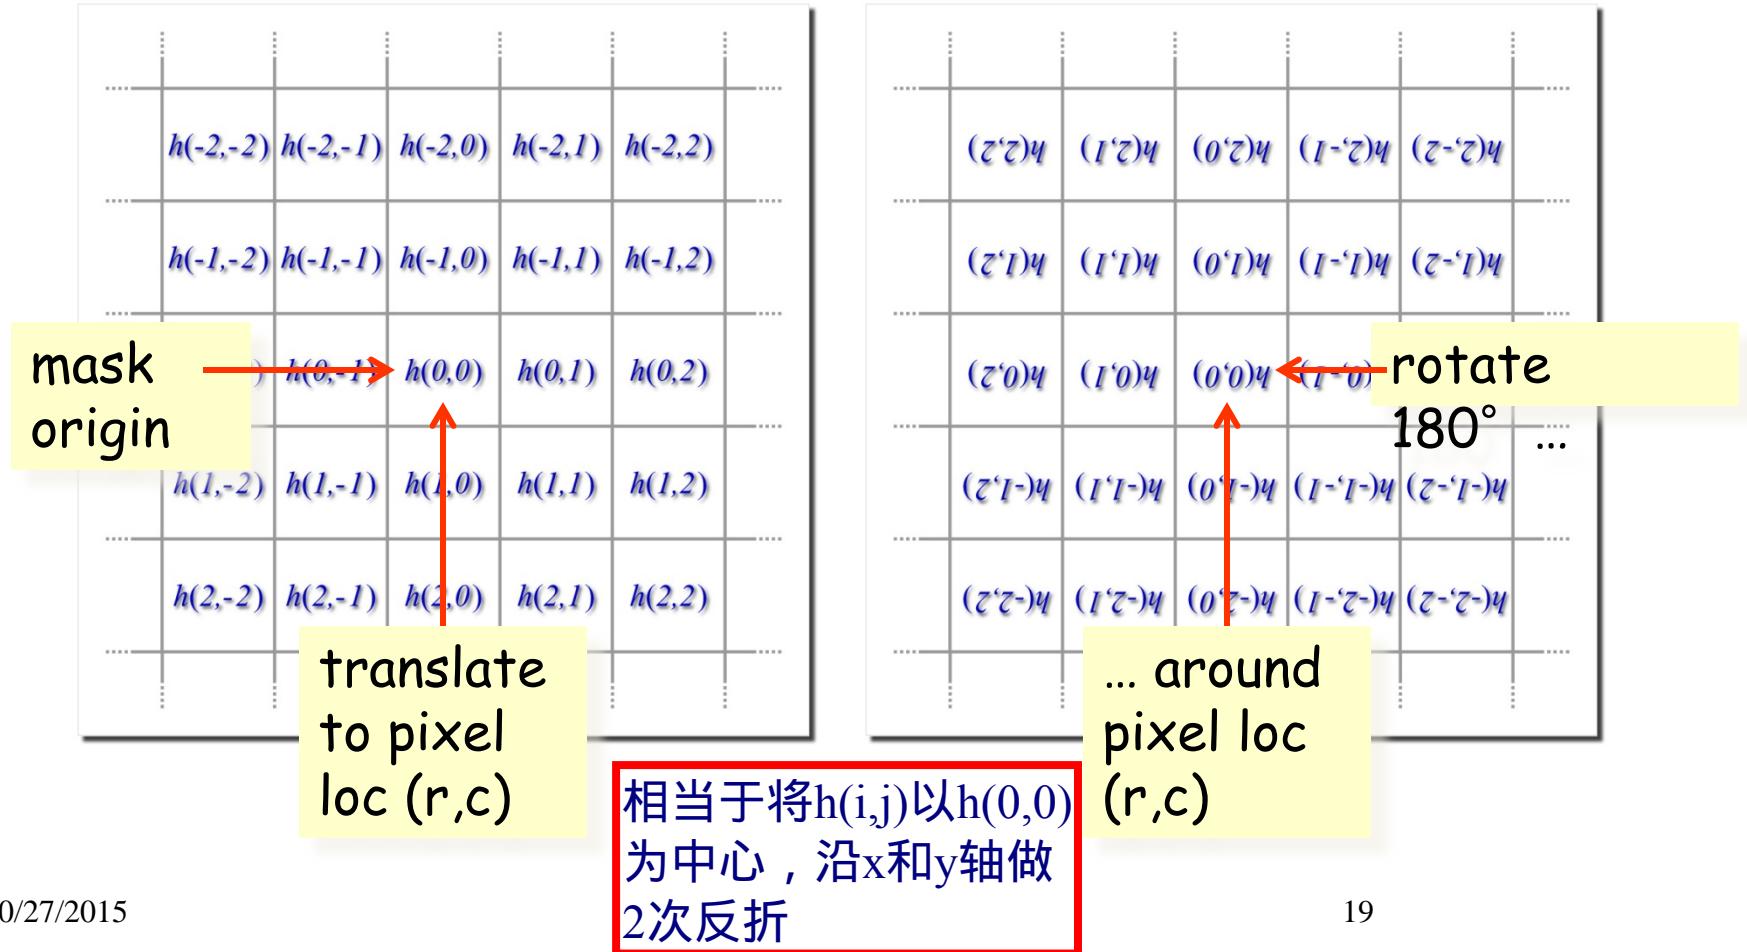
$$\mathbf{J}(r,c) = [\mathbf{I} * \mathbf{h}](r,c) = \sum_{\rho=-s}^s \sum_{\chi=-d}^d \mathbf{I}(r-\rho, c-\chi) \mathbf{h}(\rho, \chi)$$

$$\begin{array}{ll}
 h(i,j) : \text{以} 3 \times 3 \text{为例} - & \\
 \begin{array}{ccc|ccc}
 h(-1,-1) & h(-1,0) & h(-1,1) & h(1,1) & h(1,0) & h(1,-1) \\
 h(0,-1) & h(0,0) & h(0,1) & \rightarrow & h(0,1) & h(0,0) & h(0,-1) \\
 h(1,-1) & h(1,0) & h(1,1) & & h(-1,1) & h(-1,0) & h(-1,-1)
 \end{array}
 \end{array}$$

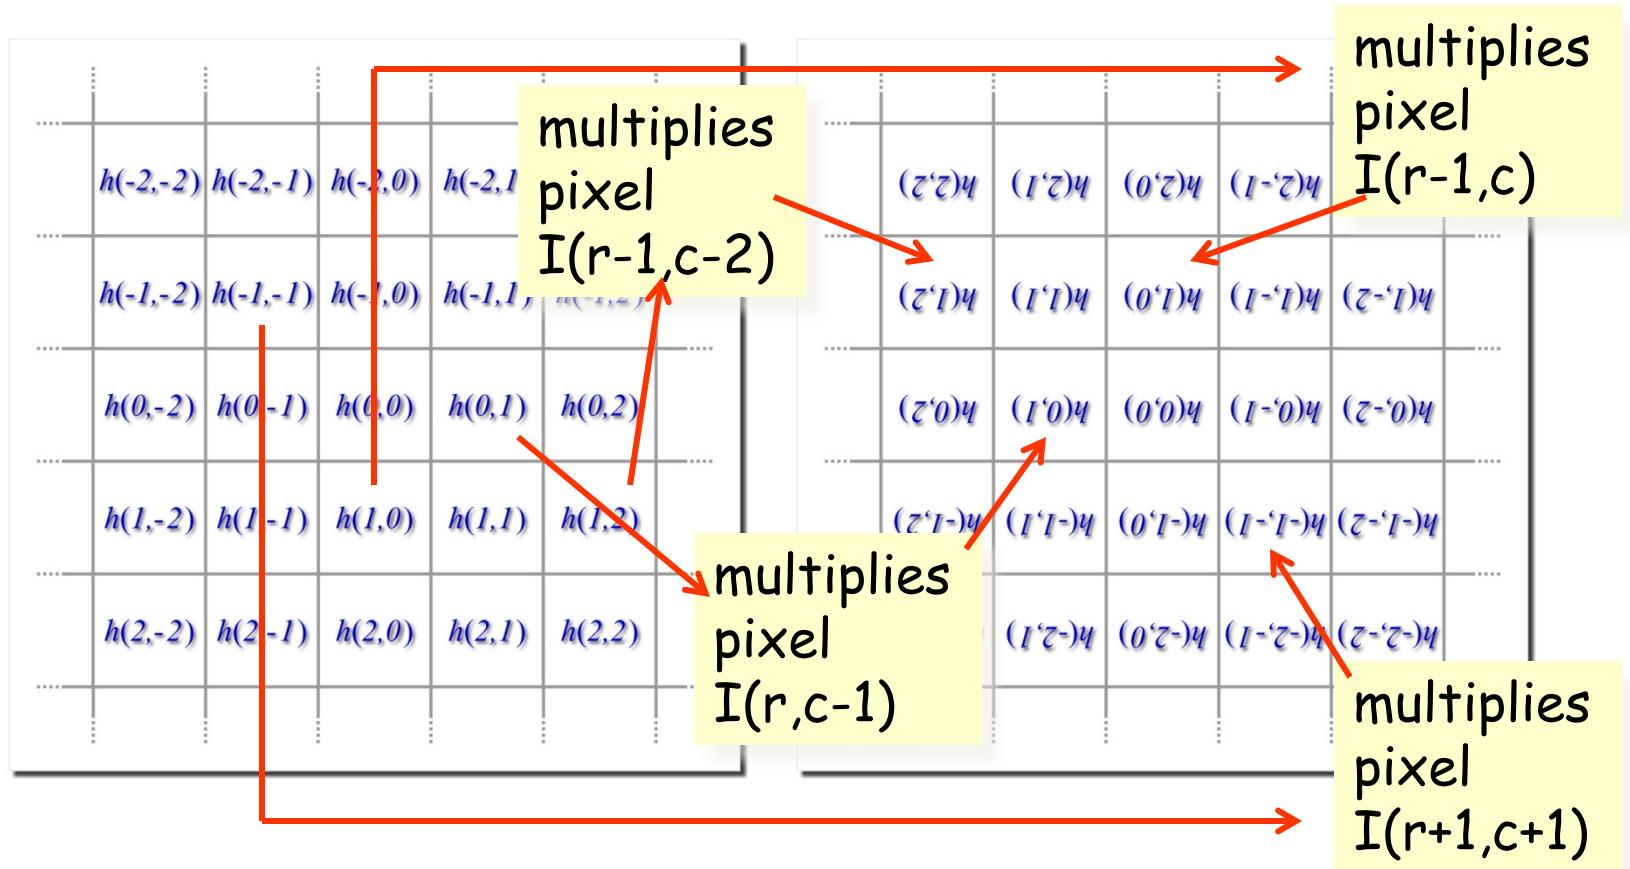
Convolution Mask (Weight Matrix)

- The object, $h(\rho, \chi)$, in the equation is a weighting function, or in the discrete case, a rectangular matrix of numbers.
- The matrix is the moving window.
- Pixel (r,c) in the output image is the weighted sum of pixels from the original image in the neighborhood of (r,c) traced by the matrix.
- Each pixel in the neighborhood of (r,c) is multiplied by the corresponding matrix value — after the matrix is rotated by 180° .
- The sum of those products is the value of pixel (r,c) in the output image

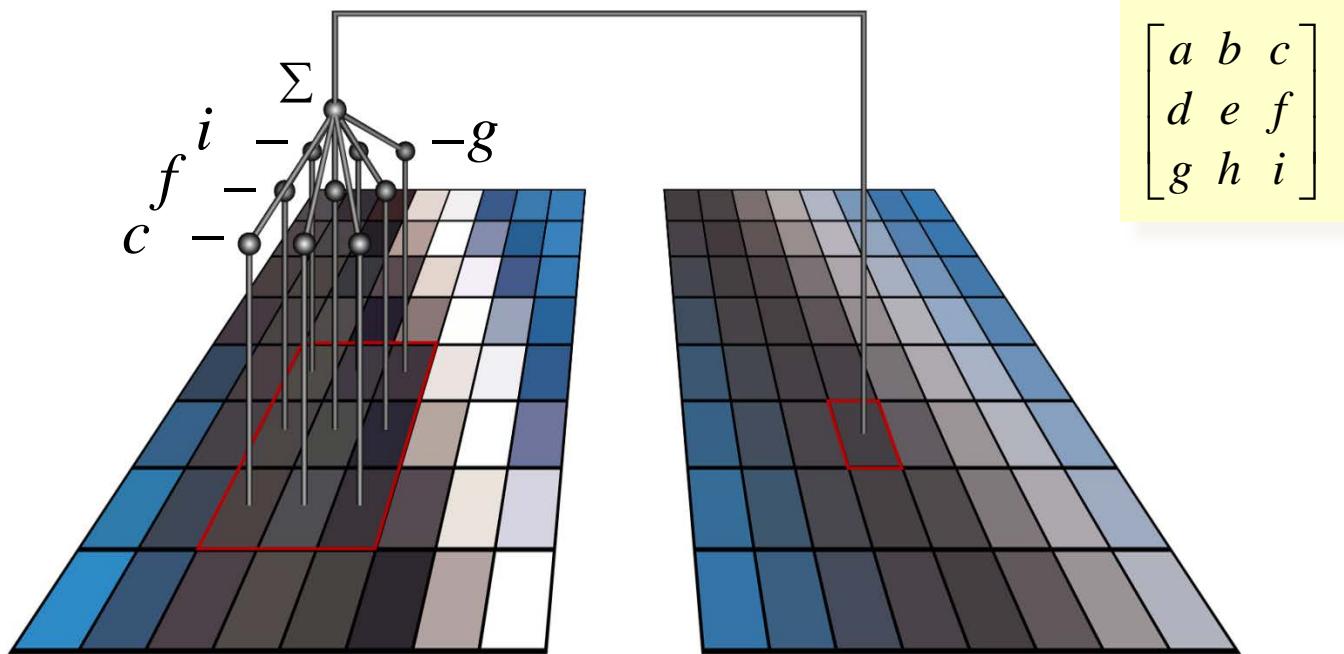
Convolution Masks: Moving Window



Convolution Masks: Moving Window

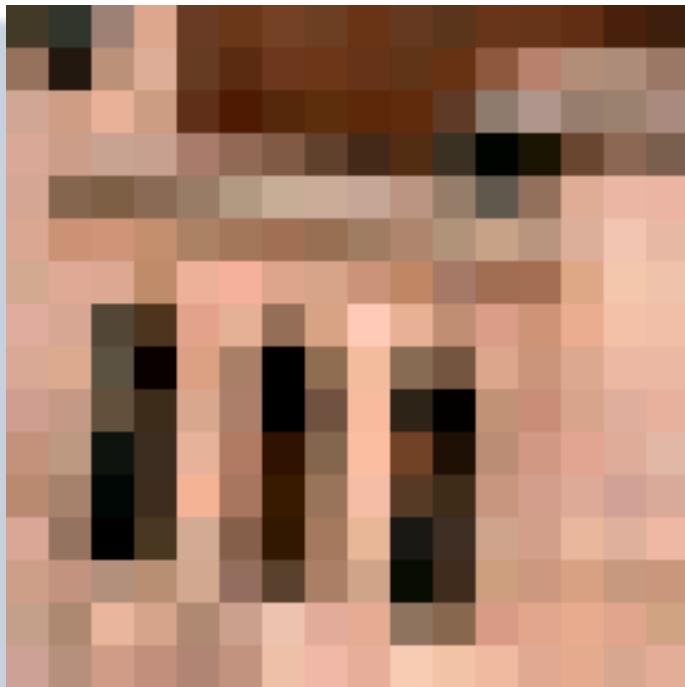


Convolution by Moving Window

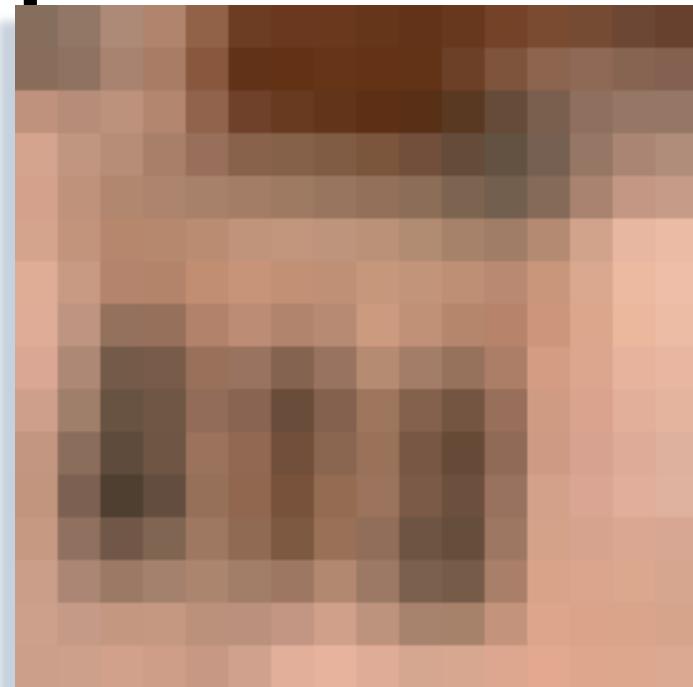


Another example

Moving Window Transform: Example

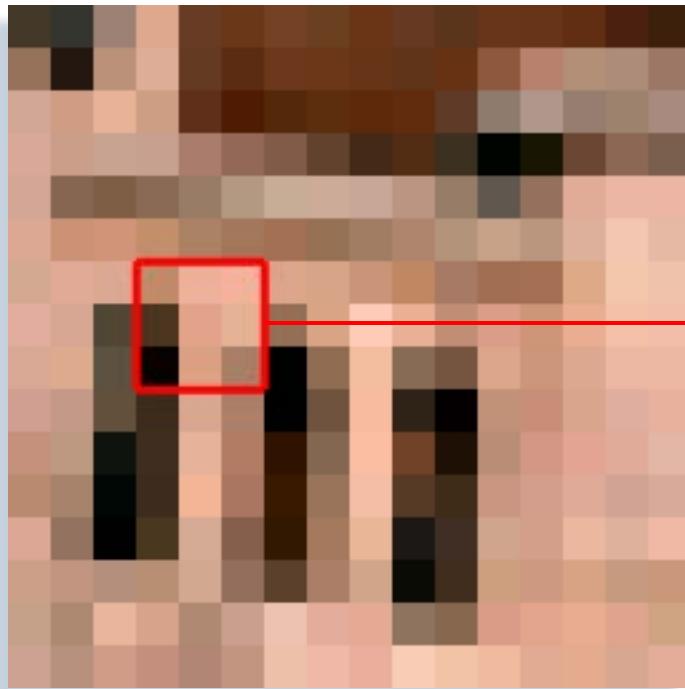


original

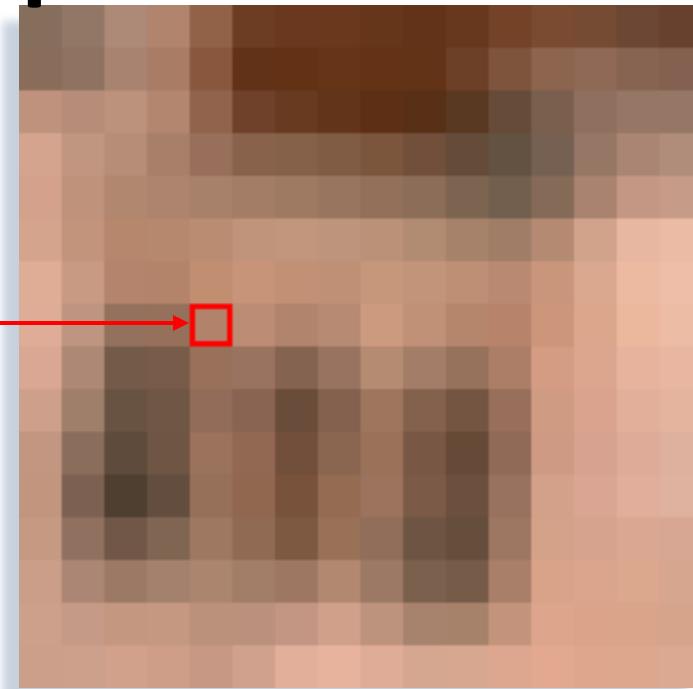


3x3 average

Moving Window Transform: Example

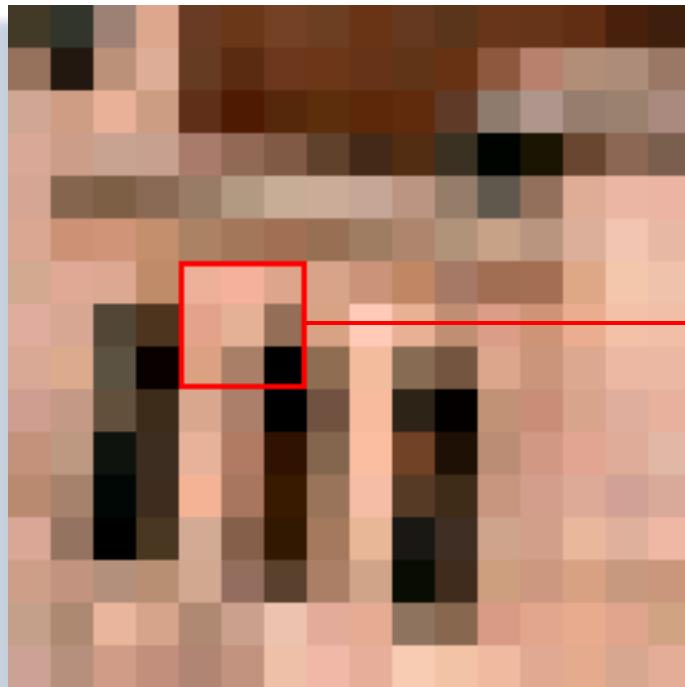


original

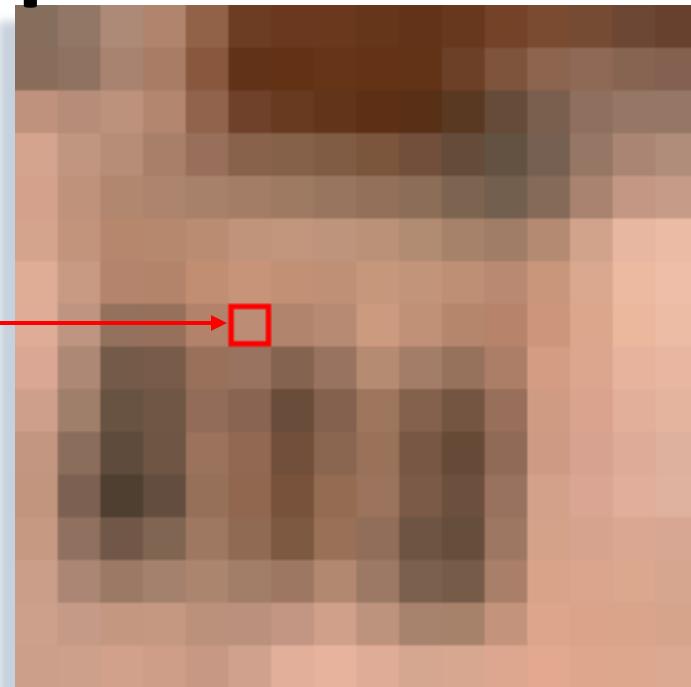


3x3 average

Moving Window Transform: Example

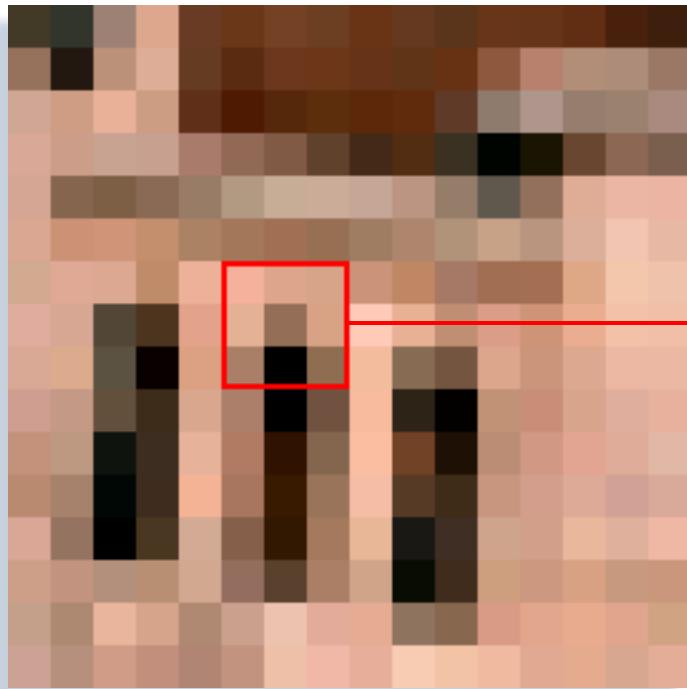


original

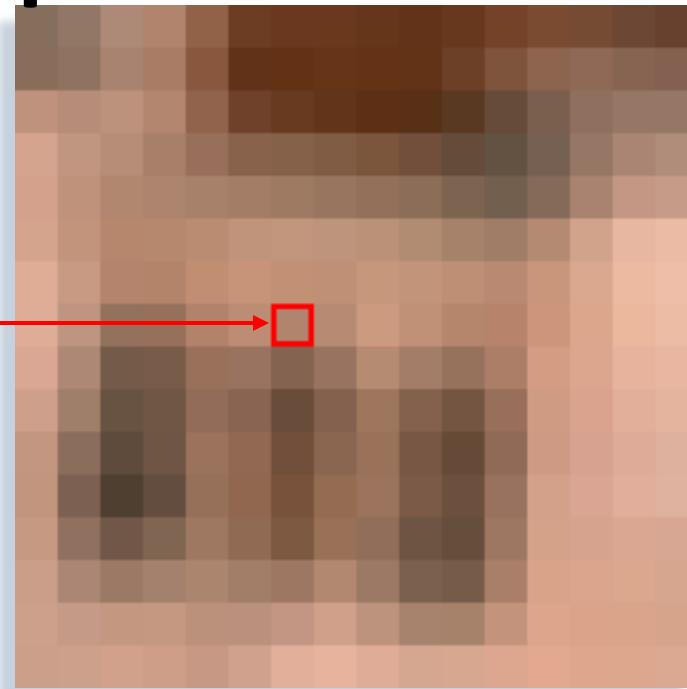


3x3 average

Moving Window Transform: Example

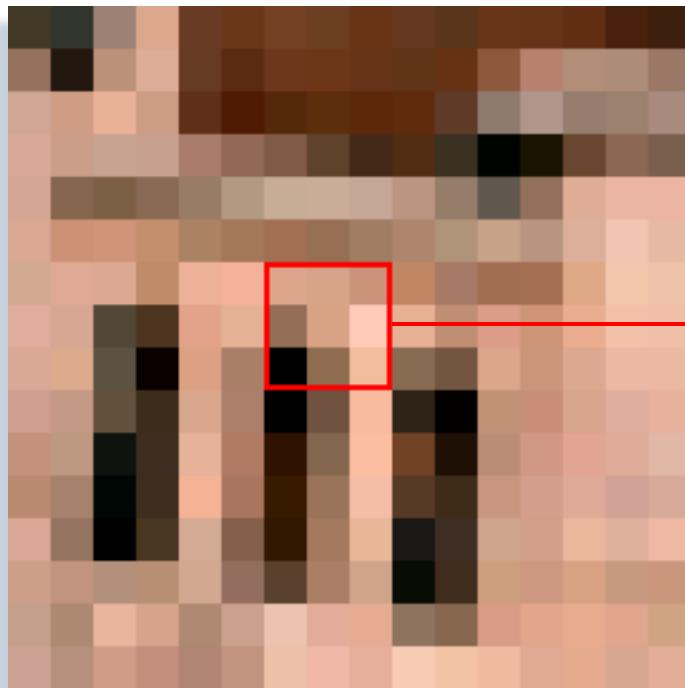


original

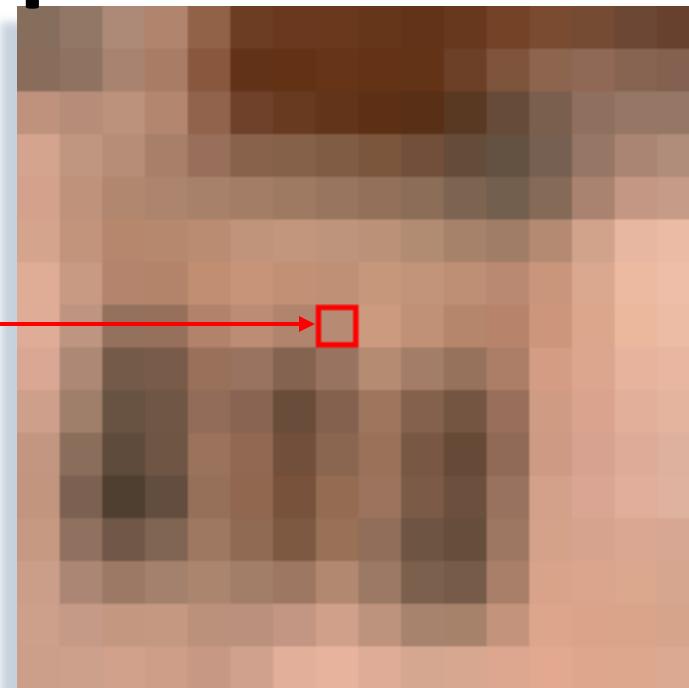


3×3 average

Moving Window Transform: Example

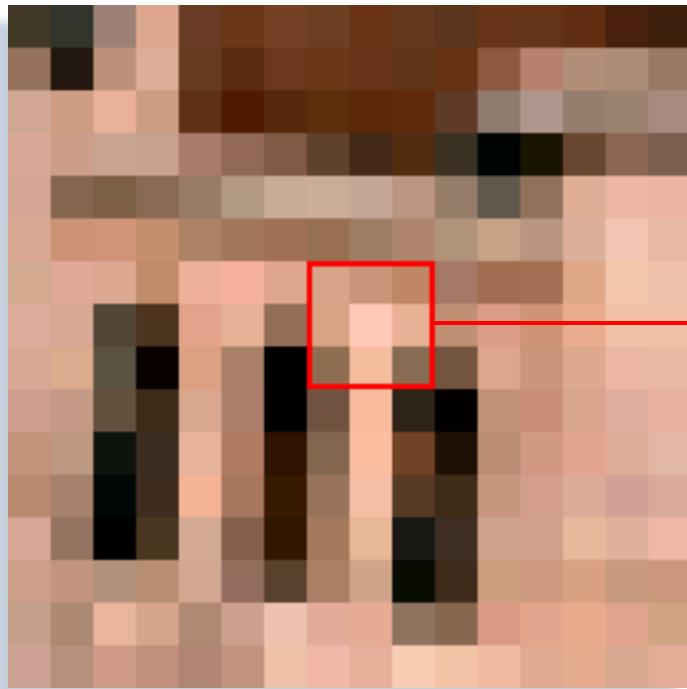


original

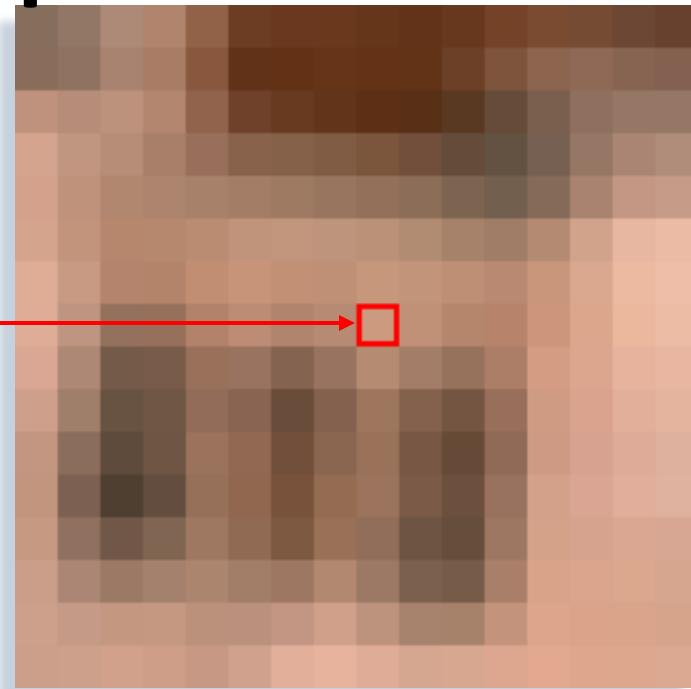


3x3 average

Moving Window Transform: Example

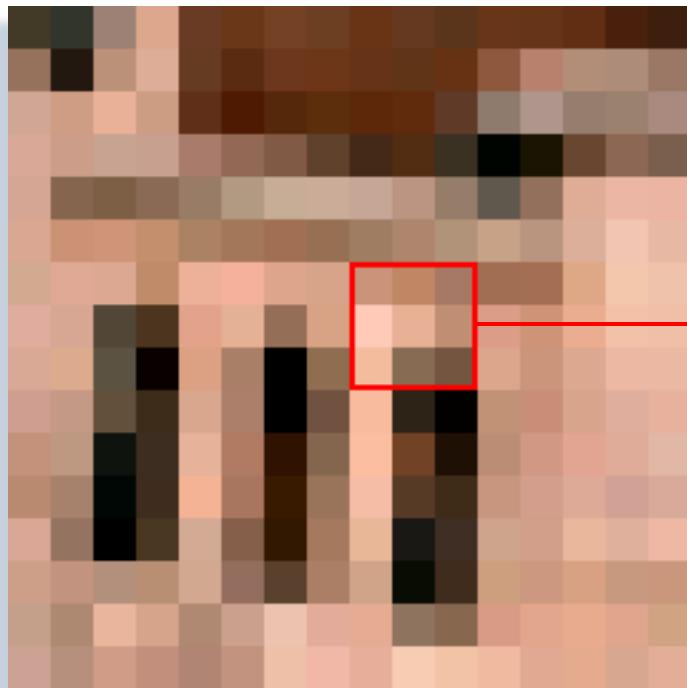


original

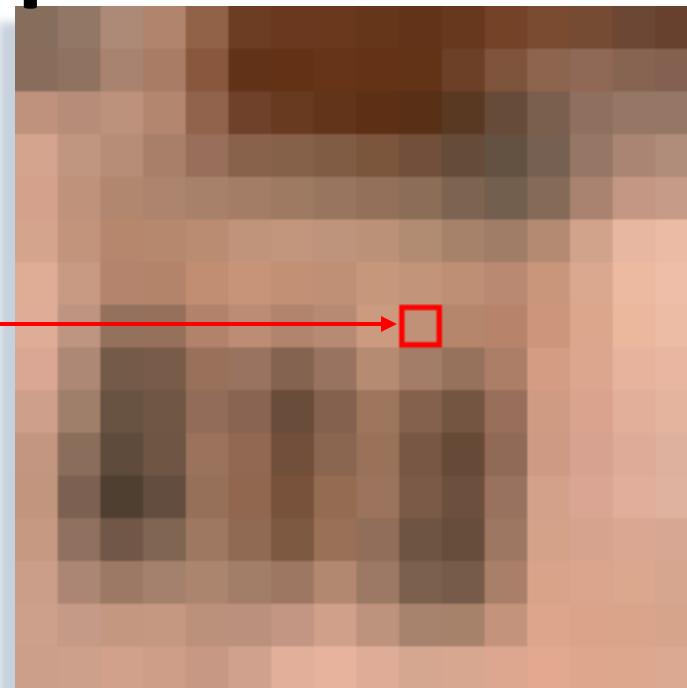


3x3 average

Moving Window Transform: Example



original



3x3 average

The Convolution

- The average filter is a particular example of a more general operation: **Image Convolution**.
- Let **A, B** be images. B is typically smaller than A.
- B is typically called the **mask** or the **kernel**.
- The convolution for 1D signal

$$(A * B)(x) = \sum_i A(i)B(x-i)$$

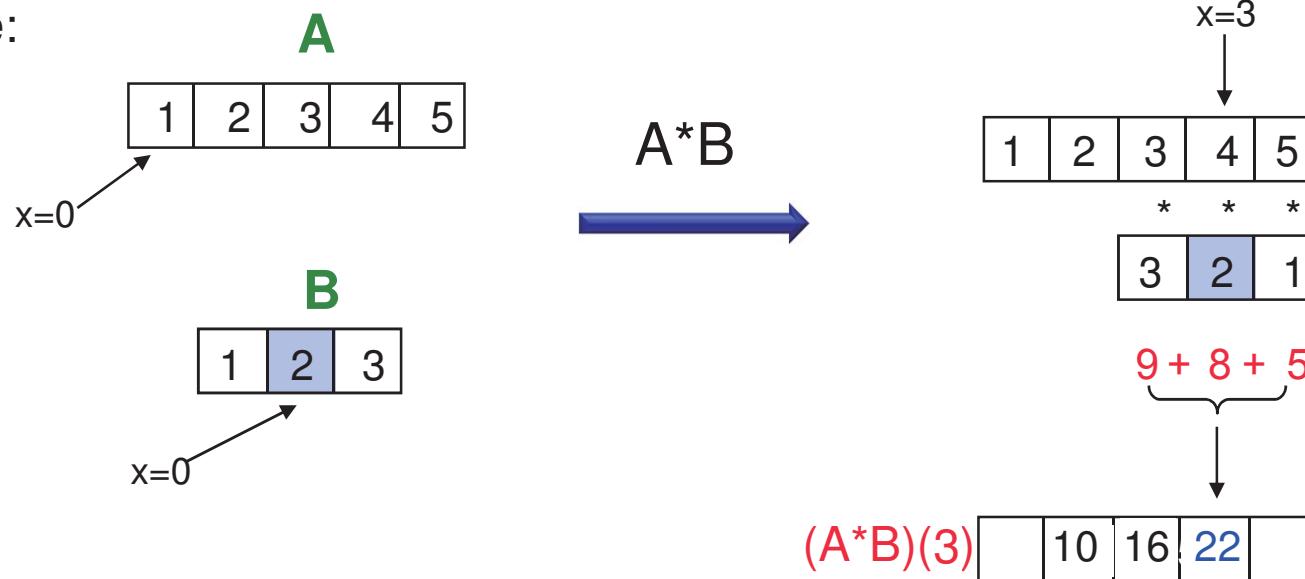
1D Convolution

$$(A * B)(x) = \sum_i A(i)B(x-i)$$

注:

- 1) 对于 $B(i)$, $-1 \leq i \leq 1$
- 2) 对于 $x=3$, $-1 \leq 3-i \leq 1$
 $2 \leq i \leq 4$
- 3) $A(2) - B(1), A(3) - B(0)$
 $A(4) - B(-1)$
 $B(0)$ 反折后与 $A()$ 相关

Example:



X

What happens near the edges?

注意：kernel需要反折—》3 2 1

- Option 1: Zero padding

$$\begin{array}{cccc} 0 & 0 & 0 & \boxed{1 \quad 2 \quad 3 \quad 4 \quad 5} \\ \times & & & \end{array} \quad \begin{array}{cccc} 0 & 0 & 0 & \boxed{1 \quad 2 \quad 3} \\ \times & & & \end{array}$$


Diagram illustrating zero padding. A 5x1 input vector [1, 2, 3, 4, 5] is multiplied by a 3x1 kernel [1, 2, 3]. The result is [4, 10, 16, 22, 22]. The first element 4 is highlighted in blue.

- Option 2: Wrap around

$$\begin{array}{ccccc} 3 & 4 & 5 & \boxed{1 \quad 2 \quad 3 \quad 4 \quad 5} & 1 \quad 2 \quad 3 \quad 4 \quad 5 \\ \times & & & & \end{array}$$

Diagram illustrating wrap-around padding. A 5x1 input vector [1, 2, 3, 4, 5] is multiplied by a 3x1 kernel [1, 2, 3]. The result is [19, 10, 16, 22, 23]. The last element 23 is highlighted in blue.

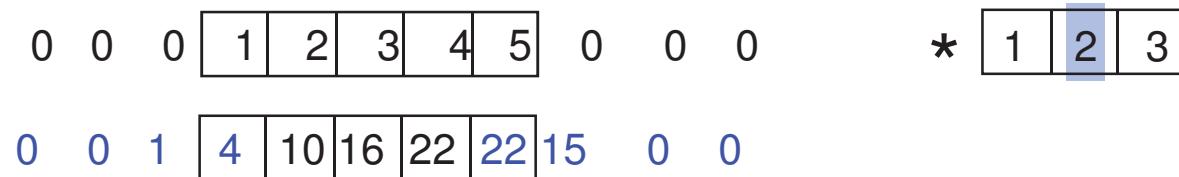
- Option 3: Reflection

$$\begin{array}{ccccccccc} 3 & 2 & 1 & \boxed{1 \quad 2 \quad 3 \quad 4 \quad 5} & 5 & 4 & 3 & 2 \\ \times & & & & & & & \end{array}$$

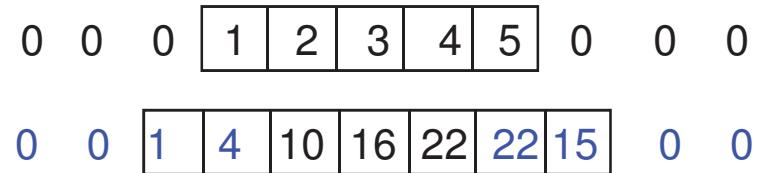

Diagram illustrating reflection padding. A 5x1 input vector [1, 2, 3, 4, 5] is multiplied by a 3x1 kernel [1, 2, 3]. The result is [7, 10, 16, 22, 27]. The last element 27 is highlighted in blue.

What is the length of the result?

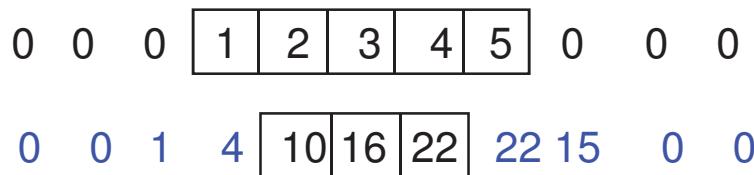
- Option 1: “same” (size A)



- Option 2: “full” (size A + size B - 1)



- Option 3: “valid” (size A – size B + 1)



Examples

Example 1:

$$\begin{array}{c} \textbf{A} \\ \boxed{1 \ 2 \ 3 \ 4 \ 5} \\ x=0 \end{array} * \begin{array}{c} \textbf{B} \\ \boxed{0 \ 1 \ 0} \\ x=0 \end{array} = \begin{array}{c} \textbf{A}^*\textbf{B} \\ \boxed{1 \ 2 \ 3 \ 4 \ 5} \end{array}$$

Example 2:

$$\begin{array}{c} \textbf{A} \\ \boxed{1 \ 2 \ 3 \ 4 \ 5} \end{array} * \begin{array}{c} \textbf{B} \\ \boxed{0 \ 2 \ 0} \end{array} = \begin{array}{c} \textbf{A}^*\textbf{B} \\ \boxed{2 \ 4 \ 5 \ 8 \ 10} \end{array}$$

Example 3:

$$\begin{array}{c} \textbf{A} \\ \boxed{1 \ 2 \ 3 \ 4 \ 5} \end{array} * \begin{array}{c} \textbf{B} \\ \boxed{1/3 \ 1/3 \ 1/3} \end{array} = \begin{array}{c} \textbf{A}^*\textbf{B} \\ \boxed{1 \ 2 \ 3 \ 4 \ 3} \end{array}$$

- Why should we flip the mask before the convolution?

With reflection:

$$\begin{array}{c} \cdots [1 \ 2 \ 3] \cdots * [0 \ 1 \ 0] \cdots = [0 \ 1 \ 2 \ 3 \ 0] \cdots \\ \downarrow \qquad \qquad \qquad \downarrow \qquad \qquad \qquad \downarrow \\ \cdots [0 \ 1 \ 0] \cdots * [1 \ 2 \ 3] \cdots = [0 \ 1 \ 2 \ 3 \ 0] \cdots \end{array}$$

Without reflection:

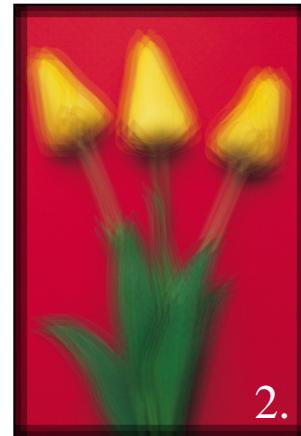
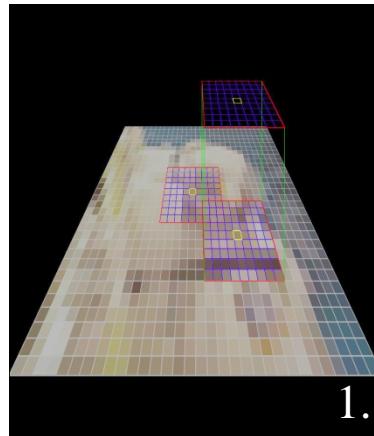
$$\begin{array}{c} \cdots [1 \ 2 \ 3] \cdots * [0 \ 1 \ 0] \cdots = [0 \ 1 \ 2 \ 3 \ 0] \cdots \\ \downarrow \qquad \qquad \qquad \downarrow \qquad \qquad \qquad \downarrow \\ \cdots [0 \ 1 \ 0] \cdots * [1 \ 2 \ 3] \cdots = [0 \ 3 \ 2 \ 1 \ 0] \cdots \end{array}$$

- Reflection is needed so that convolution is commutative:

$$A^*B = B^*A$$

Three ways to compute a convolution

1. Moving window transform as just shown.
2. Shift multiply add.
3. Fourier transform.



Shift-Multiply-Add Approach

- The image is copied 1 time for each element in the convolution mask.
- Each copy is shifted relative to the original by the displacement of its associated mask element.
- Each copy is multiplied by the value of its associated mask element.
- The set of shifted and multiplied images is summed pixel wise.

Convolution by an Impulse

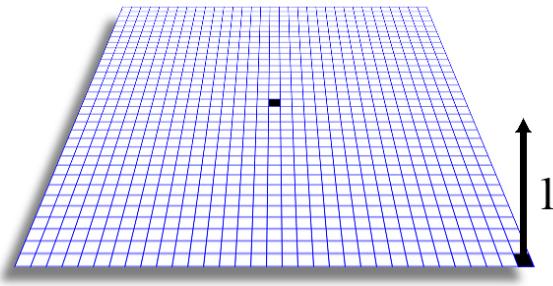
An *impulse* is a digital image, that has a single pixel with value 1; all others have value zero. An impulse at location (ρ, χ) is represented by:

$$\delta(r - \rho, c - \chi) = \begin{cases} 1, & \text{if } r = \rho \text{ and } c = \chi \\ 0, & \text{otherwise} \end{cases}$$

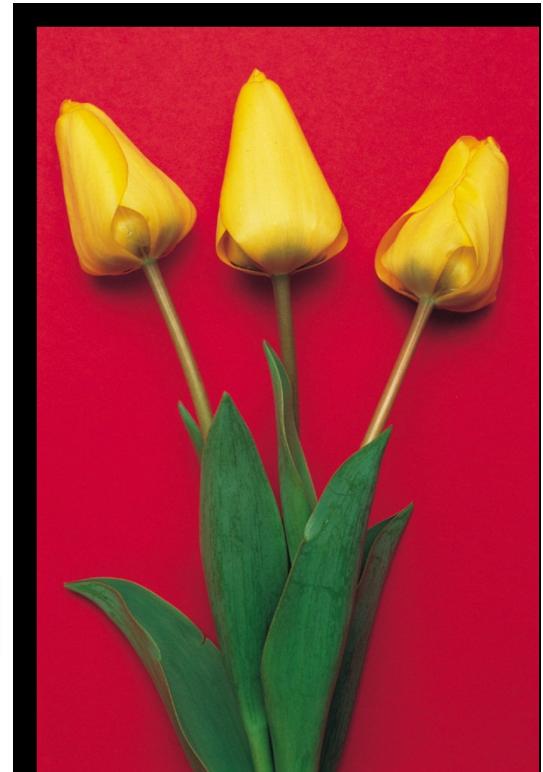
If an image is convolved with an impulse of weight w at location (ρ, χ) , then the image is multiplied by w and shifted in location down by ρ pixels and to the right by χ pixels.

$$[\mathbf{I} * w\delta(r - \rho, c - \chi)](r, c) = w\mathbf{I}(r - \rho, c - \chi).$$

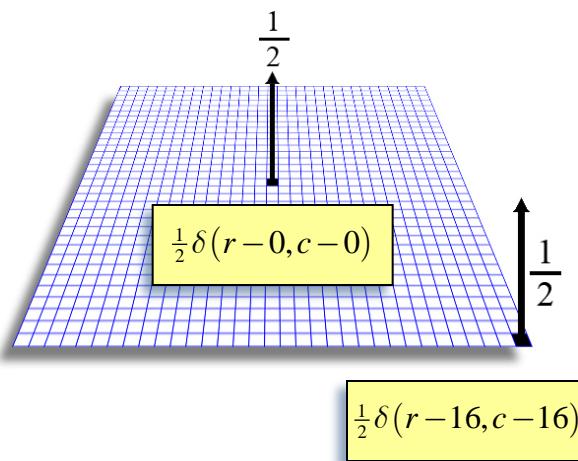
Convolution by an Impulse



Shifted down and to
the right by 16 pixels.



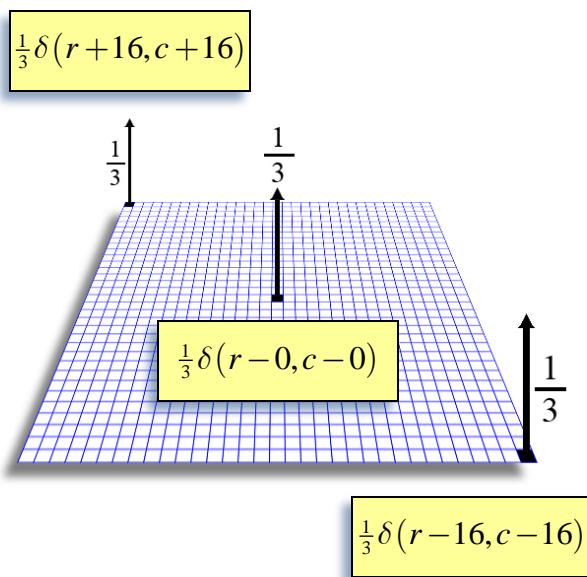
Convolution by Two Impulses



Two copies, one moved,
one not moved, averaged.



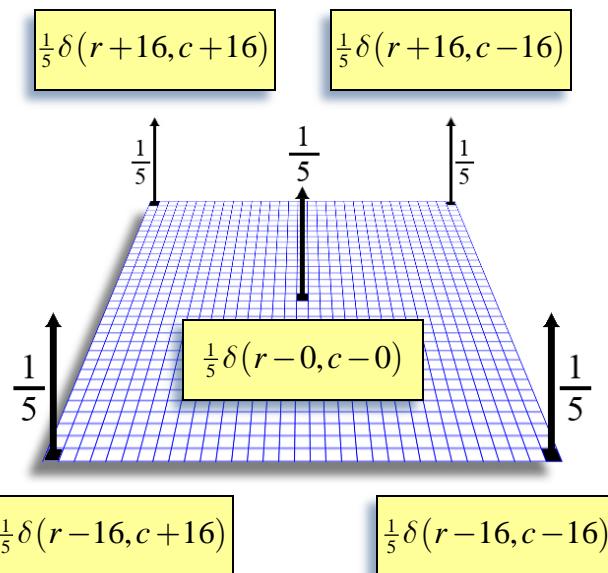
Convolution by Three Impulses



Three copies, two moved,
one not moved, averaged.

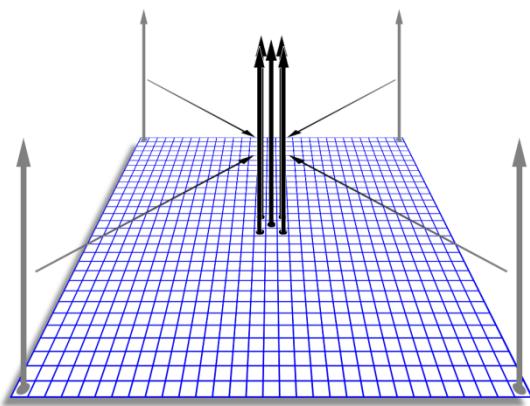


Convolution by Five Impulses



Five copies, four moved,
one not moved, averaged.

Convolution by Five Impulses

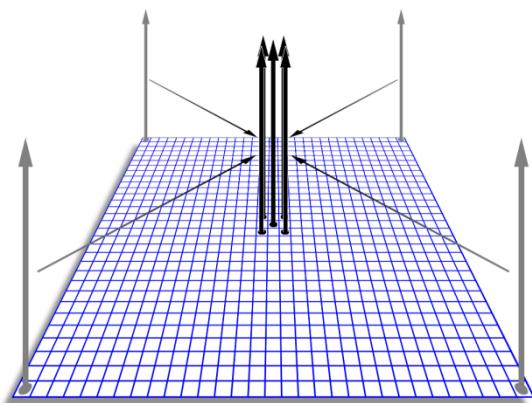
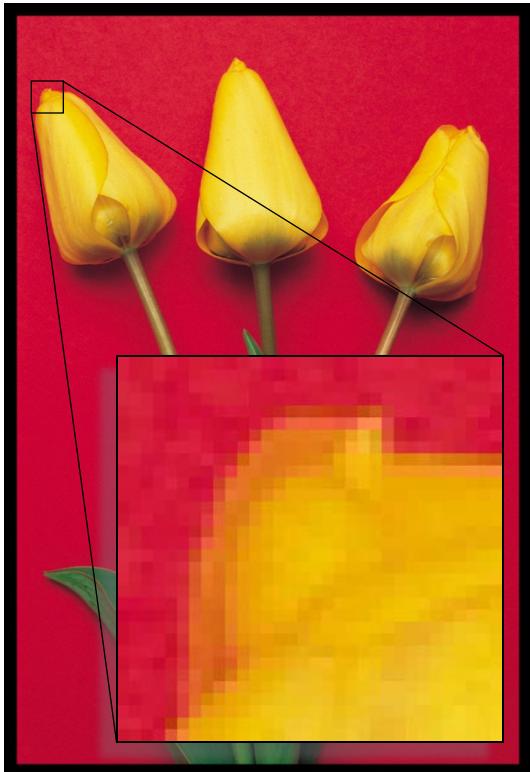


Moved adjacent to each other, the convolution becomes a blurring filter.

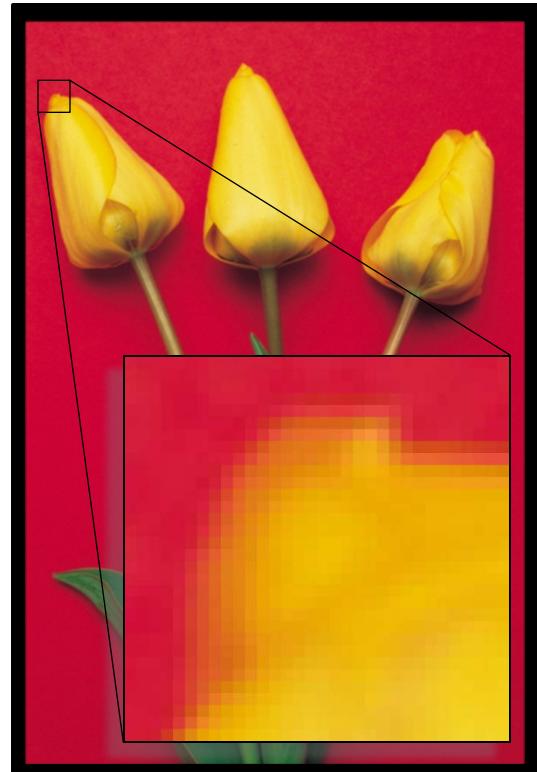


又一个例子，用于仿真图像运动模糊模型

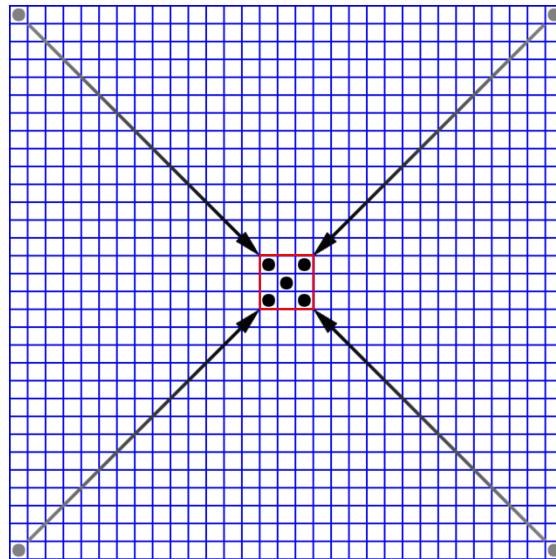
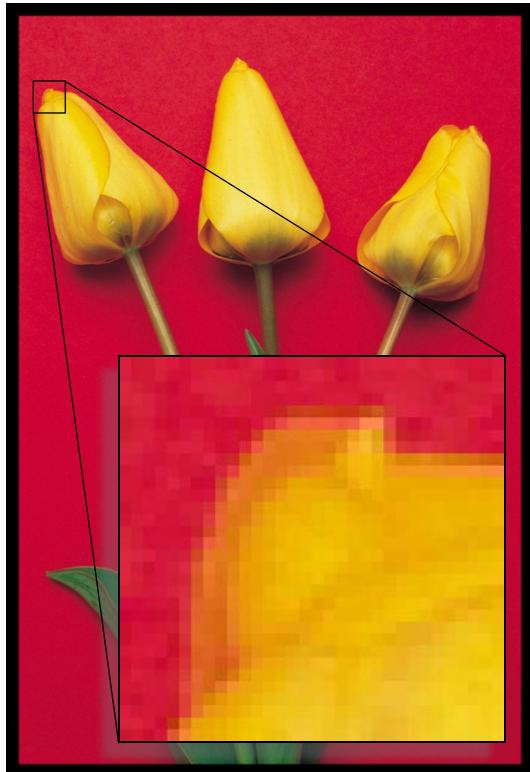
Convolution by Five Impulses



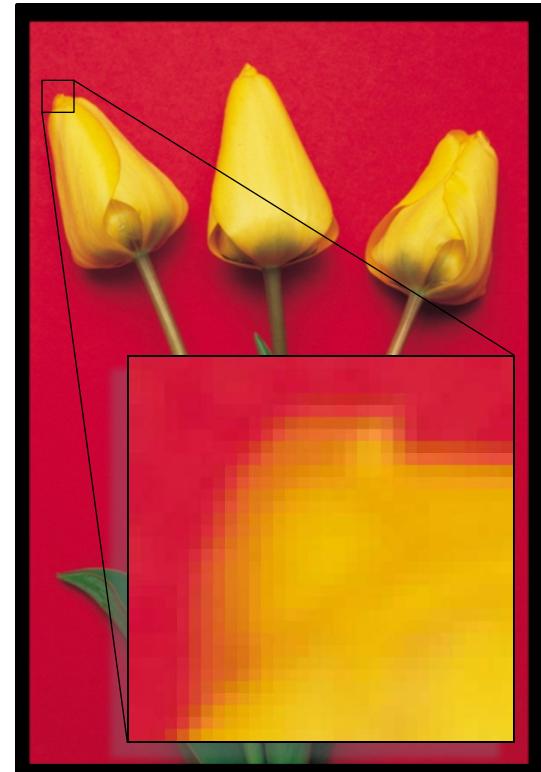
The impulses become values
in a 3×3 neighborhood.



Convolution by Five Impulses

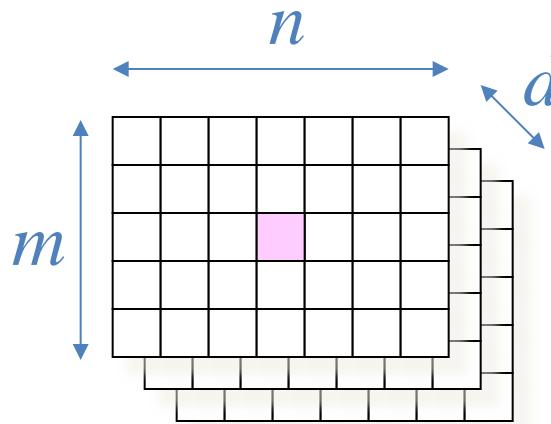


The convolution mask has five elements at 1/5 and four at 0.



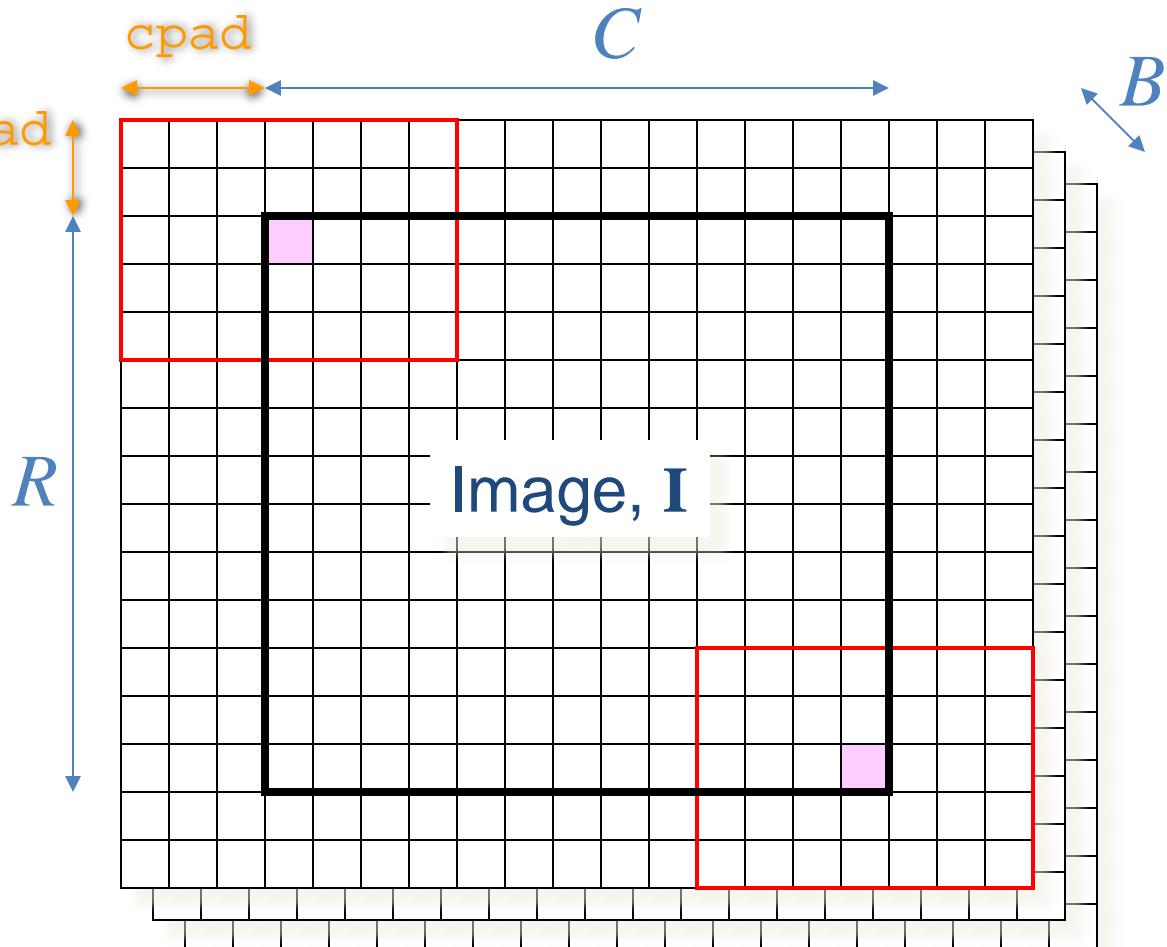
Zero Padding an Image for Convolution: Variable Names.

weight matrix, h



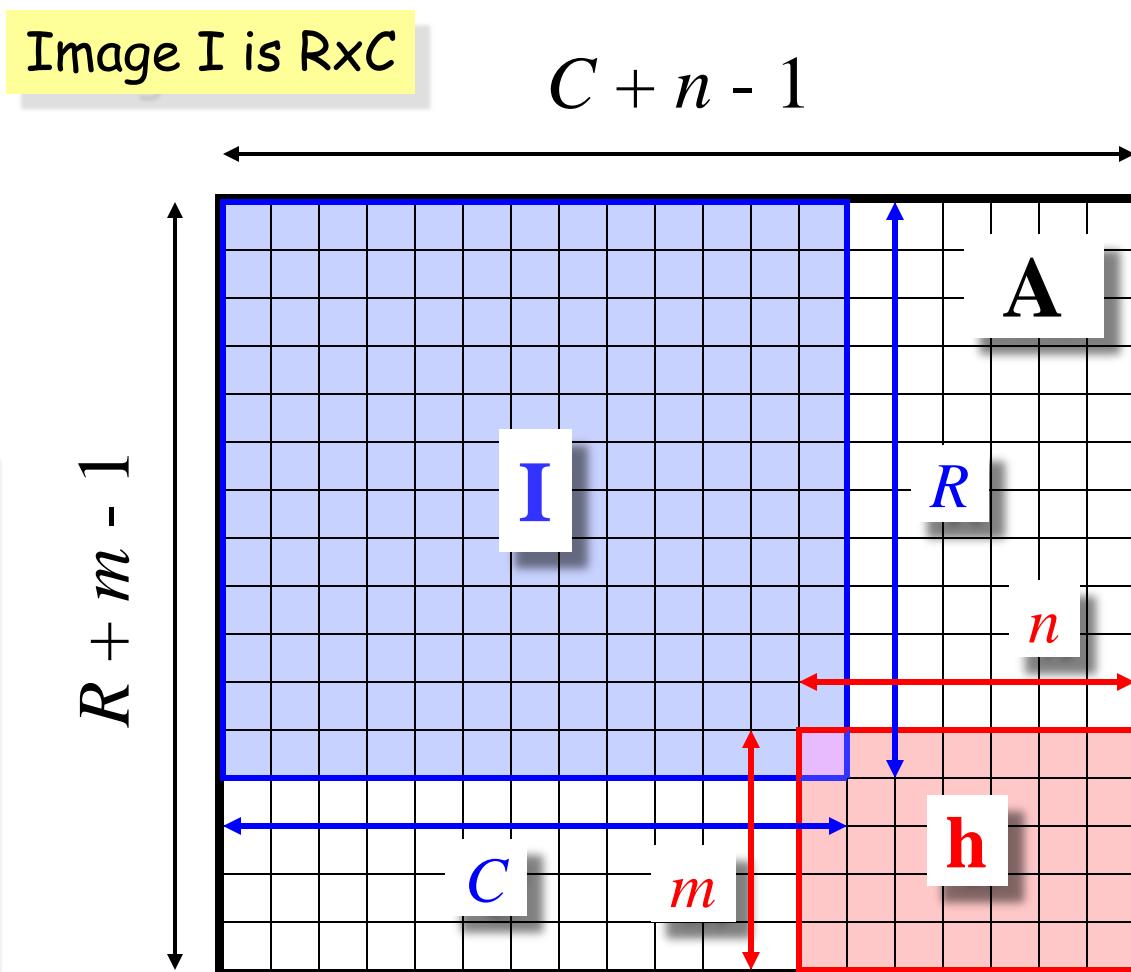
```
cpad = floor( n / 2 )
rpad = floor( m / 2 )
hcorig = cpad + 1
hrorig = rpad + 1
```

↓
hc和hr原点，
也是图像卷积模板的原点



Convolution by Copying and Shifting the Image

To use the image shift-multiply-accumulate algorithm, create an accumulator image, A , that is $R+m-1$ rows by $C+n-1$ columns

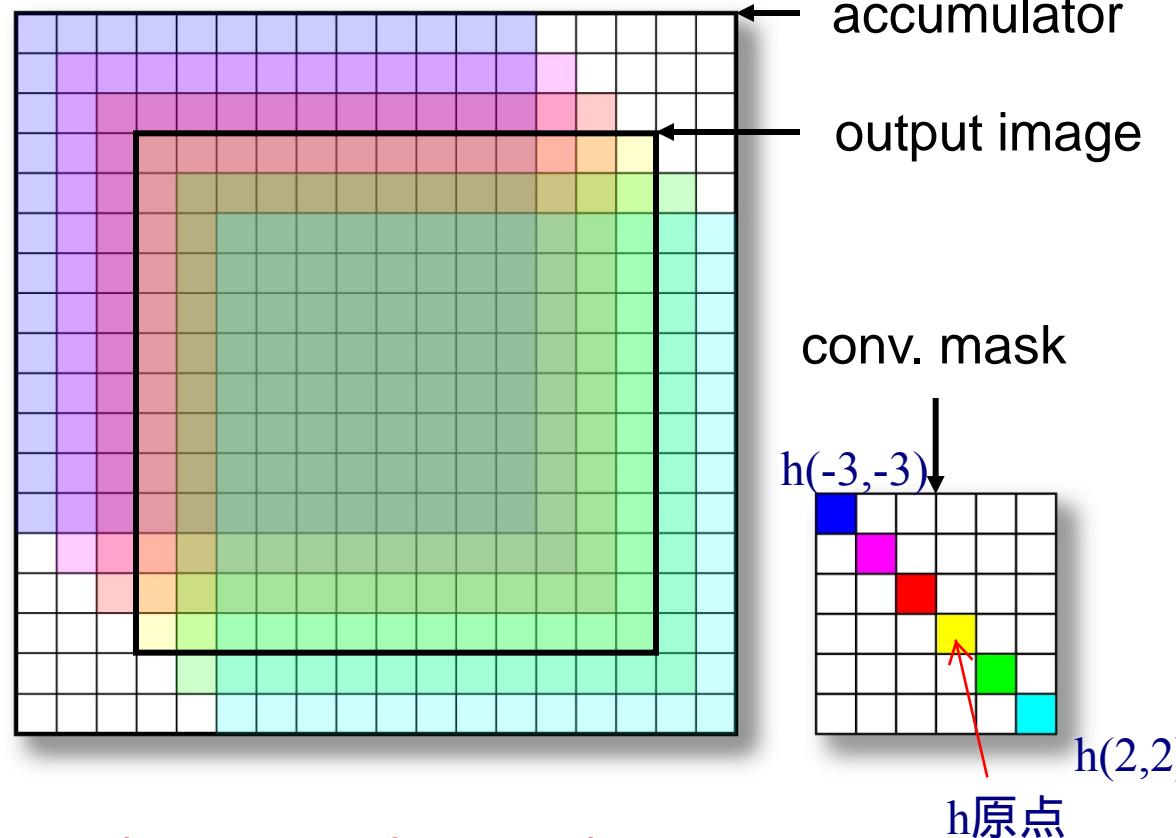


Convolution by Copying, Multiplying, and Shifting the Image

13x13 image
convolved by
6x6 mask.

Image is constant;
mask has only 6
nonzero values all
on the diagonal.

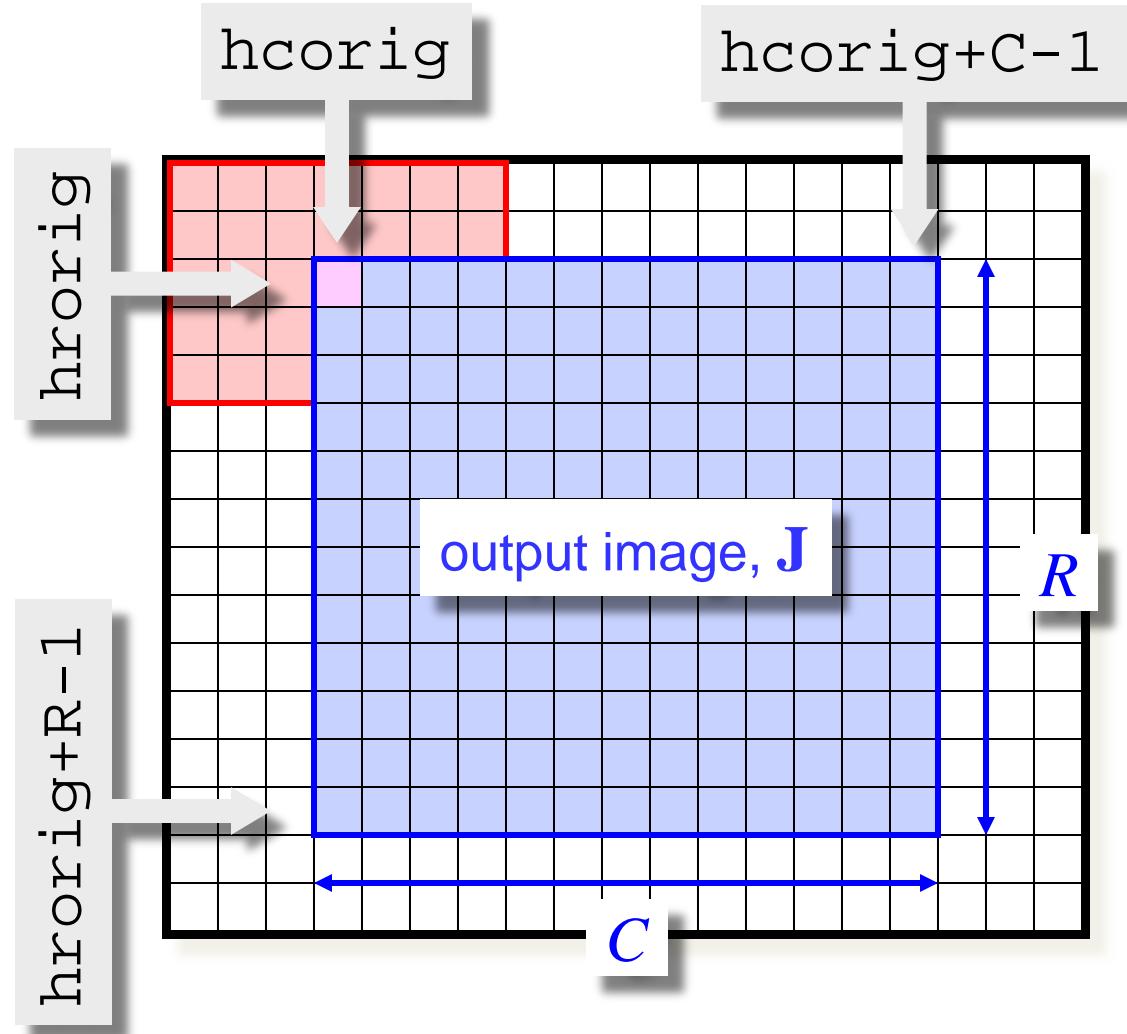
Image is shifted to
mask location,
multiplied by value,
and accumulated.



h:6x6 , cpad=3, rpad=3
模板原点 (3+1,3+1)

Convolution by Copying and Shifting the Image

When done, copy the output image from the accumulator starting at $(hrorig, hcorig)$ and ending at $(hrorig+R-1, hcorig+C-1)$

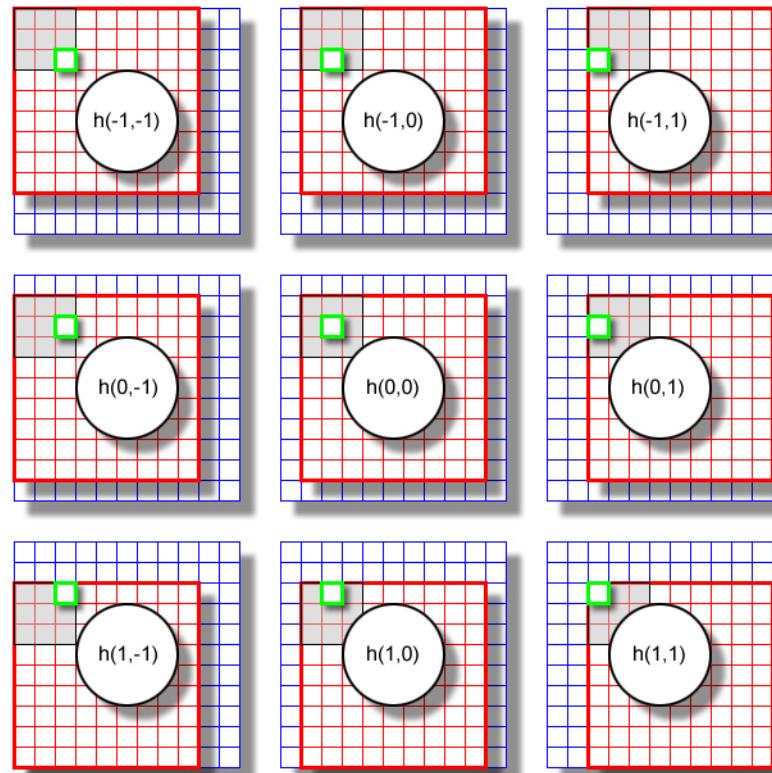


Convolution by Copying, Multiplying, and Shifting the Image

For each element $\mathbf{h}(r_h, c_h)$ in weight matrix, \mathbf{h} , image \mathbf{I} is copied into a zero-padded image, \mathbf{P} , starting at (r_h, c_h) .

Each \mathbf{P} is multiplied by the corresponding weight, $\mathbf{h}(r_h, c_h)$.

All the \mathbf{P} images are summed pixel-wise then divided by the sum of the elements of \mathbf{h} . The result is cropped out of the center of the accumulated \mathbf{Ps} .



original image, \mathbf{I}

padded image, \mathbf{P}

effective neighborhood

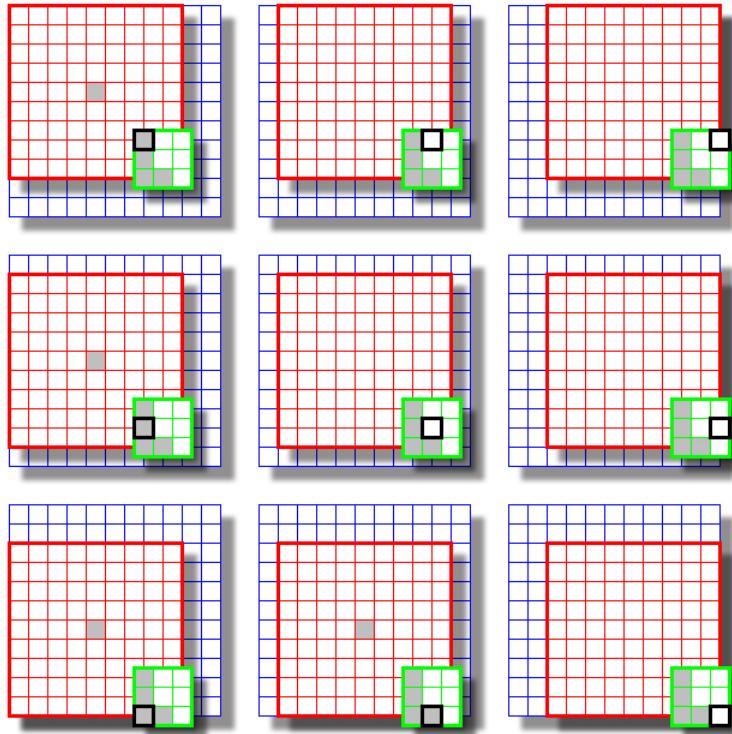
$\mathbf{h}(-1,-1)$	$\mathbf{h}(-1,0)$	$\mathbf{h}(-1,1)$
$\mathbf{h}(0,-1)$	$\mathbf{h}(0,0)$	$\mathbf{h}(0,1)$
$\mathbf{h}(1,-1)$	$\mathbf{h}(1,0)$	$\mathbf{h}(1,1)$

weight matrix

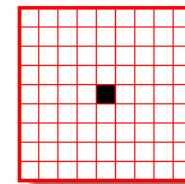
aligned pixels to be summed

weight for image

Convolution by Copying, Multiplying, and Shifting the Image



original image, \mathbf{I}
effective neighborhood
padded image, \mathbf{P}

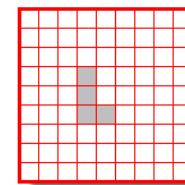


The original image has a black impulse at the center and zeros (white) elsewhere.

*



=



The weight matrix has a gray 'L' at its left and zeros (white) elsewhere.

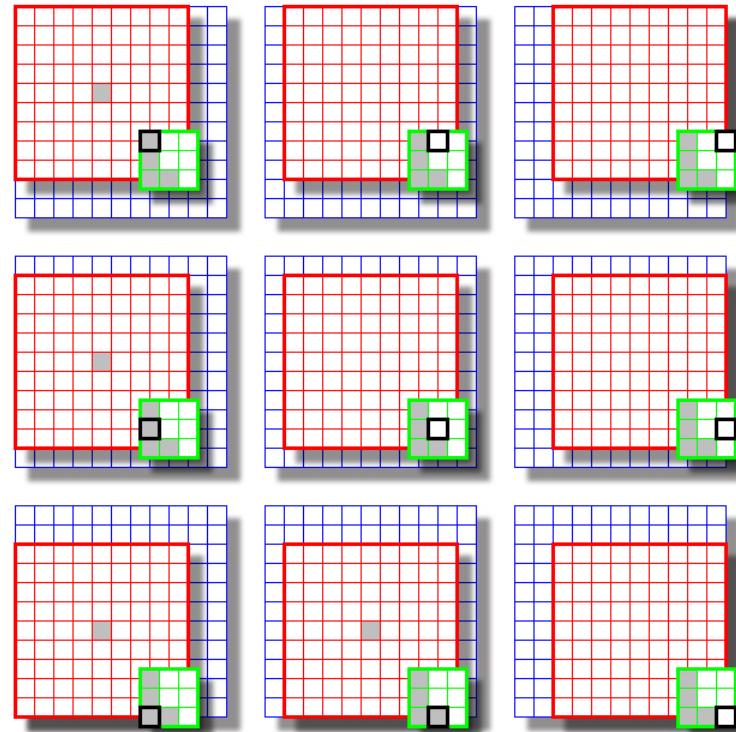
The resulting image has a copy of the weight matrix pegged to the impulse location.

In the result, the origin of the weight matrix coincides with the original location of the impulse.

Convolution by Copying, Multiplying, and Shifting the Image

Each copy of the (entire) image is multiplied by the value of the weight matrix in black square (here, white = 0) before being accumulated (pixelwise) in the padded image

The position of the black square relative the center of the weight matrix indicates the shift of the original image relative to the middle of the padded image.

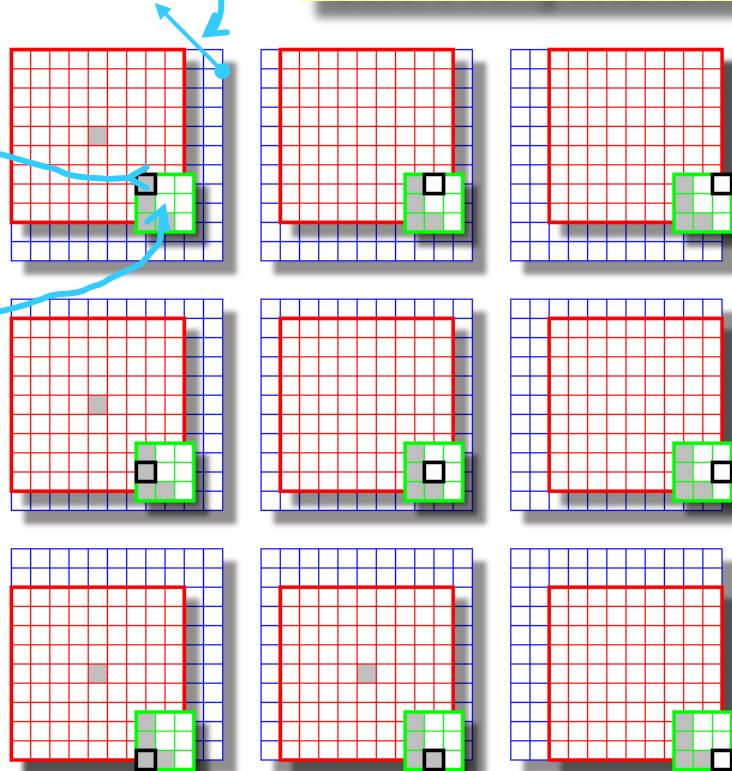


$$\begin{matrix} & \ast \\ \begin{matrix} & o \\ \odot & \end{matrix} & = \\ & \end{matrix}$$

In this image, only the pixel in the center is nonzero so only it shows a result when the image is multiplied by a nonzero value

Convolution by Copying, Multiplying, and Shifting the Image

The position of the black square relative the center of the weight matrix indicates the shift of the original image relative to the middle of the padded image.

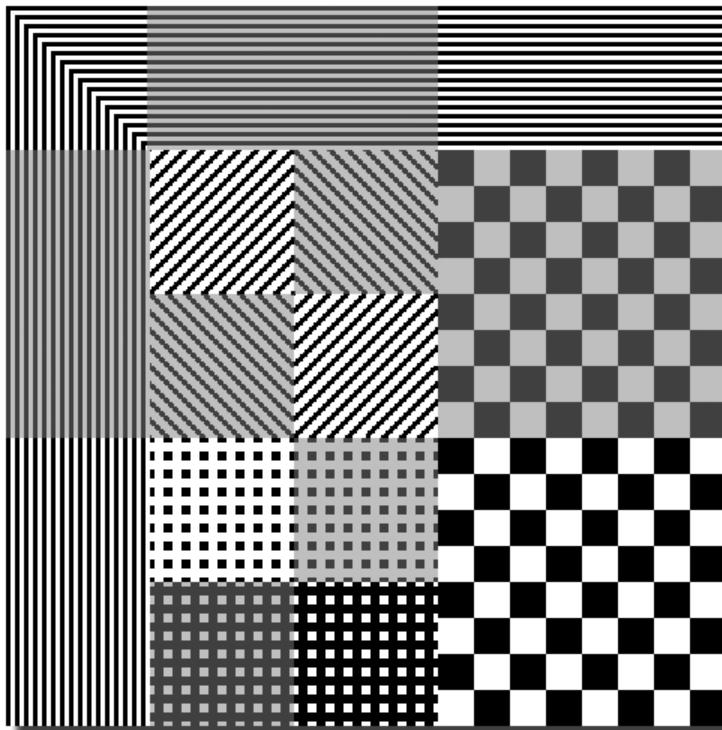


Each copy of the (entire) image is multiplied by the value of the weight matrix in black square (here, white = 0) before being accumulated (pixelwise) in the padded image

$$\begin{matrix} & \text{---} \\ & * \\ \begin{matrix} & & 0 \\ & & \end{matrix} & = \\ & & \text{---} \end{matrix}$$

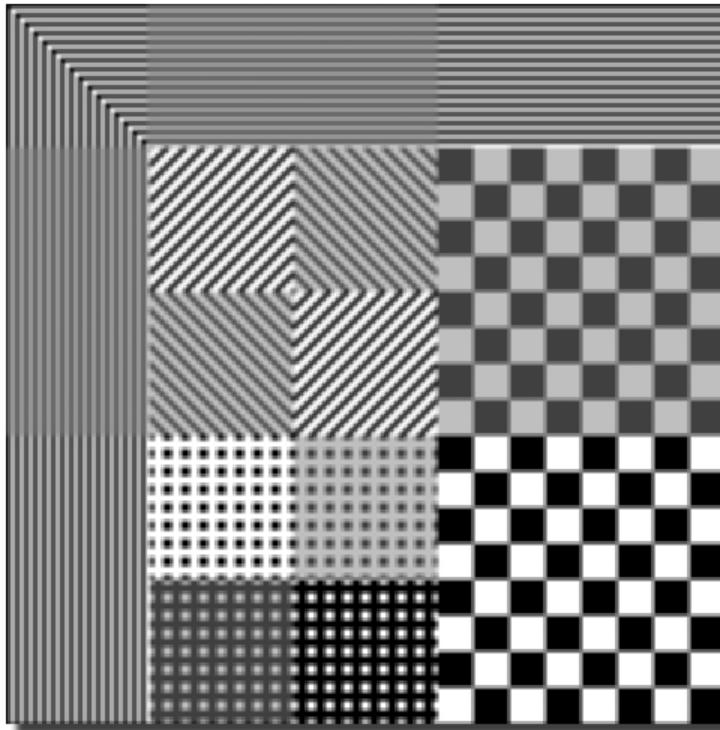
In this image, only the pixel in the center is nonzero so only it shows a result when the image is multiplied by a nonzero value

Convolution Examples: Original Images



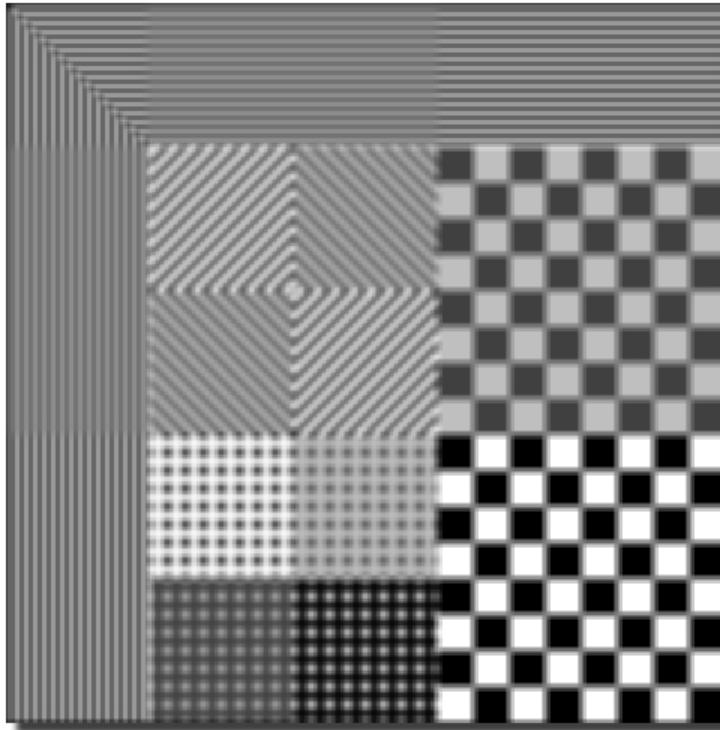
$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Convolution Examples: 3×3 Blur



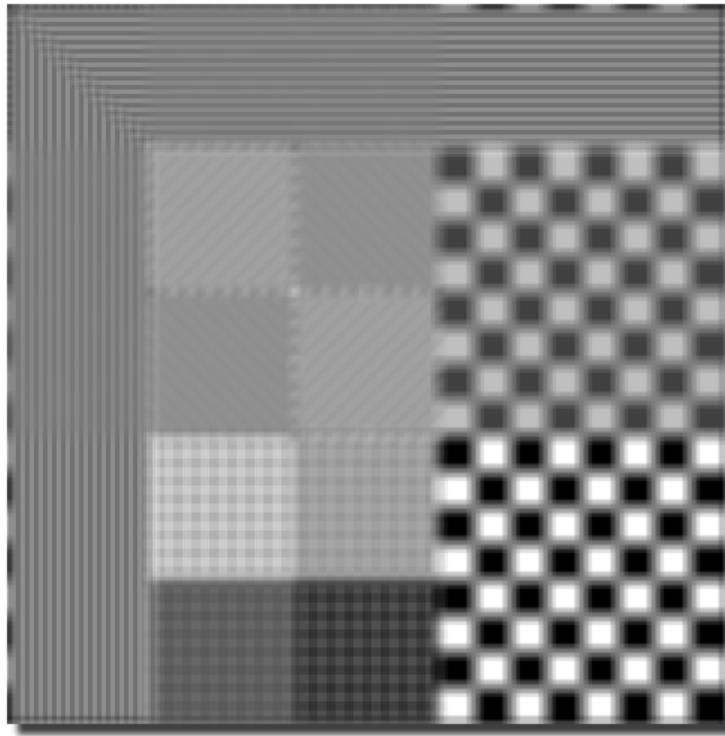
$$\frac{1}{25} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

Convolution Examples: 5×5 Blur



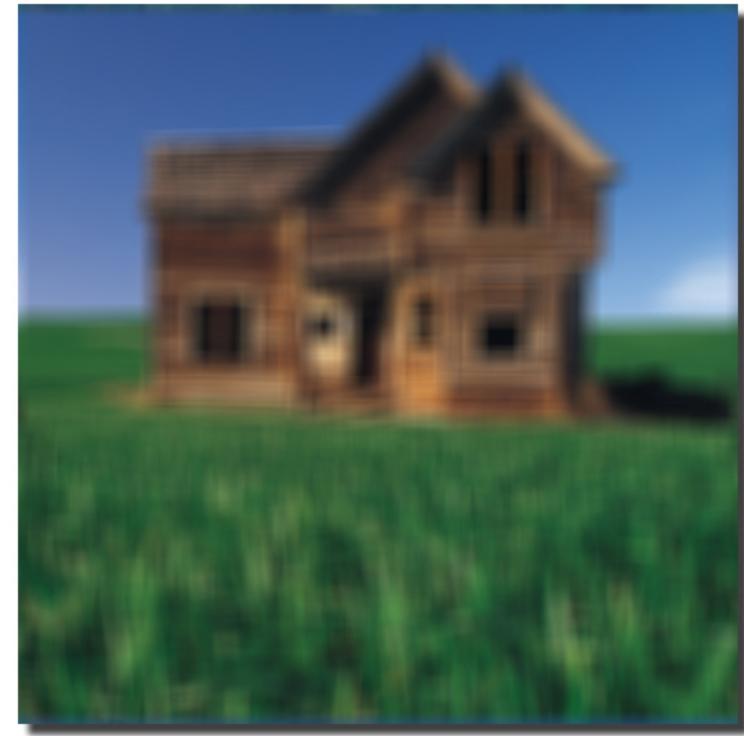
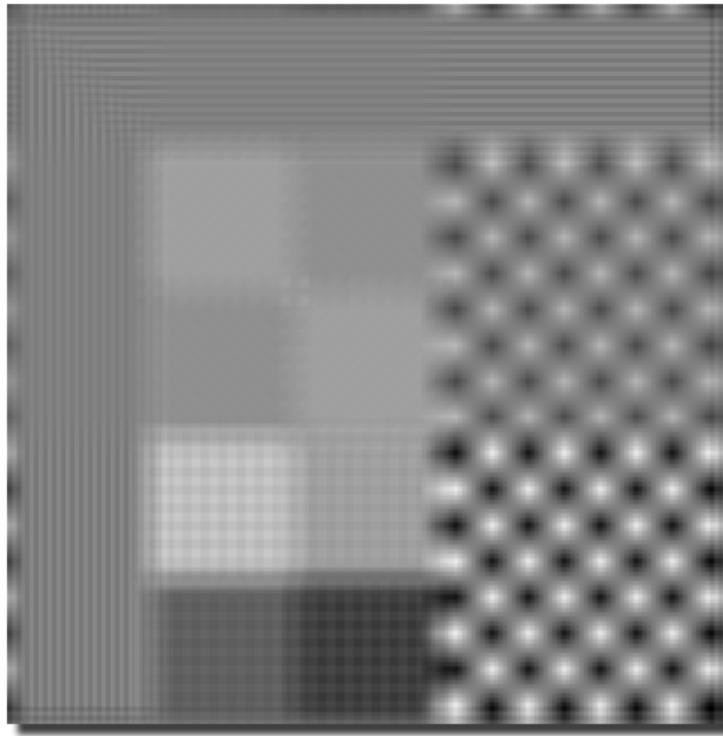
$$\frac{1}{81} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

Convolution Examples: 9×9 Blur

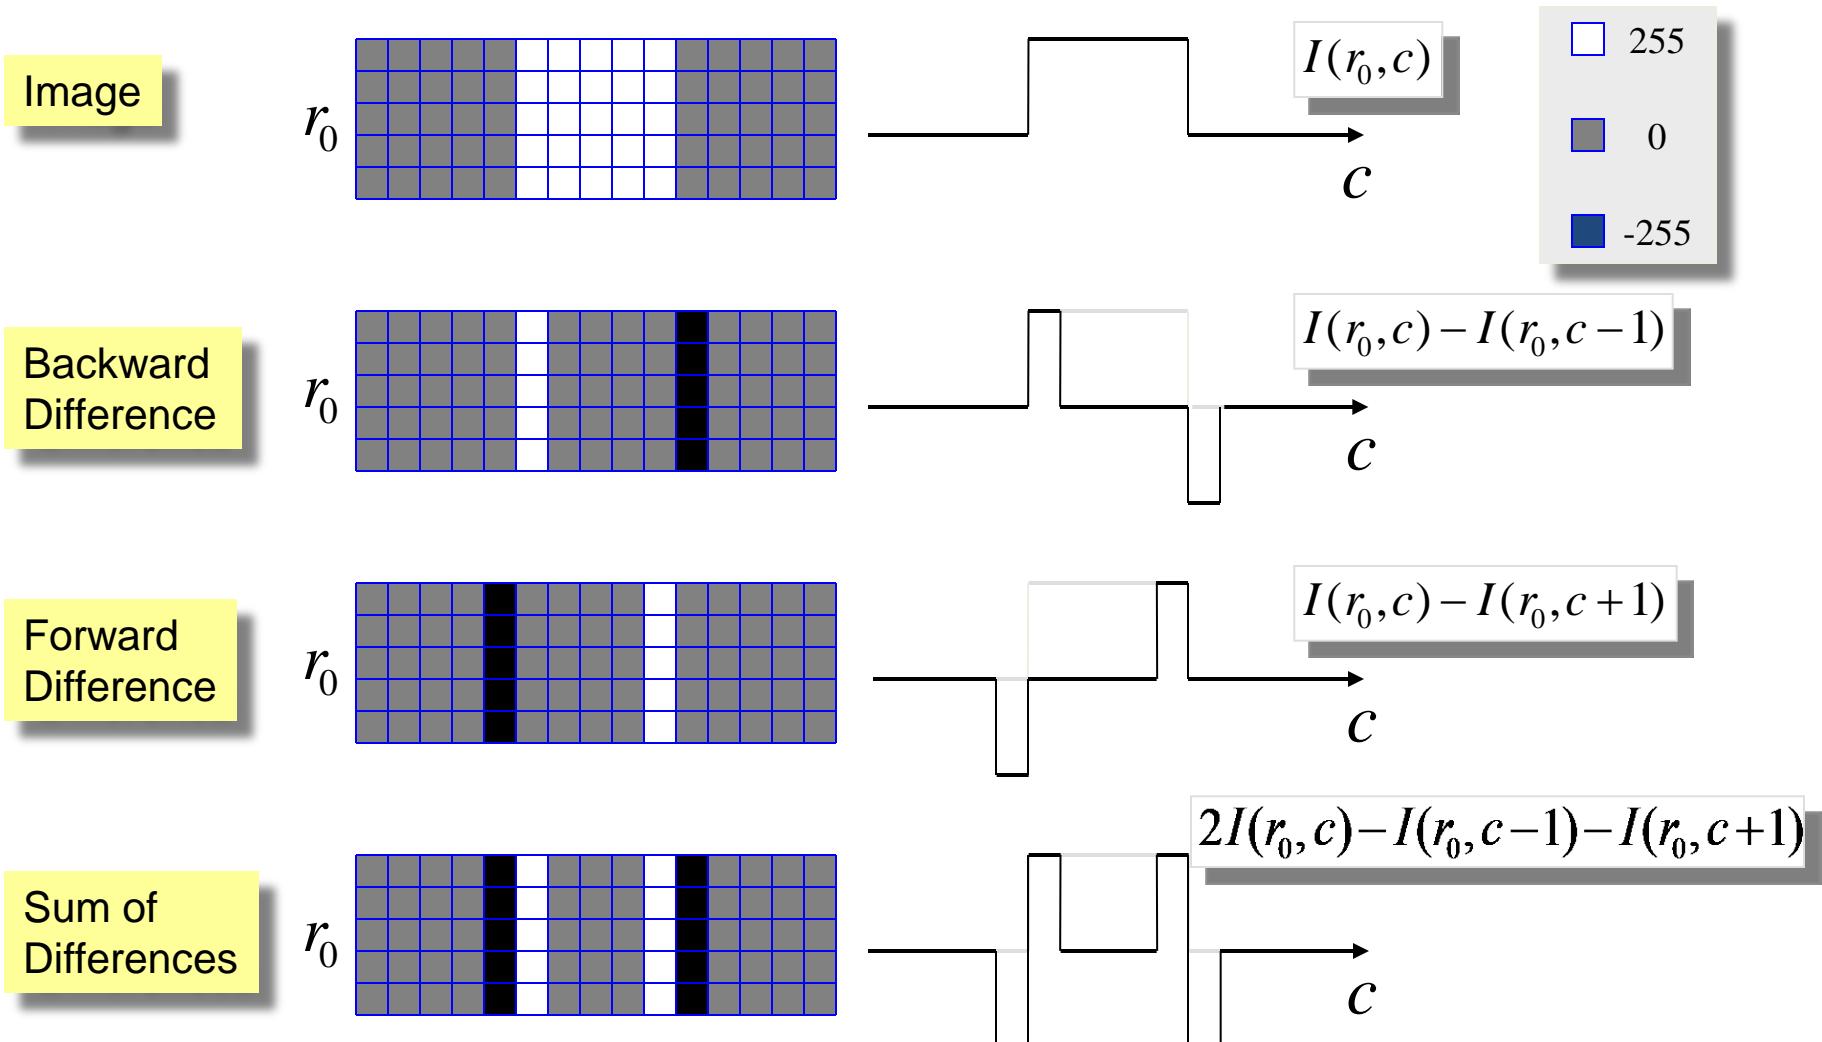


$$\frac{1}{289} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

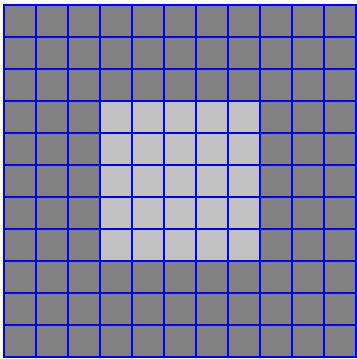
Convolution Examples: 17×17 Blur



Vertical Edge Detection



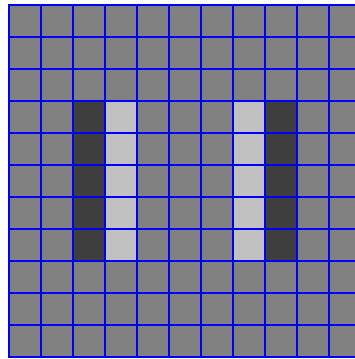
Symmetric Edge Detection (如何构造检测模板)



$I(r,c)$

用到3x3邻域

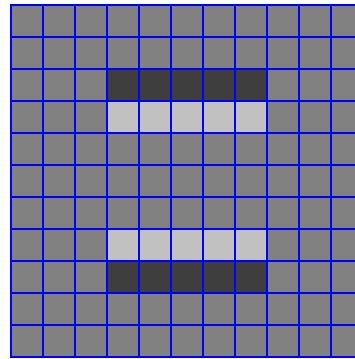
- 510
- 255
- 0
- 255



$$2I(r,c) - I(r,c-1) - I(r,c+1)$$

垂直边缘检测

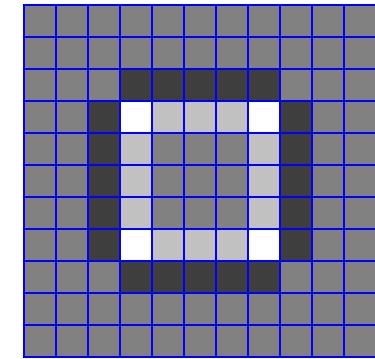
-1	2	-1



$$2I(r,c) - I(r-1,c) - I(r+1,c)$$

水平边缘检测

	-1	
	2	
	-1	

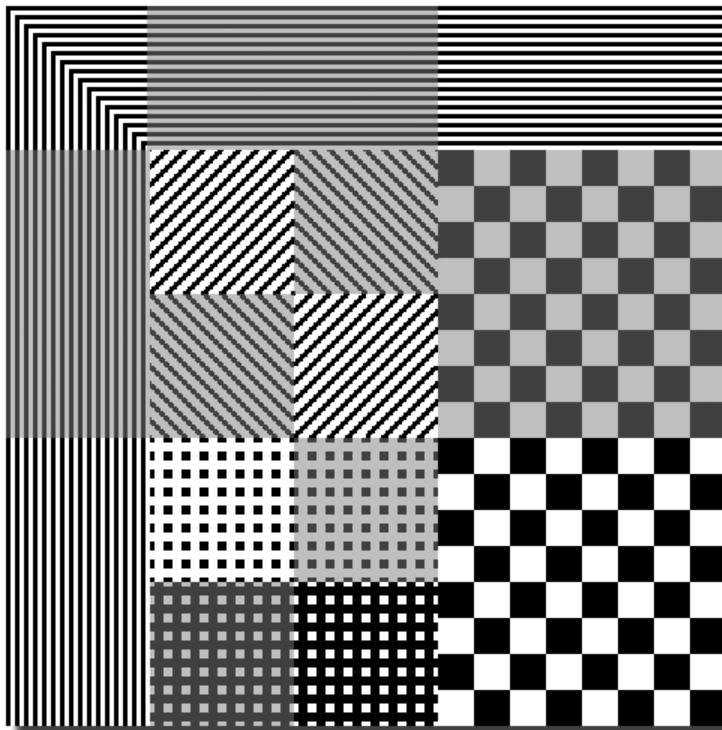


$$4I(r,c) - I(r-1,c) - I(r+1,c) - I(r,c-1) - I(r,c+1)$$

联合边缘检测

	-1	
-1	4	-1
	-1	

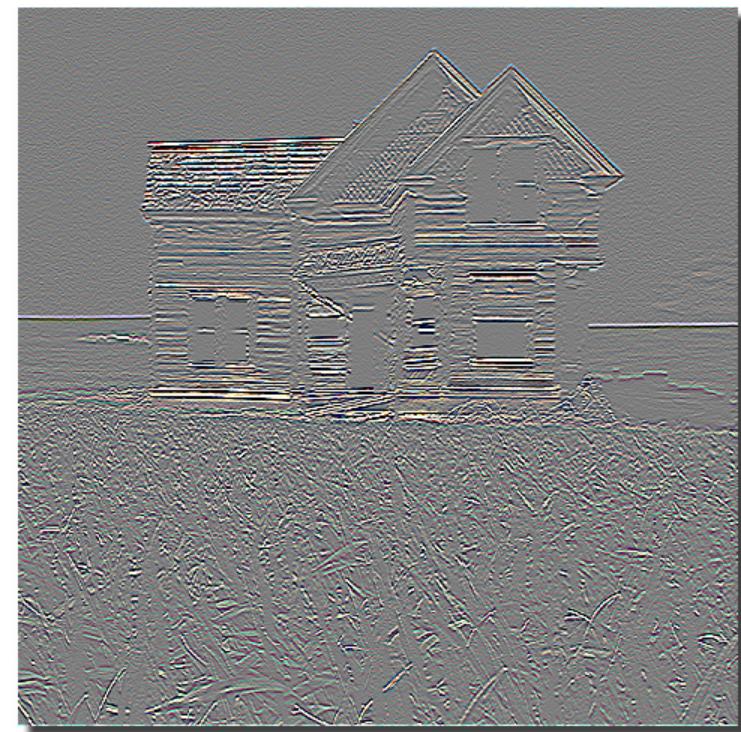
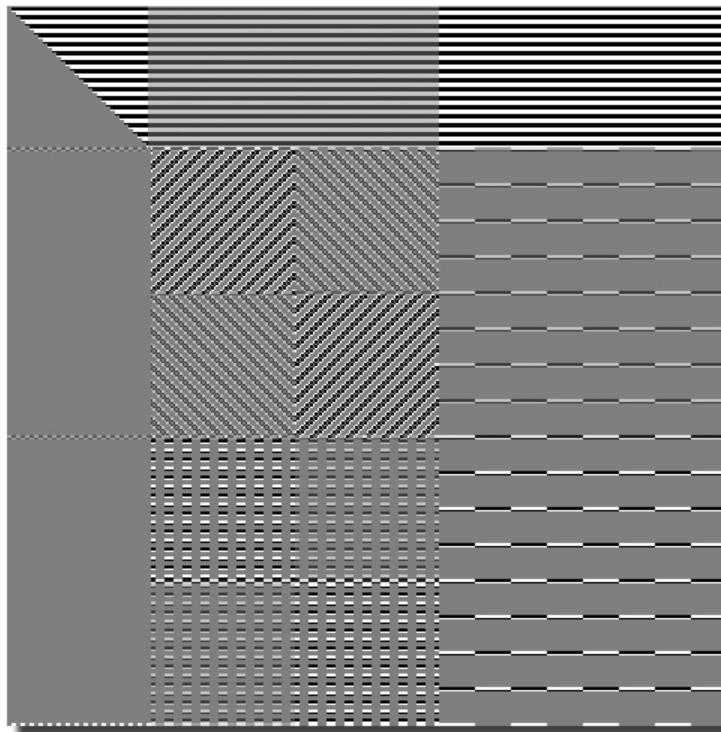
Convolution Examples: Original Images



$$\begin{bmatrix} -1 \\ 2 \\ -1 \end{bmatrix}$$

对于水平边缘检测（垂直变化），其滤波器核是一维的，可以将图像分解为多个列，按列作一维滤波

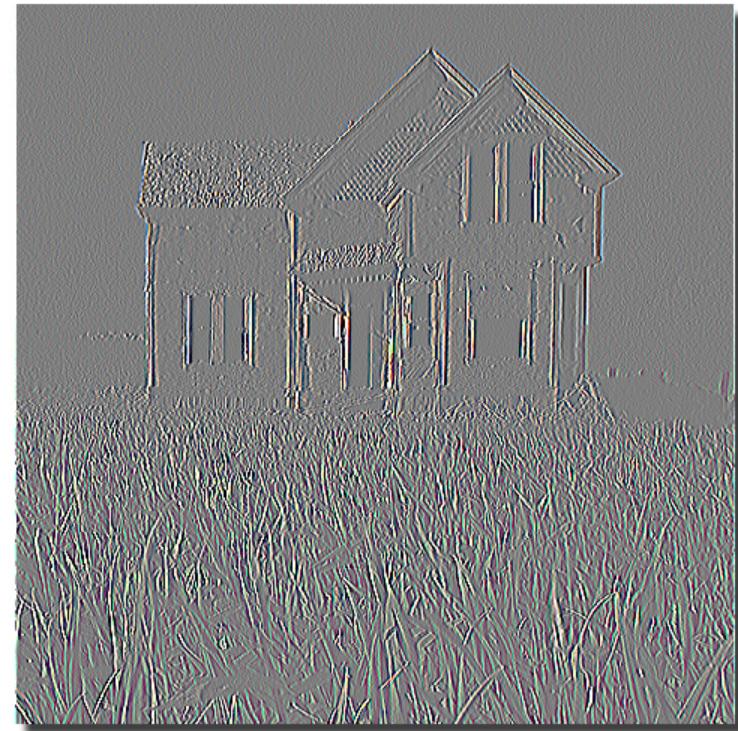
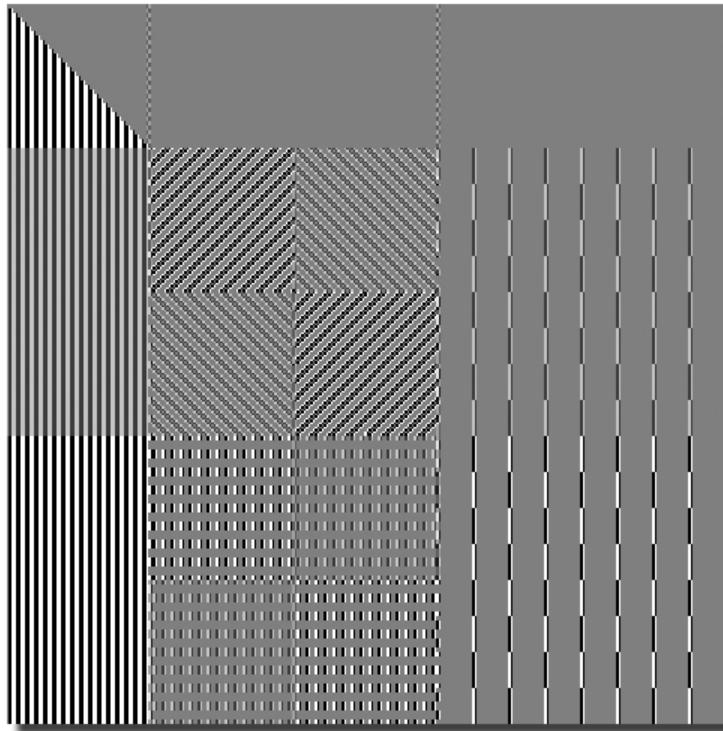
Convolution Examples: Vertical Difference



$[-1 \ 2 \ -1]$

按行作一维滤波

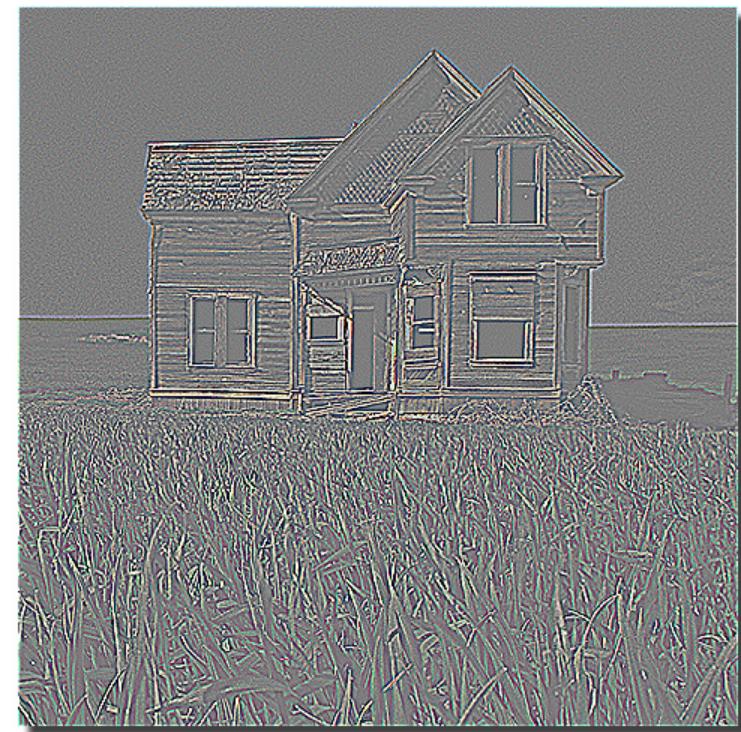
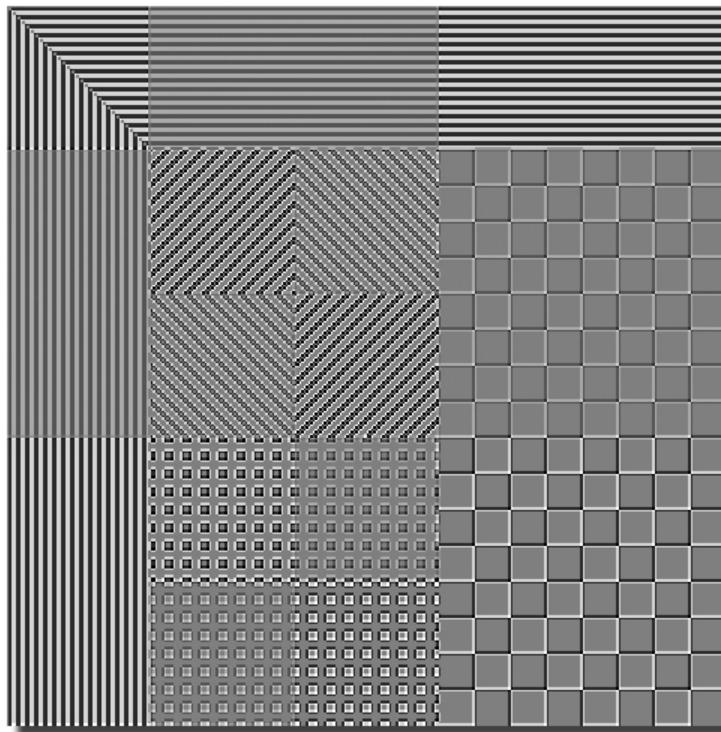
Convolution Examples: Horizontal Difference



$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

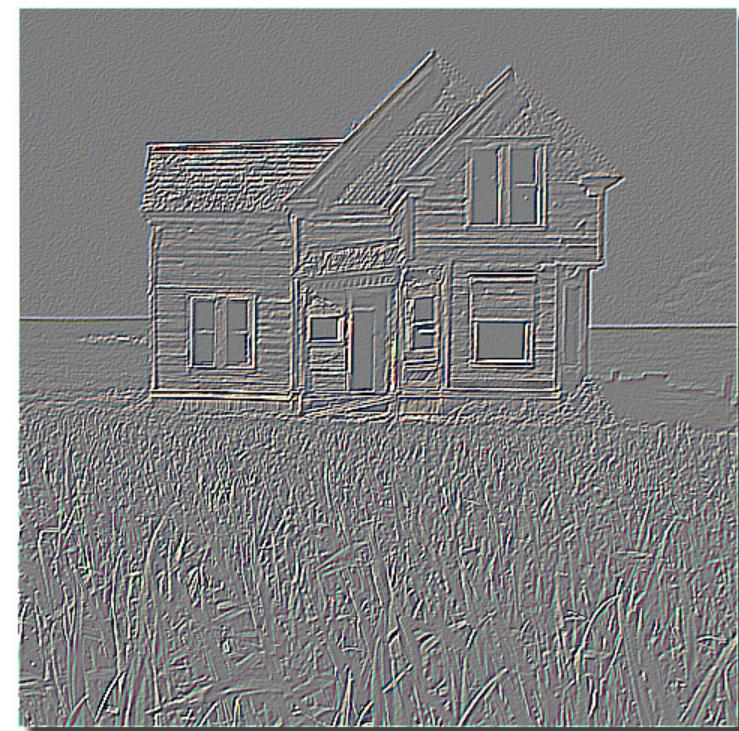
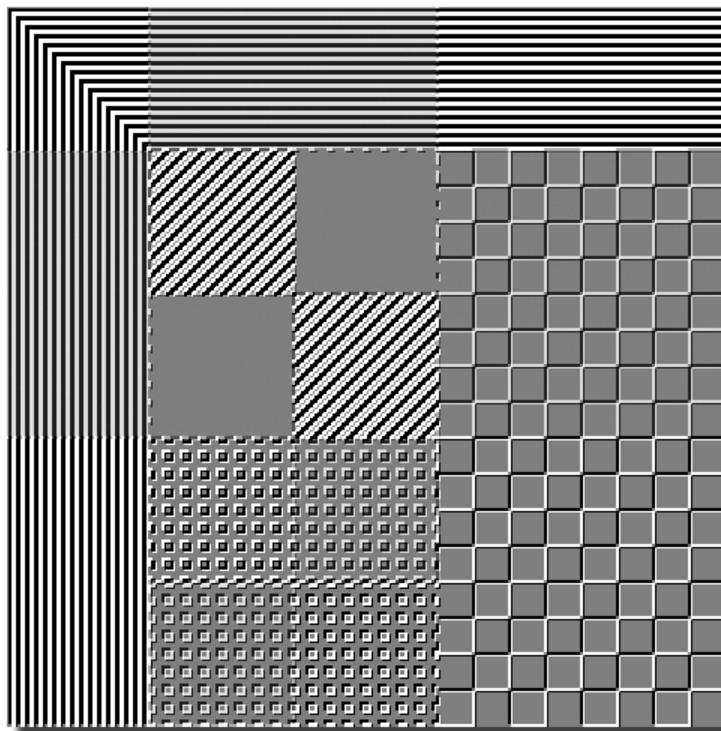
作二维滤波

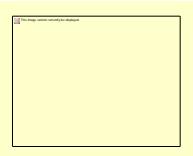
Convolution Examples: H + V Diff.



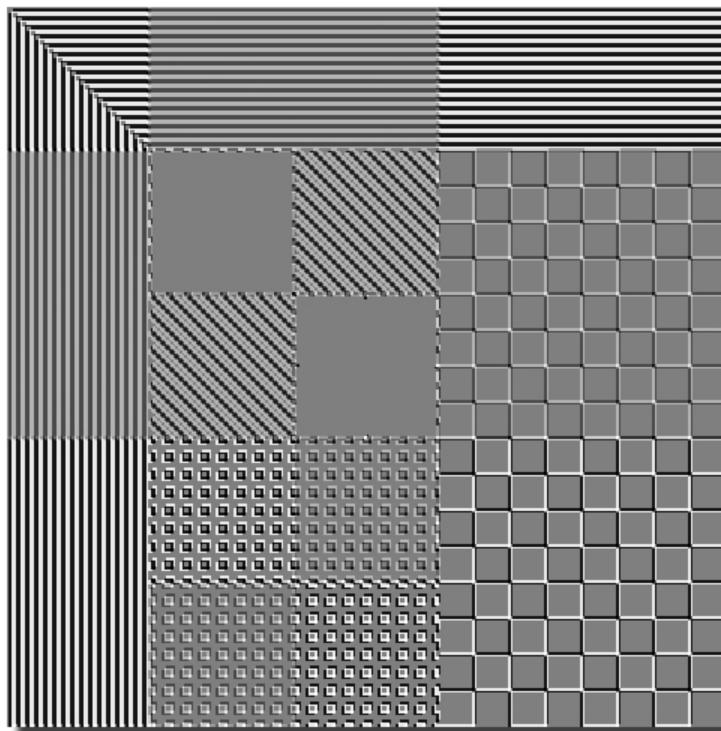
$$\begin{bmatrix} -1 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & -1 \end{bmatrix}$$

Convolution Examples: Diagonal Difference



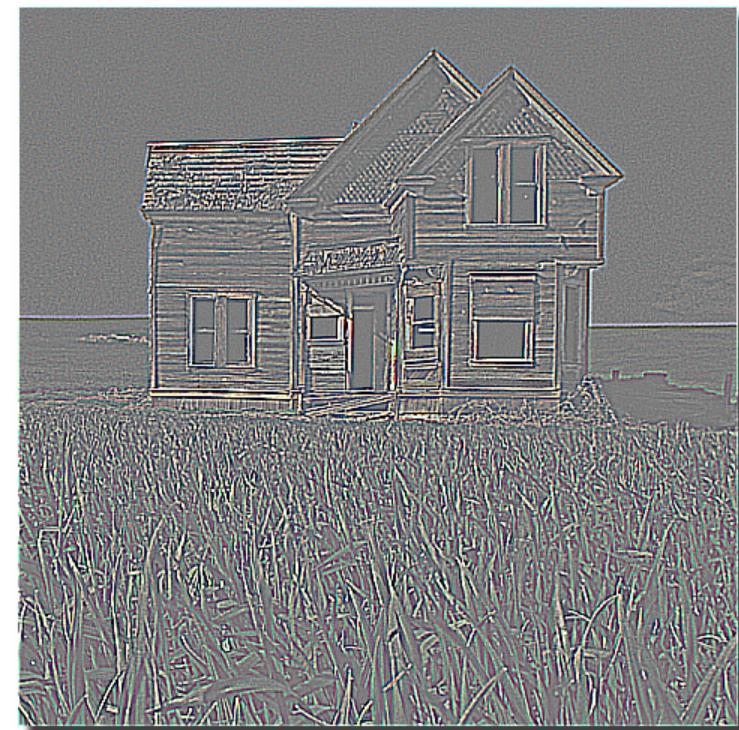
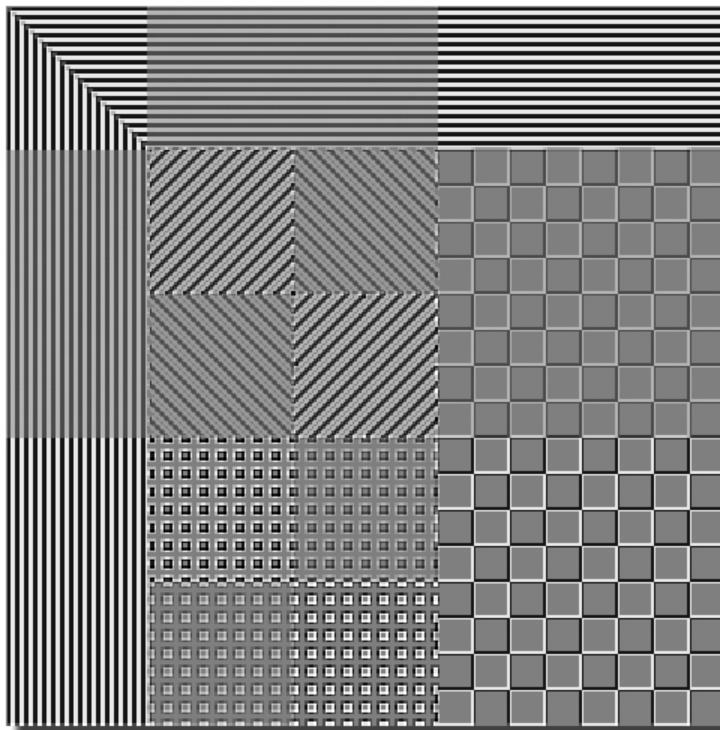


Convolution Examples: Diagonal Difference



$$\begin{bmatrix} -1 & 0 & -1 \\ 0 & 4 & 0 \\ -1 & 0 & -1 \end{bmatrix}$$

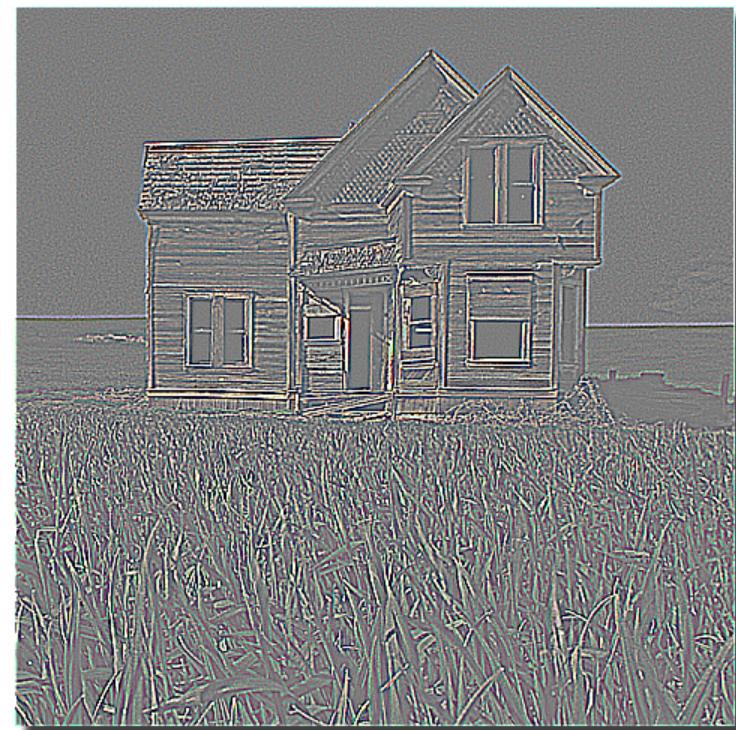
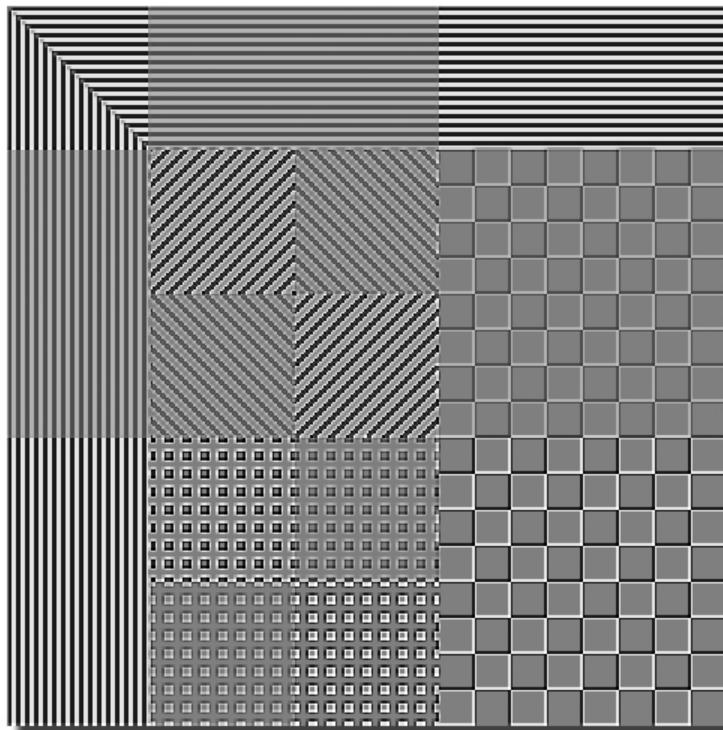
Convolution Examples: D + D Difference



$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

注意：高通滤波器各元素的和为0

Convolution Examples: H + V + D Diff.



Thanks!