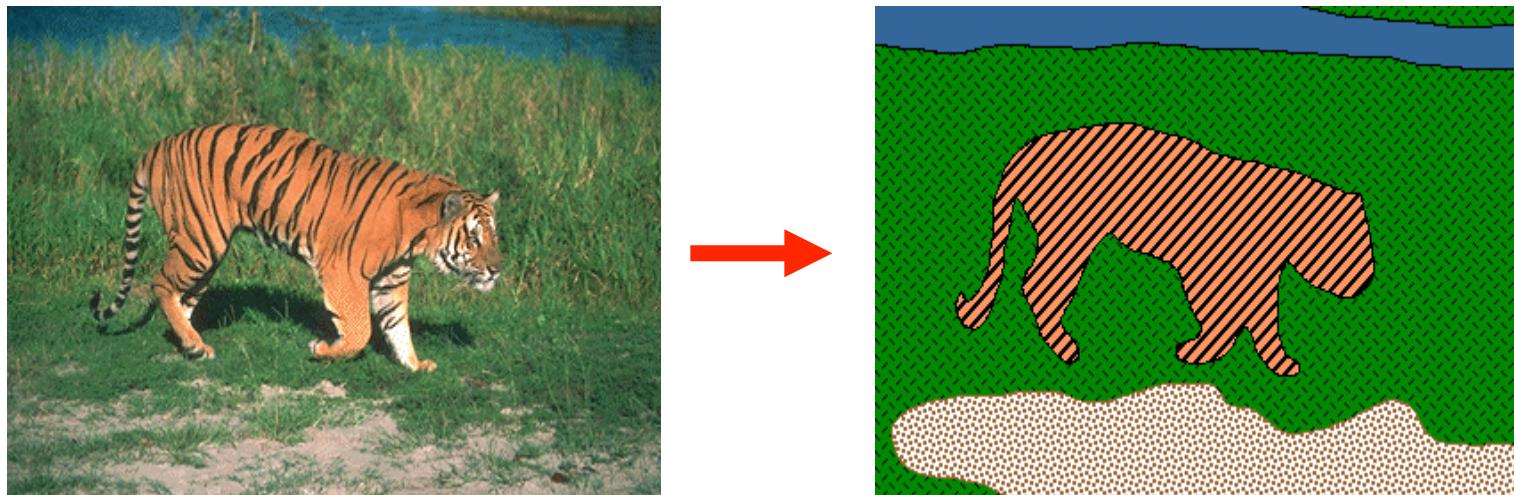


Image Processing

Segmentation – Part 2

Review- Image segmentation

- Goal: identify groups of pixels that go together

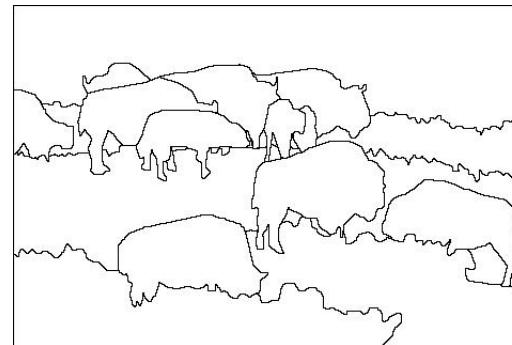


Review- The goals of segmentation

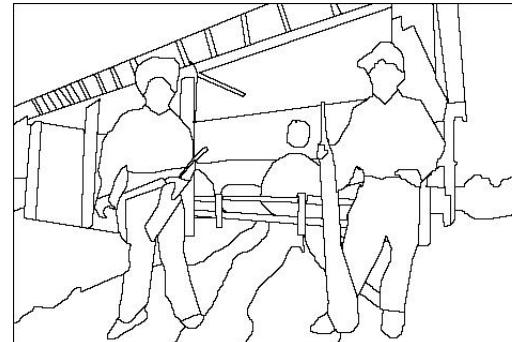
- Separate image into coherent “objects”



image



human segmentation

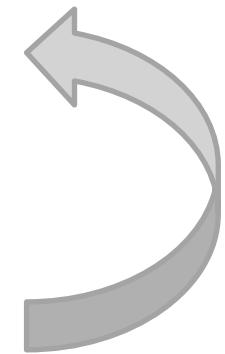


Review- What is segmentation?

- Clustering image elements that “belong together”
 - Partitioning
 - Divide into regions/sequences with coherent internal properties
 - Grouping
 - Identify sets of coherent tokens in image

Review- K-means clustering

- Basic idea: randomly initialize the k cluster centers, and iterate between the two steps we just saw.
 1. Randomly initialize the cluster centers, c_1, \dots, c_K
 2. Given cluster centers, determine points in each cluster
 - For each point p , find the closest c_i . Put p into cluster i
 3. Given points in each cluster, solve for c_i
 - Set c_i to be the mean of points in cluster i
 4. If c_i have changed, repeat Step 2



Properties

- Will always converge to some solution
- Can be a “local minimum”
 - does not always find the global minimum of objective function:

$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

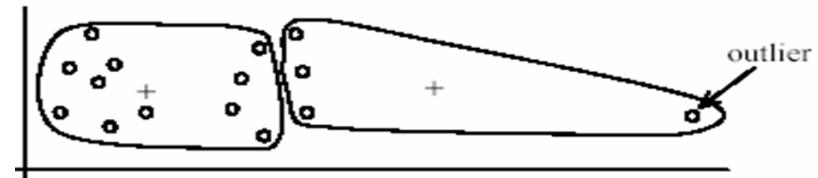
Review - K-means: pros and cons

Pros

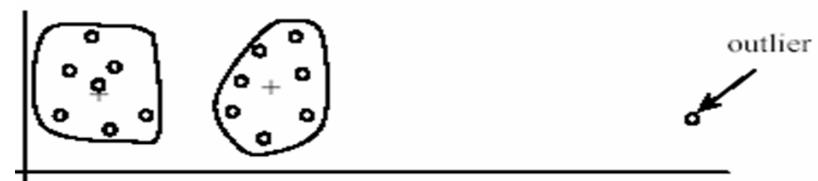
- Simple, fast to compute
- Converges to local minimum of within-cluster squared error

Cons/issues

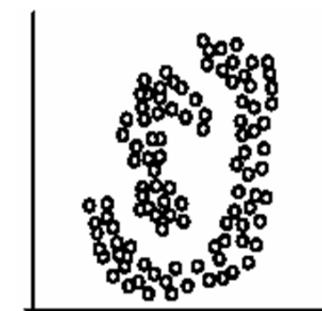
- Setting k?
- Sensitive to initial centers
- Sensitive to outliers
- Detects spherical clusters
- Assuming means can be computed



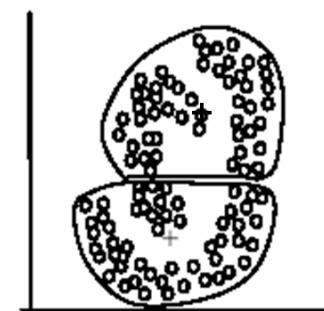
(A): Undesirable clusters



(B): Ideal clusters



(A): Two natural clusters



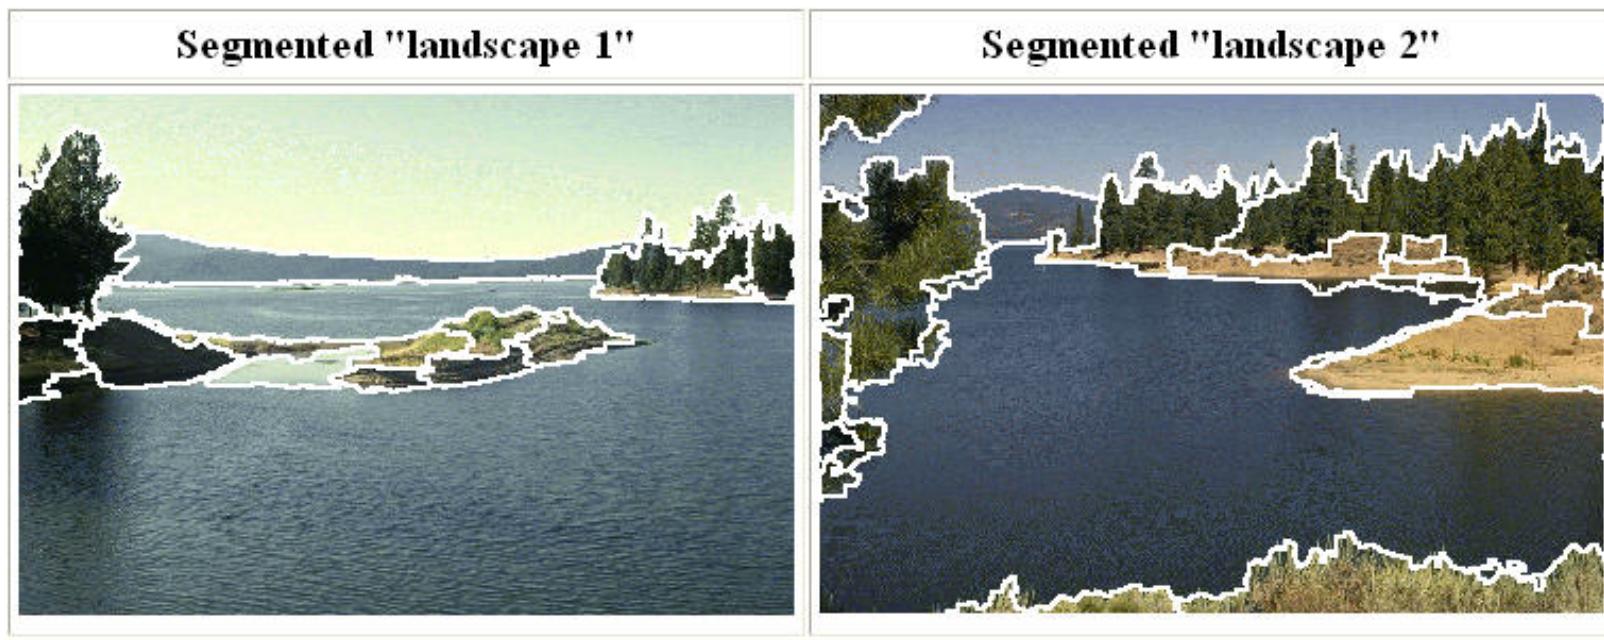
(B): k -means clusters

Segmentation methods

- Segment foreground from background
- Histogram-based segmentation
- Segmentation as clustering
 - K-means clustering
 - Mean-shift segmentation
- Graph-theoretic segmentation
 - Min cut
 - Normalized cuts
- Interactive segmentation

Mean shift clustering and segmentation

- An advanced and versatile technique for clustering-based segmentation

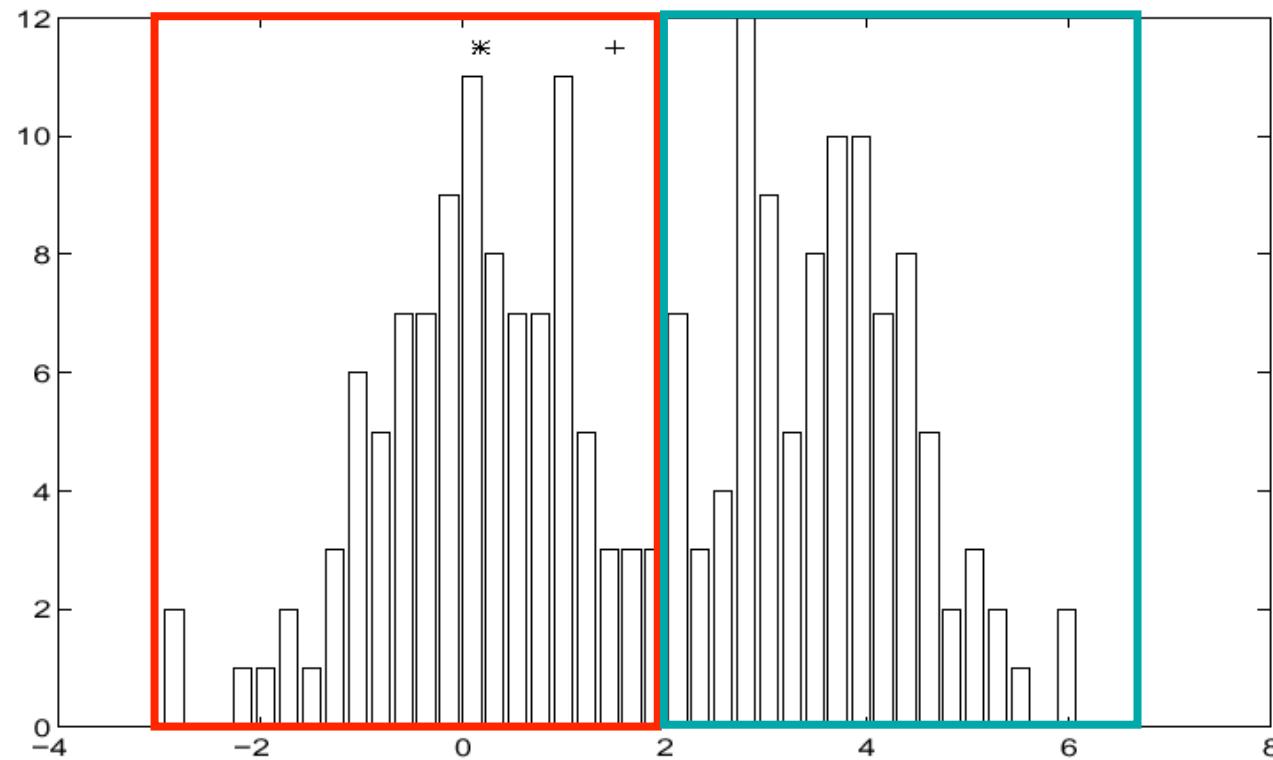


<http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html>

D. Comaniciu and P. Meer, [Mean Shift: A Robust Approach toward Feature Space Analysis](#),
PAMI 2002.

Slide credit: S. Lazebnik

Finding Modes in a Histogram



- How Many Modes Are There?
 - Easy to see, hard to compute

Mean shift algorithm

- The mean shift algorithm seeks *modes* or local maxima of density in the feature space

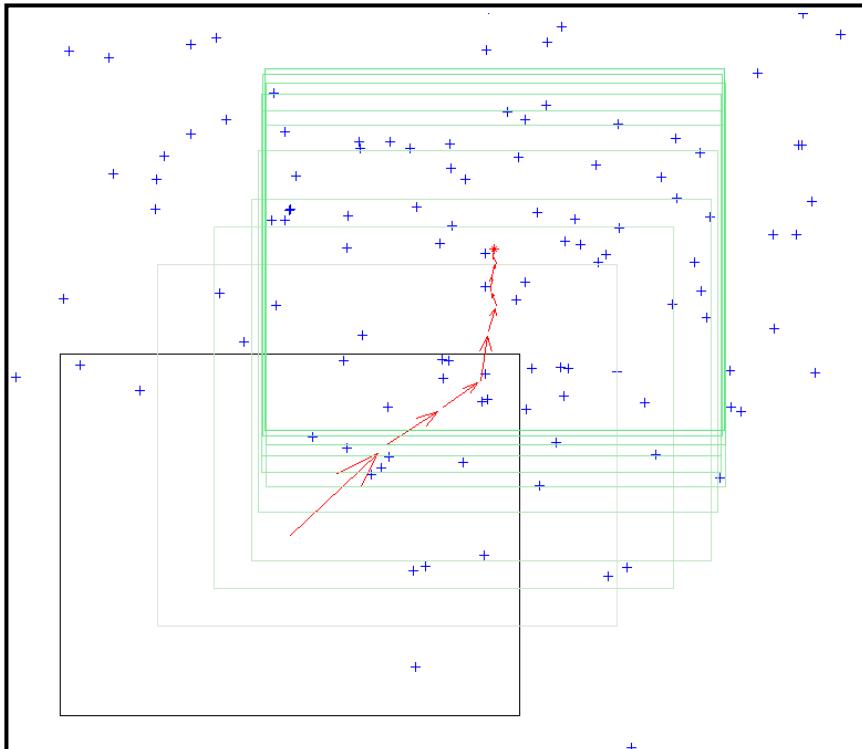
For n-D feature space, find local maxima or modes is extremely hard to compute

Mean shift algorithm

Mean Shift Algorithm

1. Choose a search window size.
2. Choose the initial location of the search window.
3. Compute the mean location (centroid of the data) in the search window.
4. Center the search window at the mean location computed in Step 3.
5. Repeat Steps 3 and 4 until convergence.

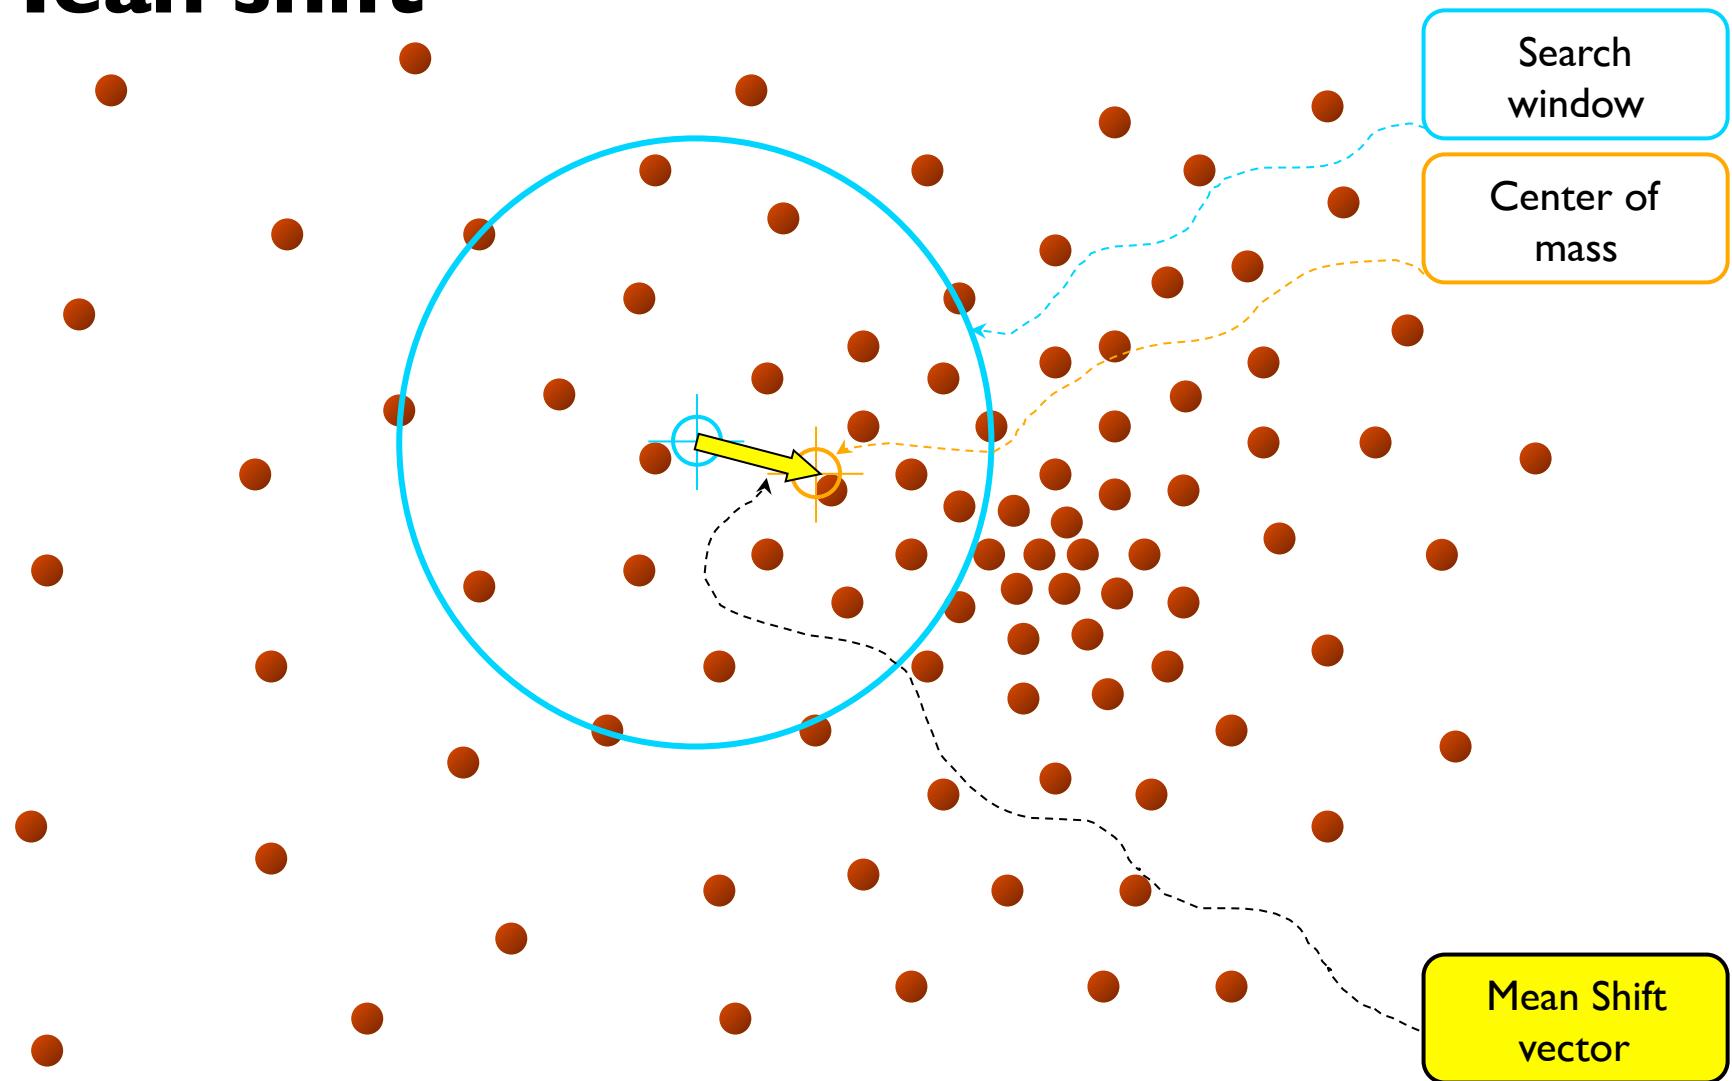
The mean shift algorithm seeks the “mode” or point of highest density of a data distribution:



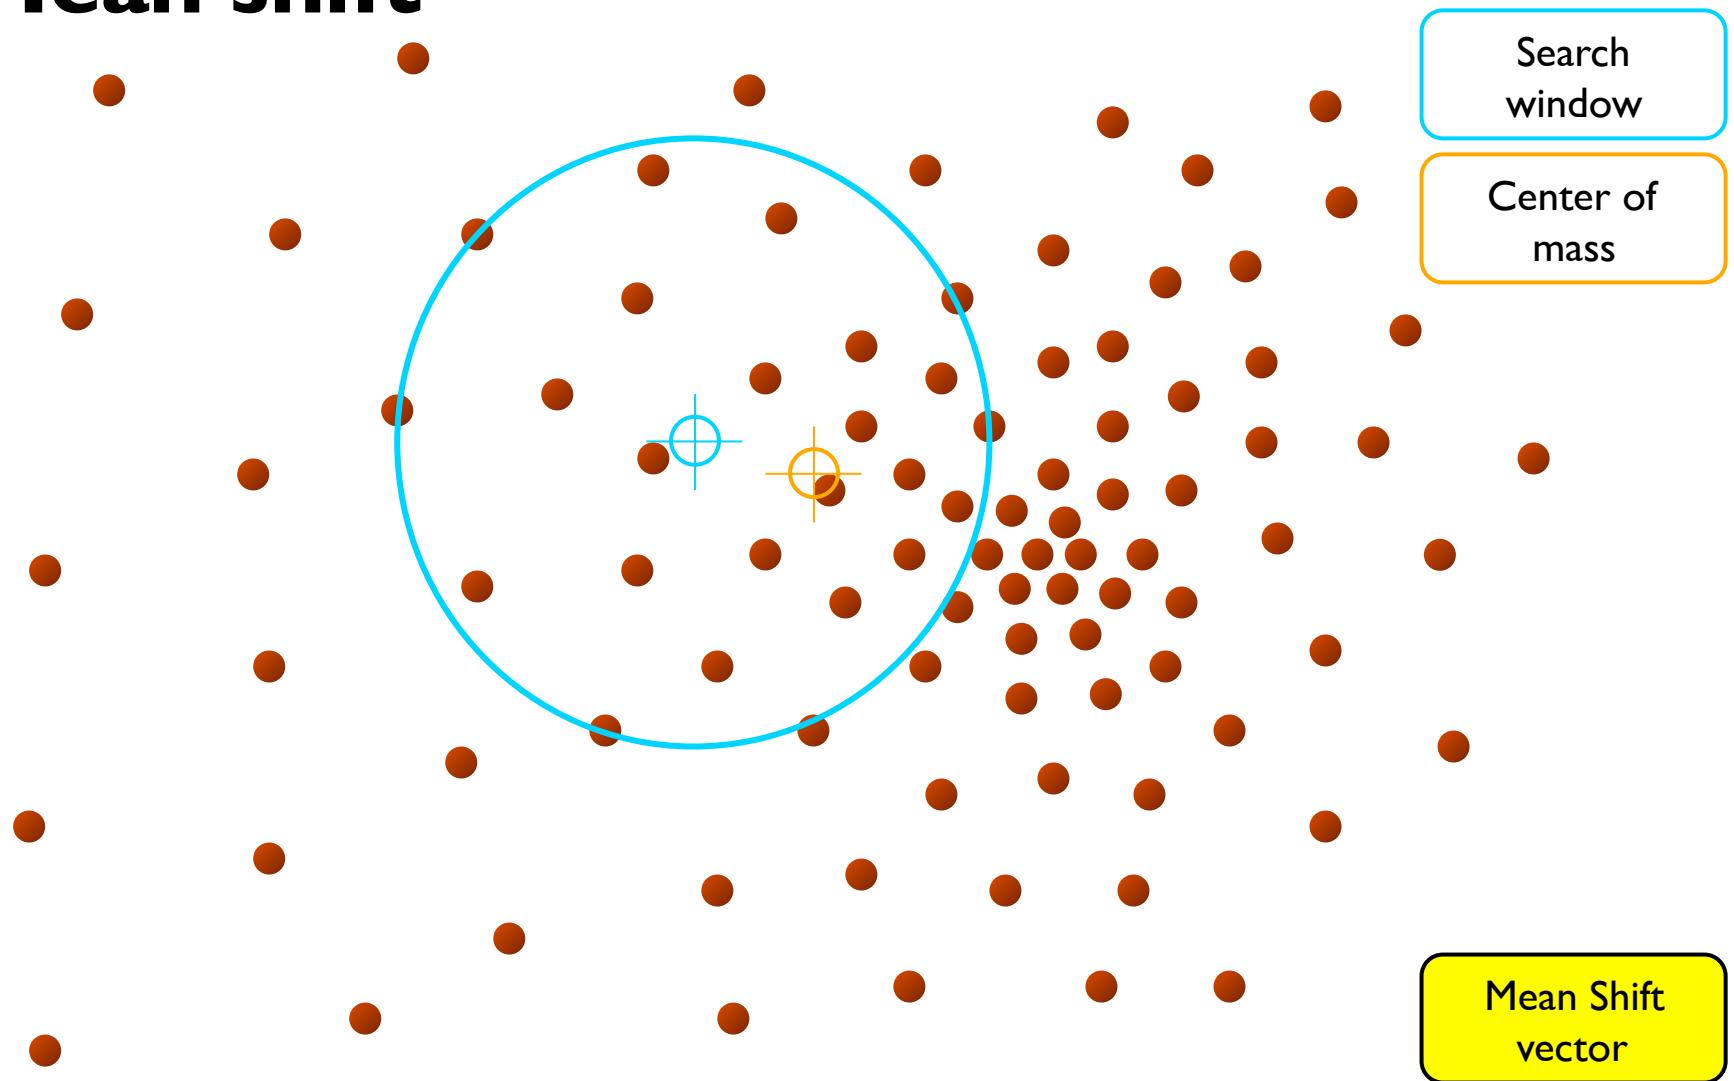
Two issues:

- (1) Kernel to interpolate density based on sample positions.
- (2) Gradient ascent to mode.

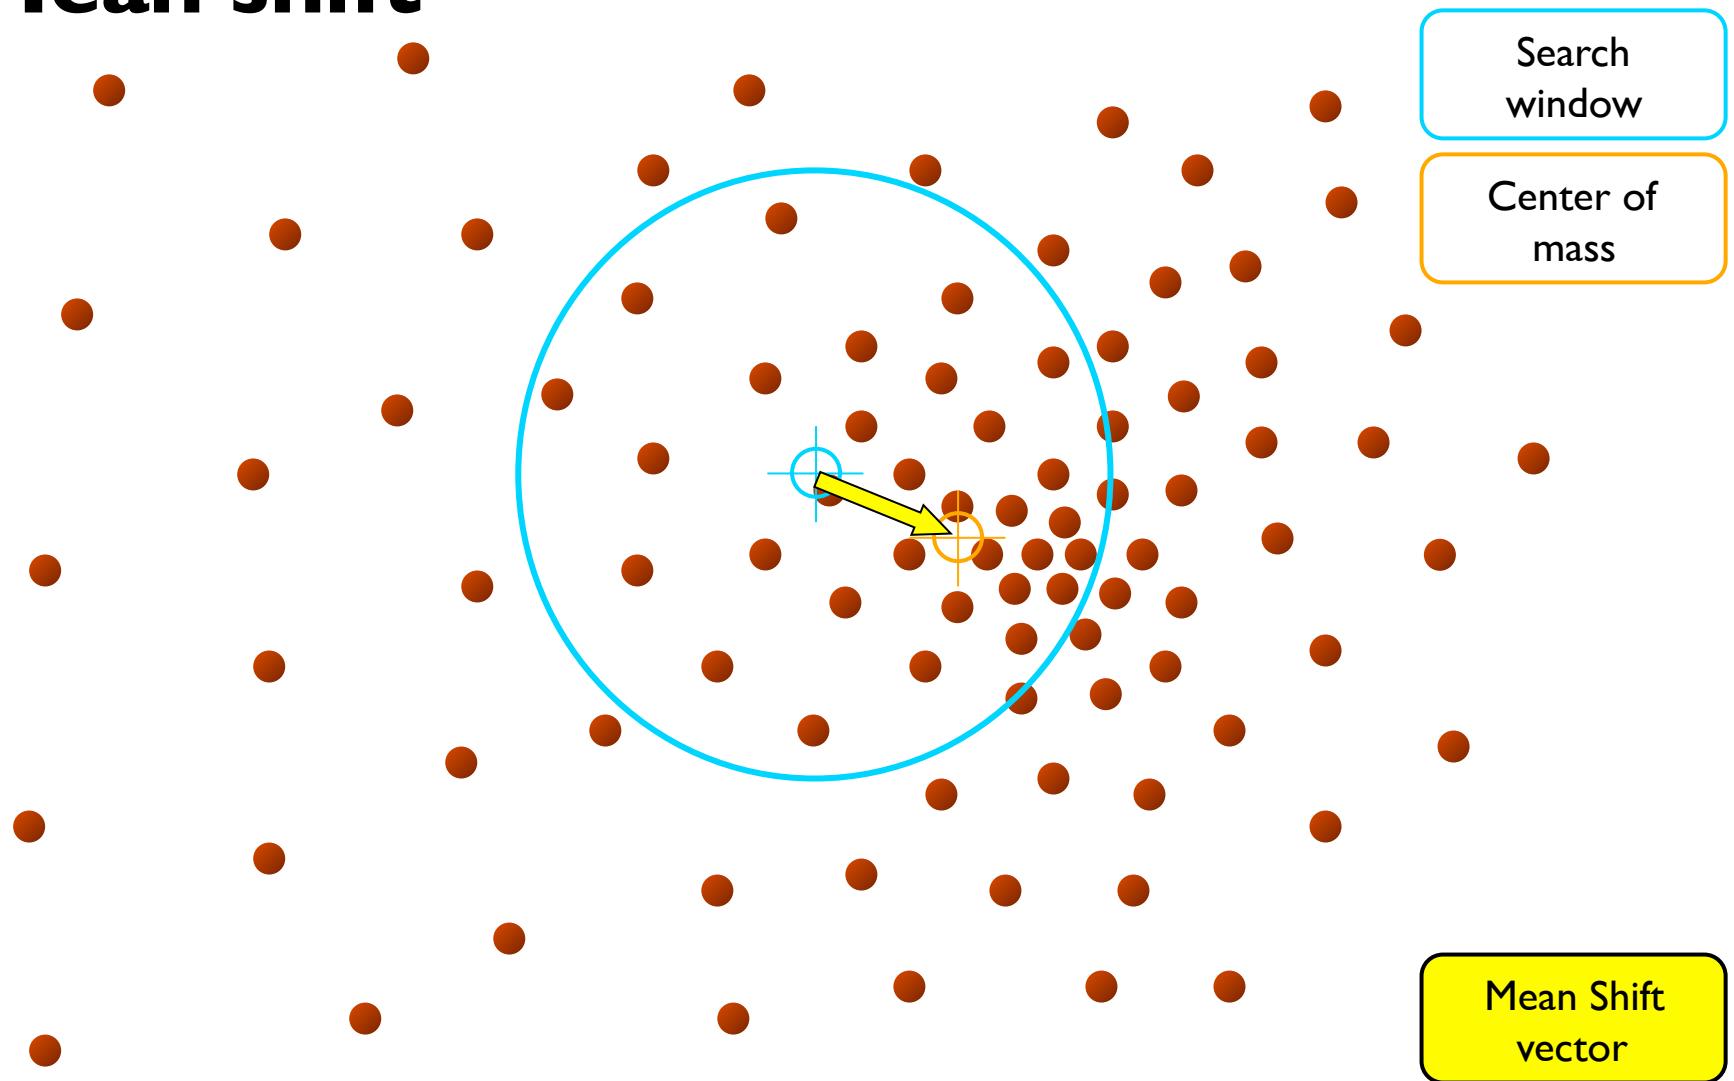
Mean shift



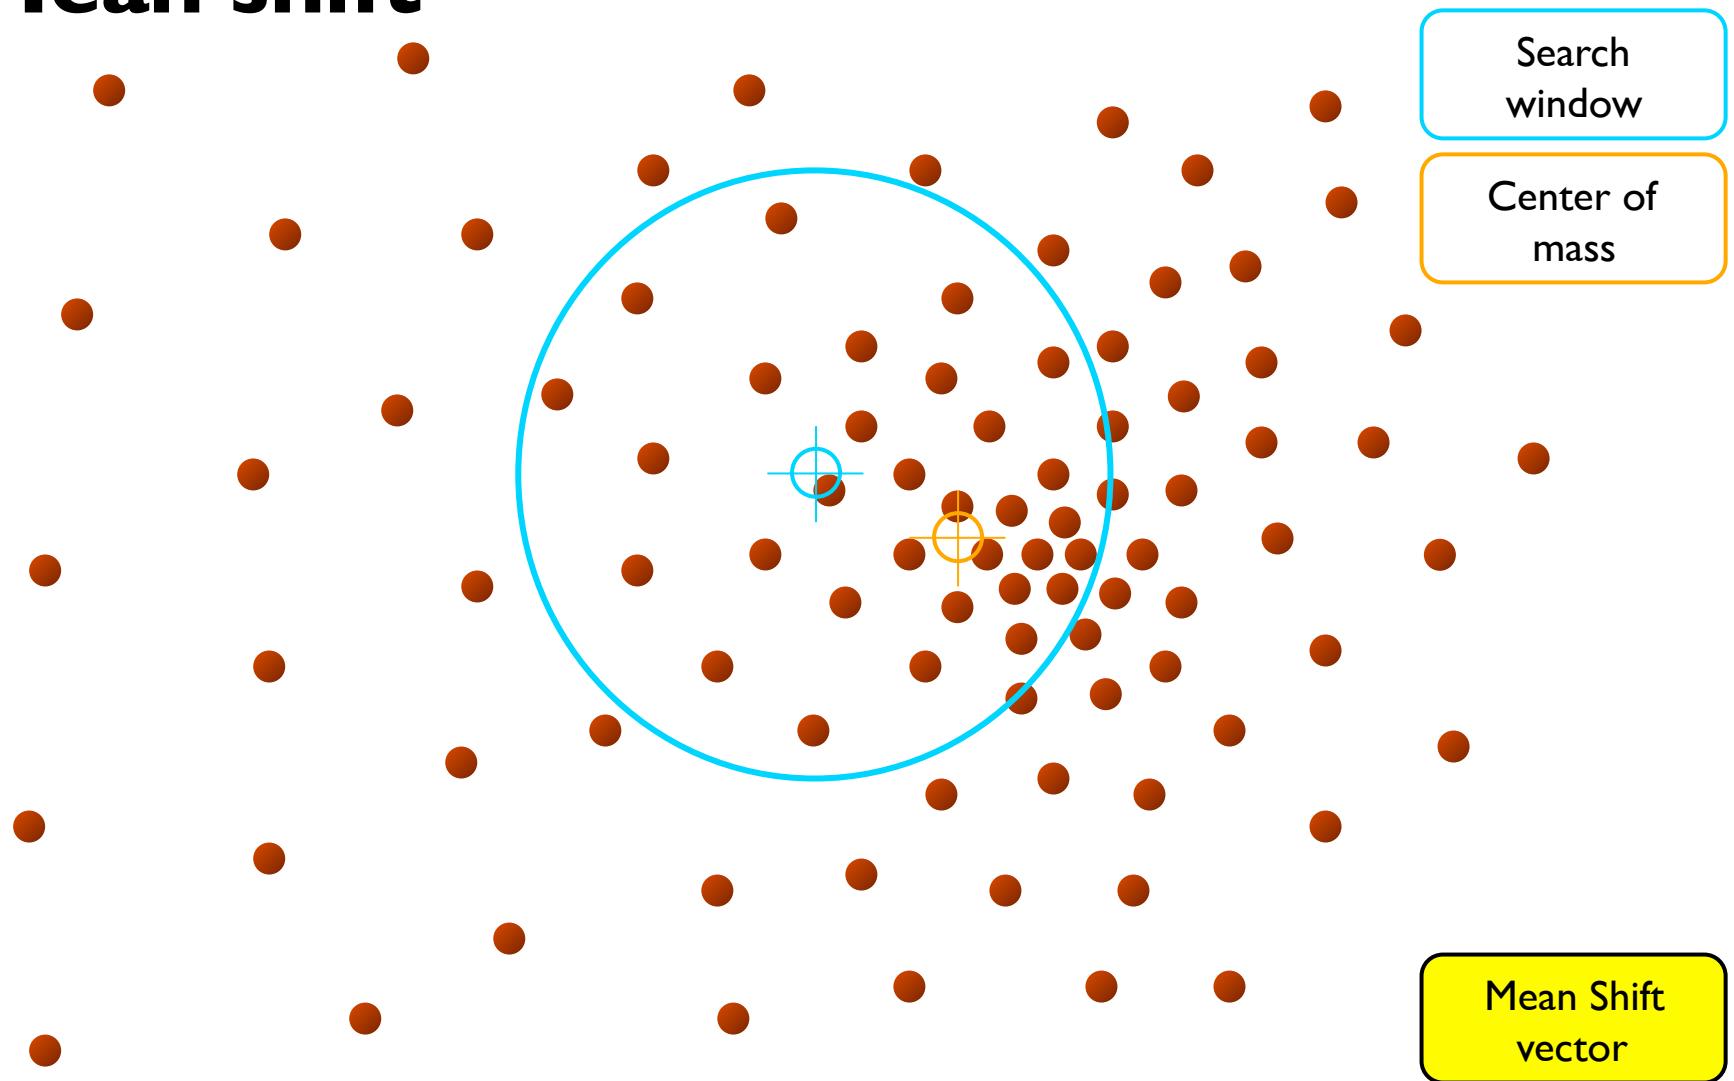
Mean shift



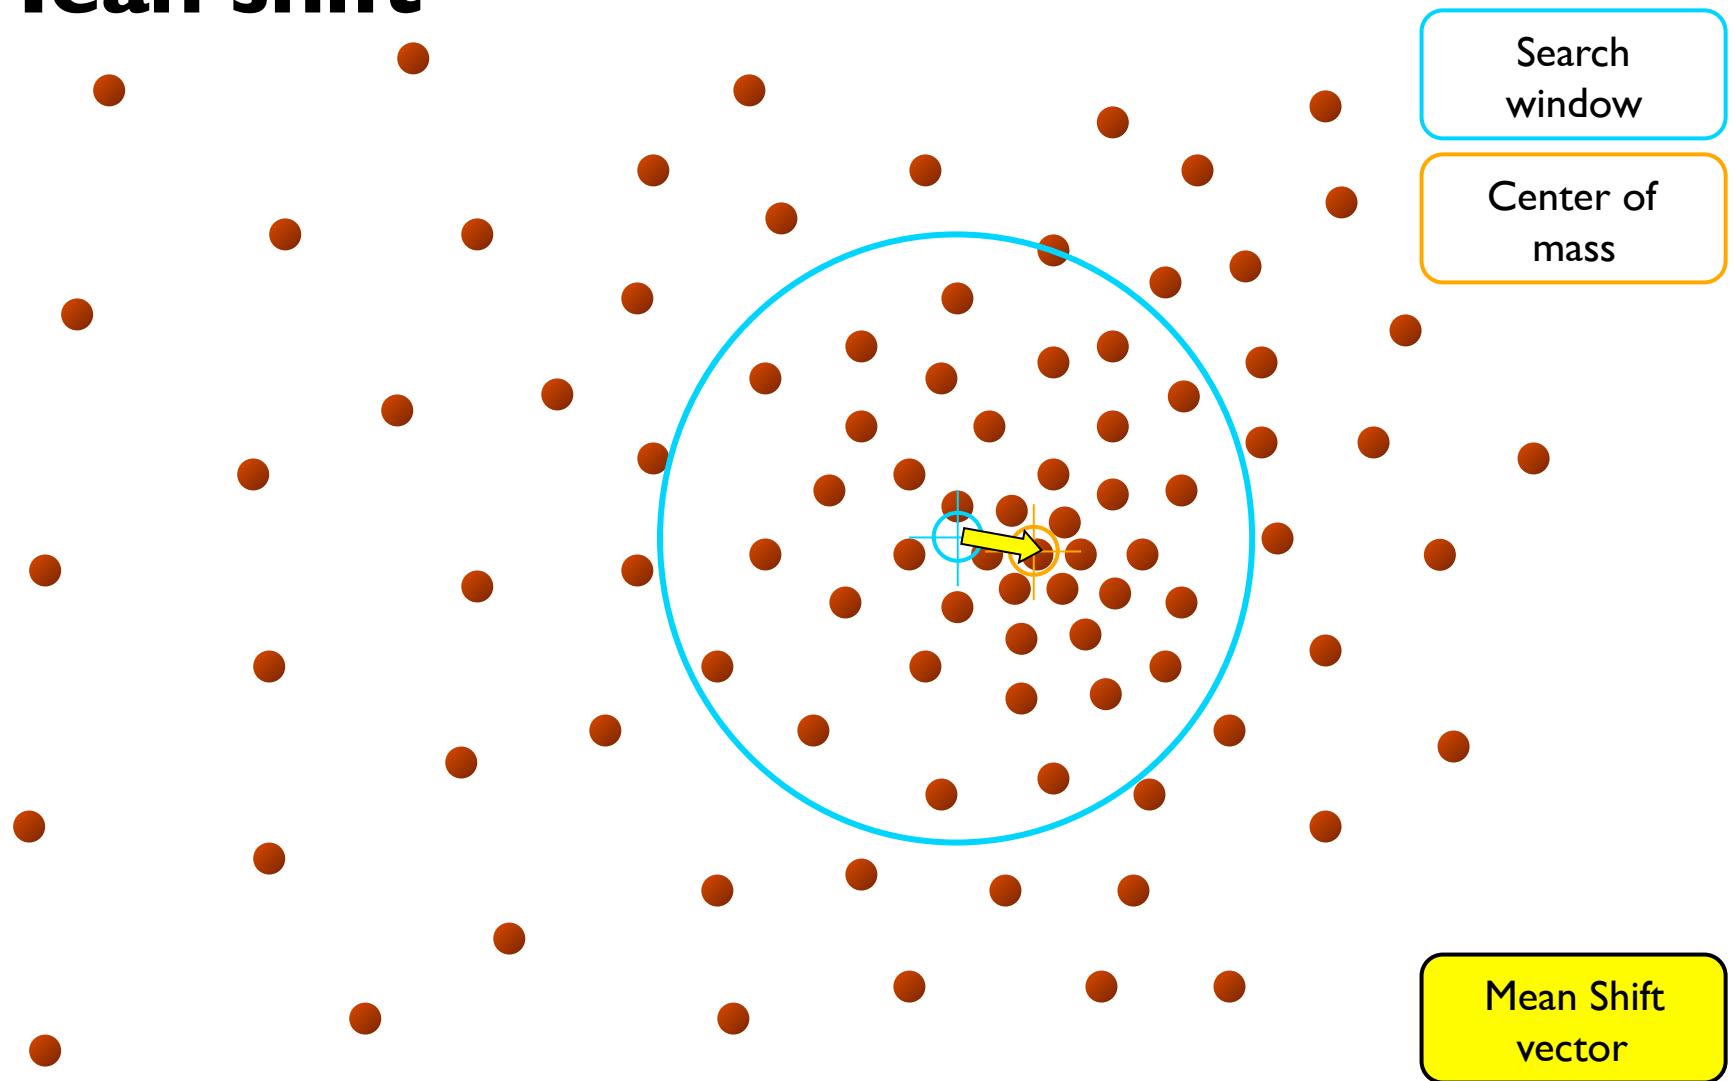
Mean shift



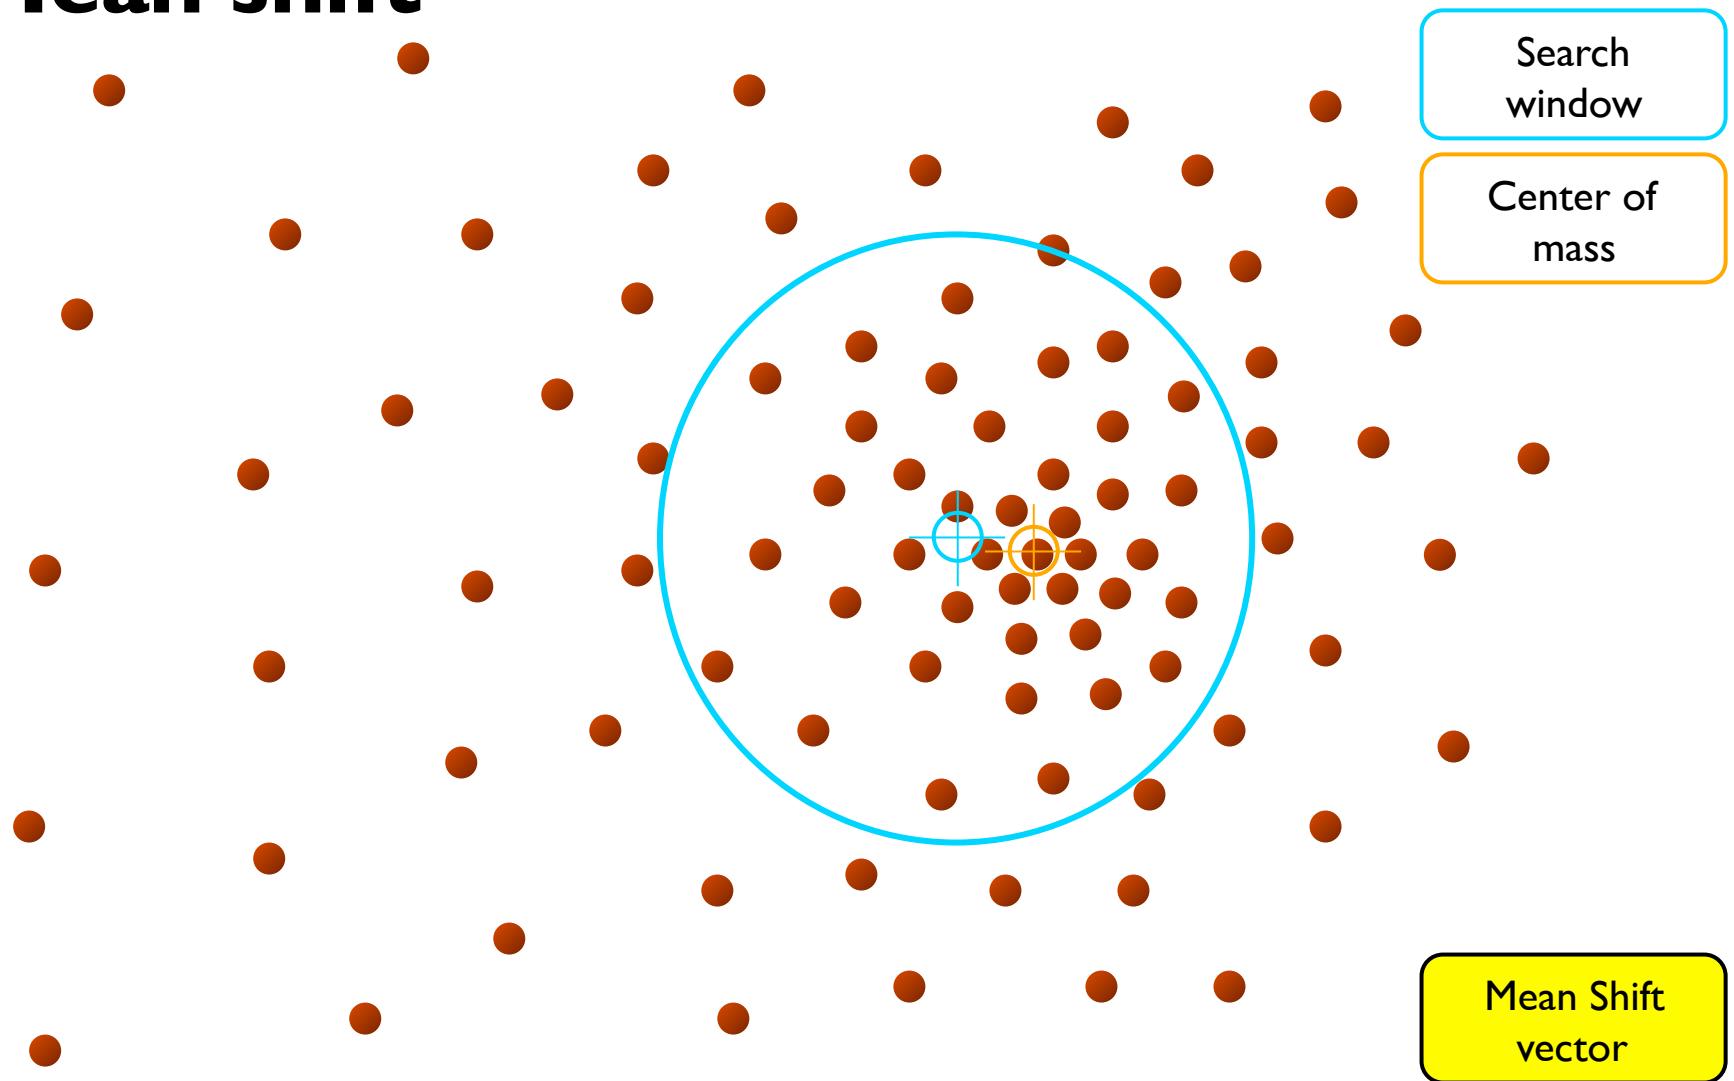
Mean shift



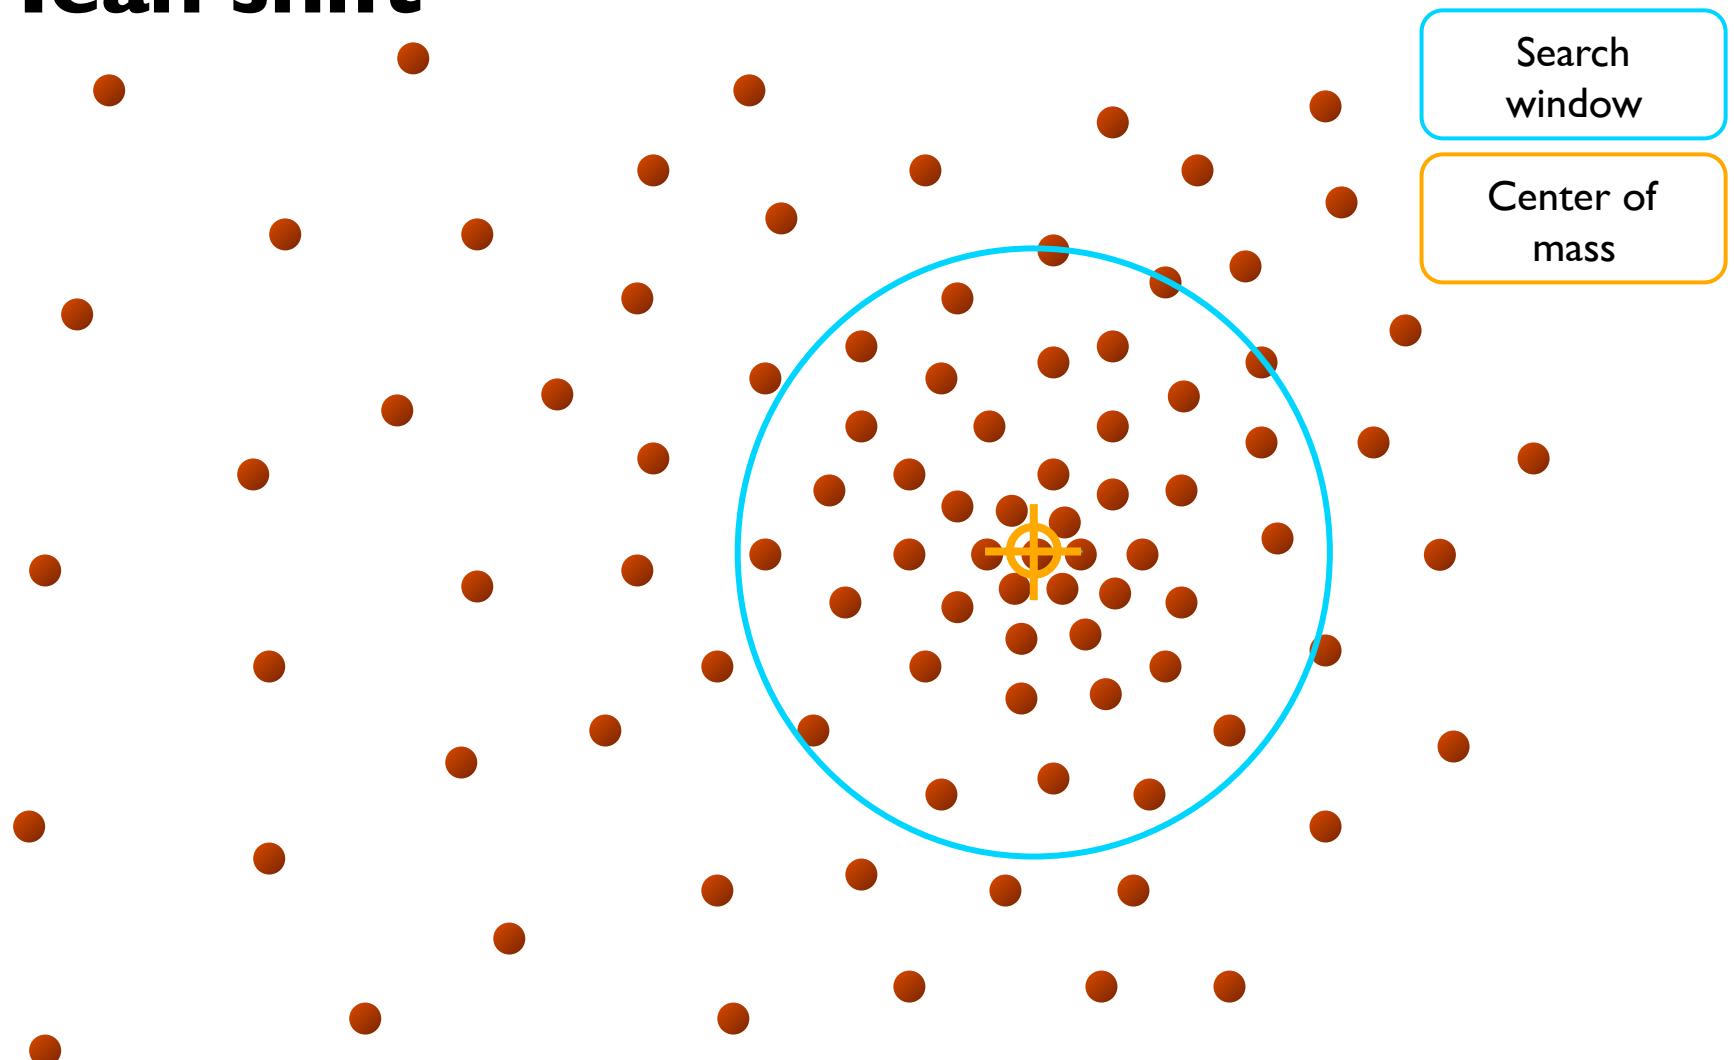
Mean shift



Mean shift

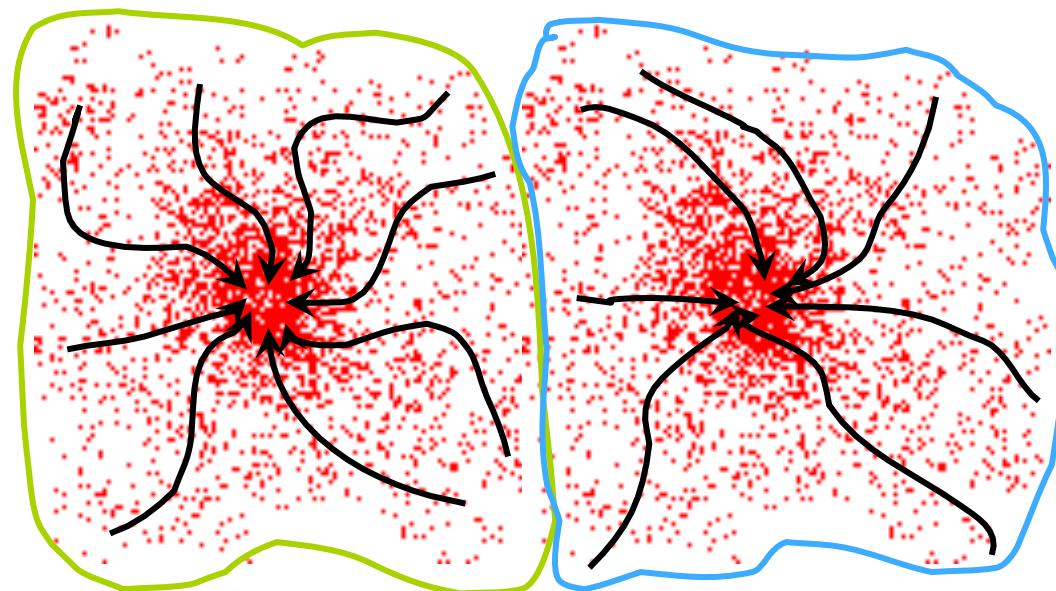


Mean shift



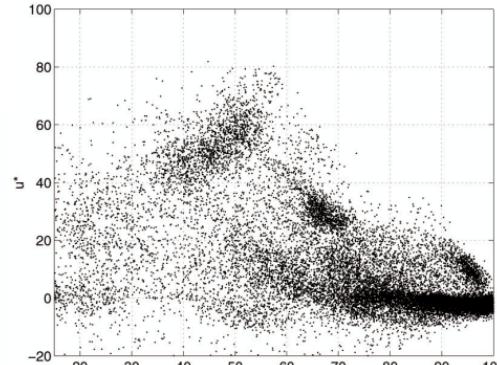
Mean shift clustering

- Cluster: all data points in the attraction basin of a mode
- Attraction basin: the region for which all trajectories lead to the same mode



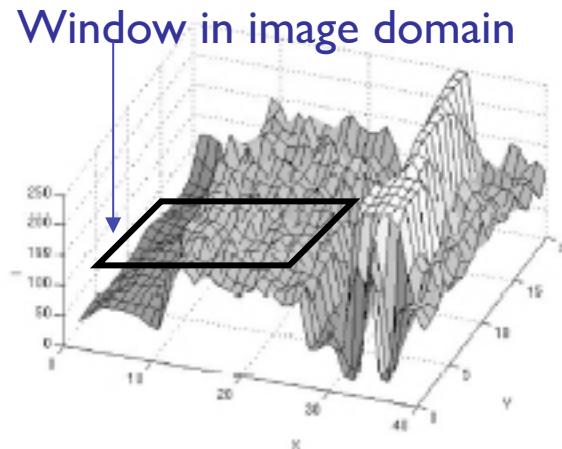
Mean shift clustering/segmentation

- Find features (color, gradients, texture, etc)
- Initialize windows at individual feature points
- Perform mean shift for each window until convergence
- Merge windows that end up near the same “peak” or mode



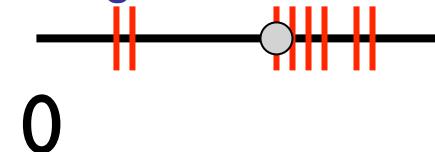
(a)

1



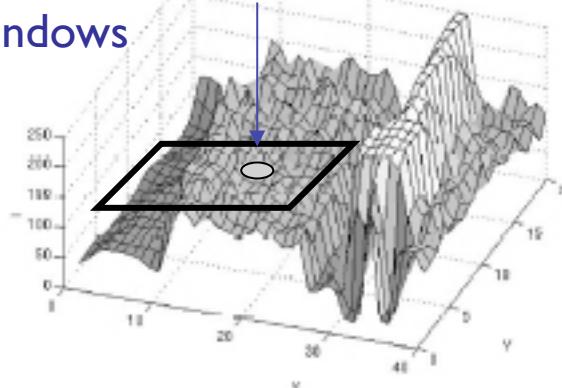
Apply mean shift jointly in the image (left col.) and range (right col.) domains

Intensities of pixels within image domain window

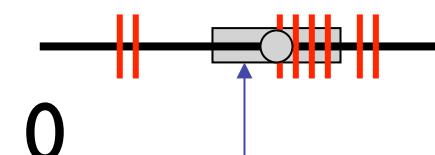


2

Center of mass of pixels within both image and range domain windows

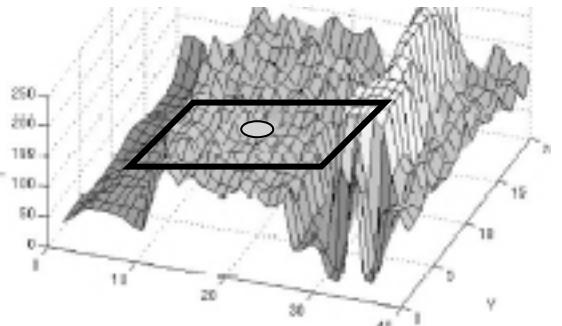


3

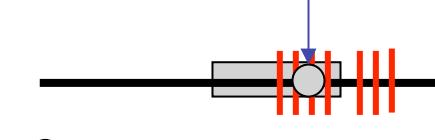


Window in range domain

4

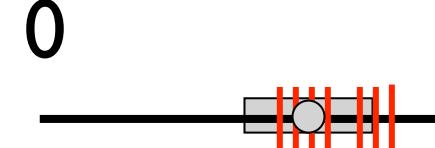


Center of mass of pixels within both image and range domain windows



6

5



7

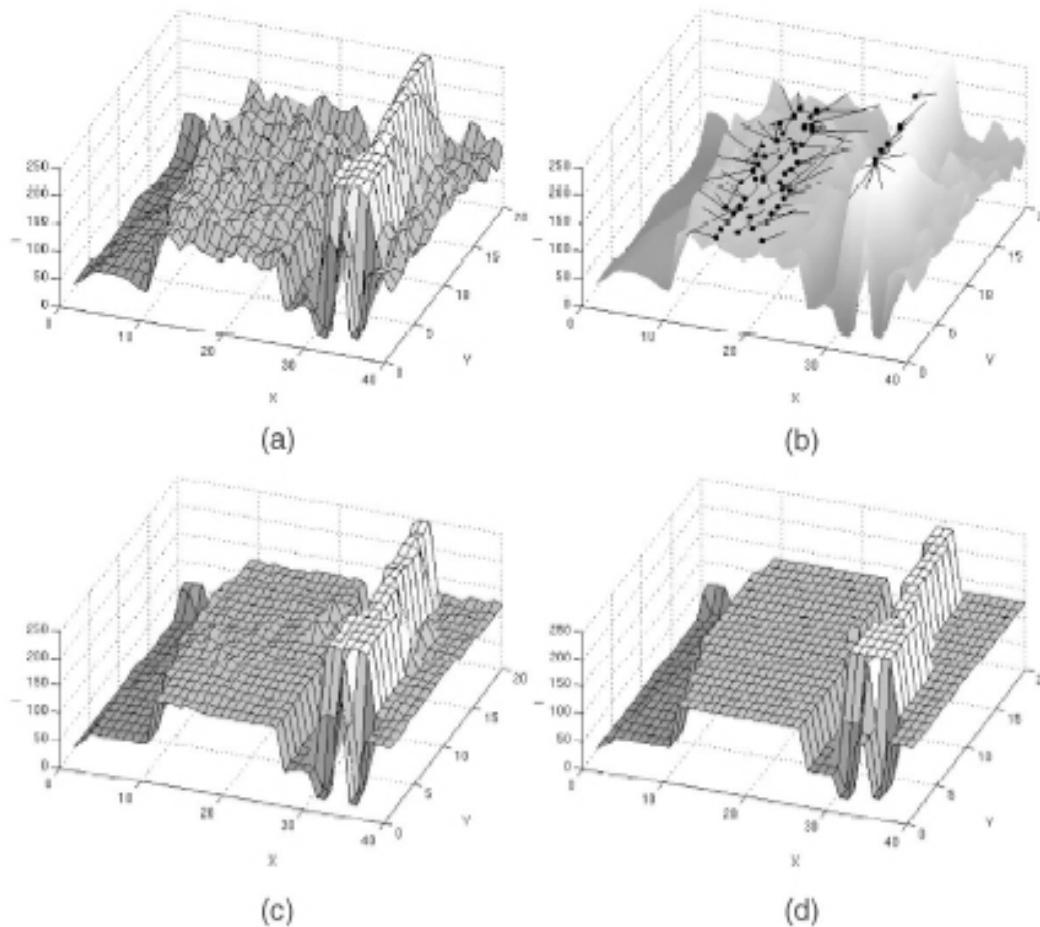


Fig. 4. Visualization of mean shift-based filtering and segmentation for gray-level data. (a) Input. (b) Mean shift paths for the pixels on the plateau and on the line. The black dots are the points of convergence. (c) Filtering result $(h_s, h_r) = (8, 4)$. (d) Segmentation result.

Mean shift segmentation results



<http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html>

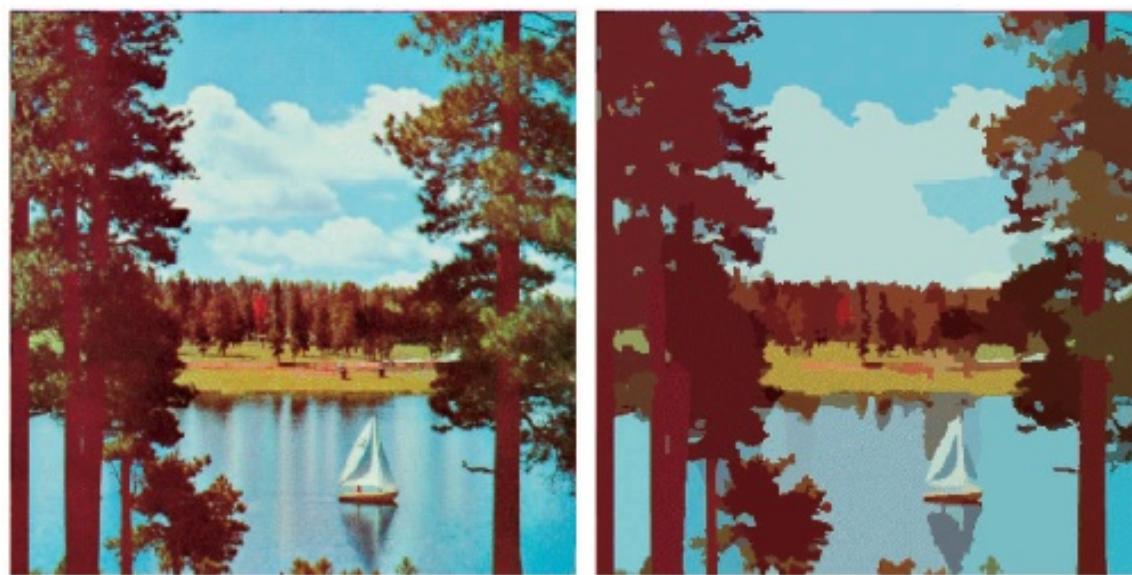
Slide credit: S. Lazebnik

More results



Slide credit: S. Lazebnik

More results



Slide credit: S. Lazebnik

Mean shift pros and cons

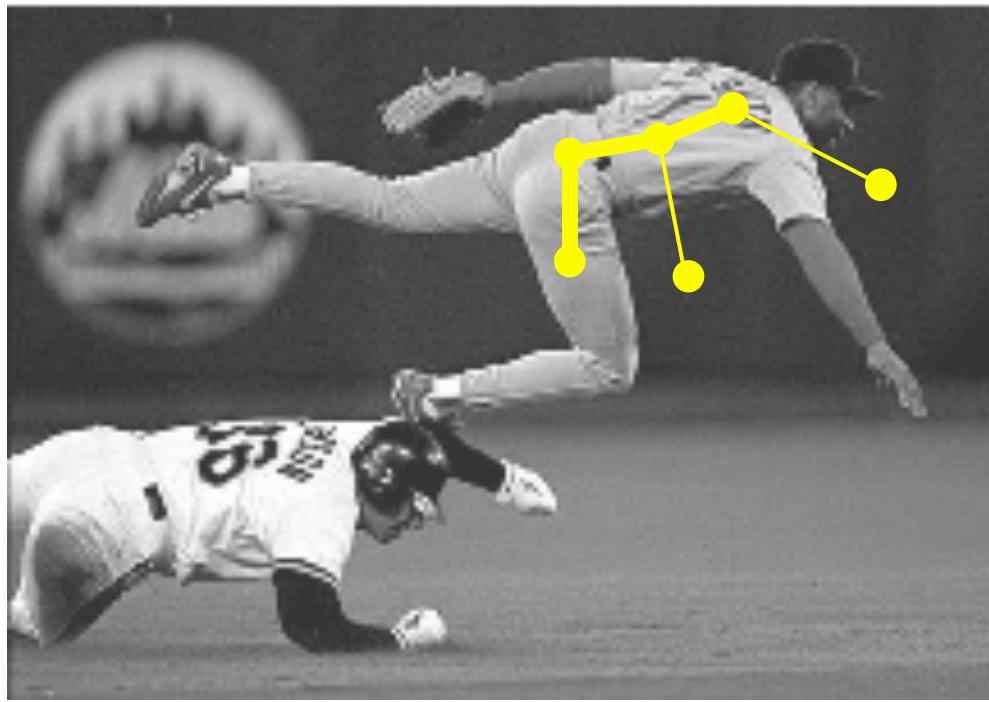
- Pros
 - Does not assume spherical clusters
 - Just a single parameter (window size)
 - Finds variable number of modes
 - Robust to outliers
- Cons
 - Output depends on window size
 - Computationally expensive
 - Does not scale well with dimension of feature space

Segmentation methods

- Segment foreground from background
- Histogram-based segmentation
- Segmentation as clustering
 - K-means clustering
 - Mean-shift segmentation
- **Graph-theoretic segmentation**
 - Min cut
 - Normalized cuts
- Interactive Segmentation

Graph-Theoretic Image Segmentation

Build a weighted graph $G=(V,E)$ from image



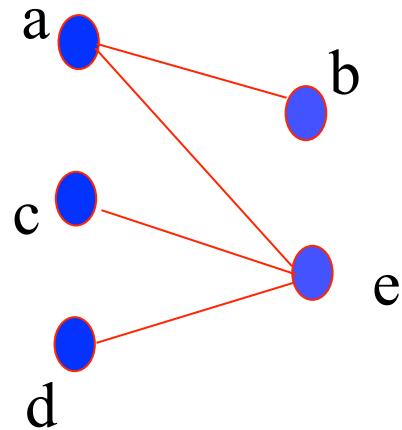
V : image pixels

E : connections between pairs of nearby pixels

W_{ij} : probability that i & j belong to the same region

Segmentation = graph partition

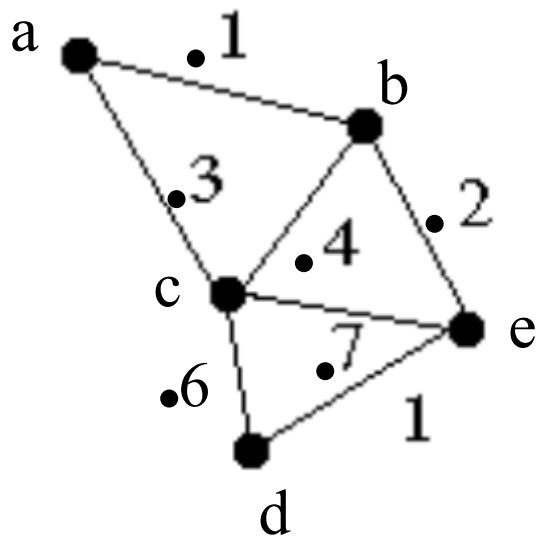
Graphs Representations



	a	b	c	d	e
a	0	1	0	0	1
b	1	0	0	0	0
c	0	0	0	0	1
d	0	0	0	0	1
e	1	0	1	1	0

Adjacency Matrix

A Weighted Graph and its Representation

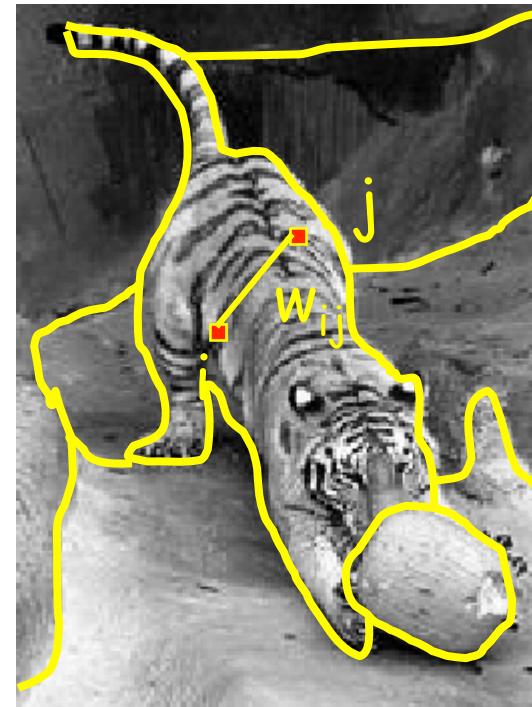
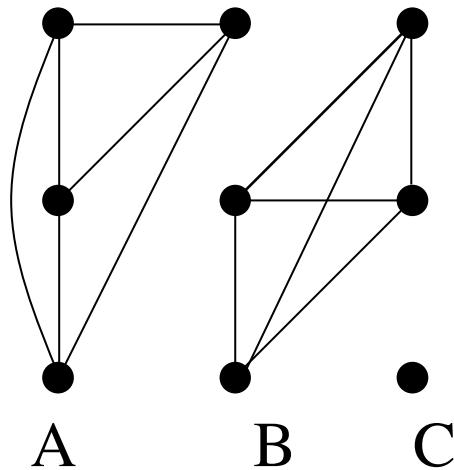


Affinity Matrix

$$W = \begin{bmatrix} 1 & .1 & .3 & 0 & 0 \\ .1 & 1 & .4 & 0 & .2 \\ .3 & .4 & 1 & .6 & .7 \\ 0 & 0 & .6 & 1 & 1 \\ 0 & .2 & .7 & 1 & 1 \end{bmatrix}$$

W_{ij} : probability that i &j belong to the same region

Segmentation by graph partitioning



- Break graph into segments
 - Delete links that cross between segments
 - Easiest to break links that have low affinity
 - similar pixels should be in the same segments
 - dissimilar pixels should be in different segments

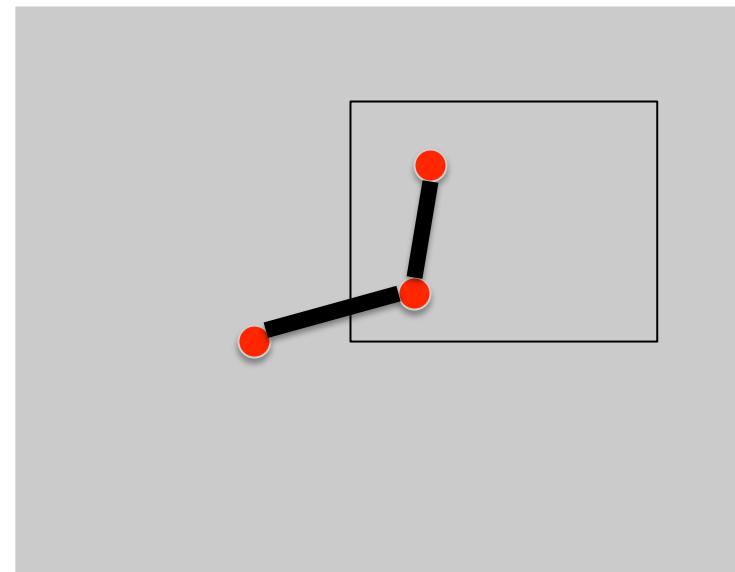
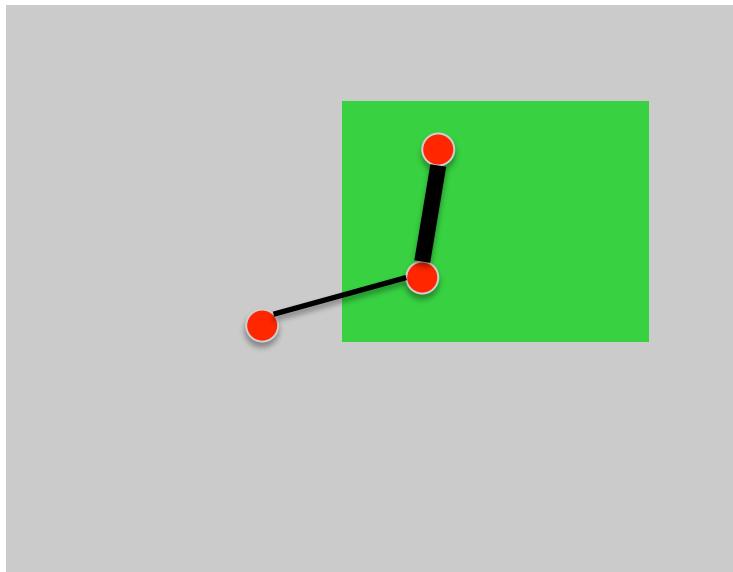
Affinity between pixels

像素 p_i 的特征

Similarities among pixel descriptors

$$W_{ij} = \exp(-\|z_i - z_j\|^2 / \sigma^2)$$

← σ = Scale factor...
it will hunt us later



Slide credit: B. Freeman and A. Torralba

Affinity between pixels

Similarities among pixel descriptors

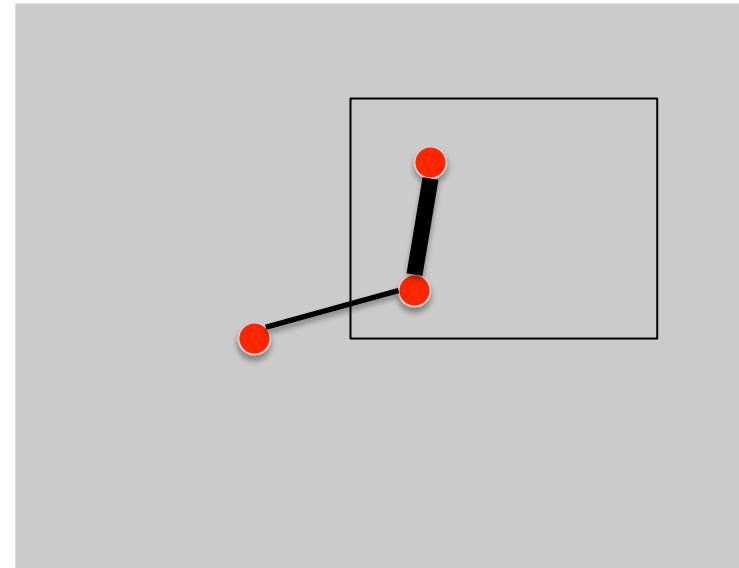
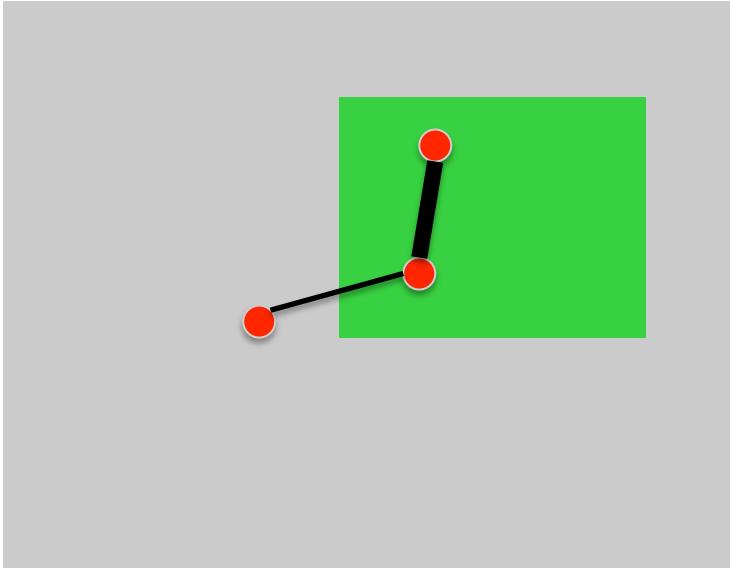
$$W_{ij} = \exp(-\|z_i - z_j\|^2 / \sigma^2)$$

Interleaving edges

σ = Scale factor...
it will hunt us later

$$W_{ij} = 1 - \max_{\text{Line between } i \text{ and } j} P_b$$

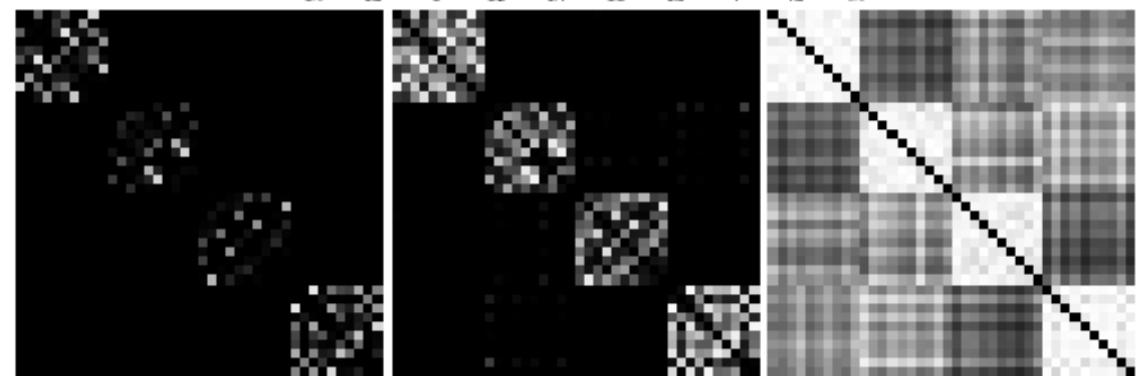
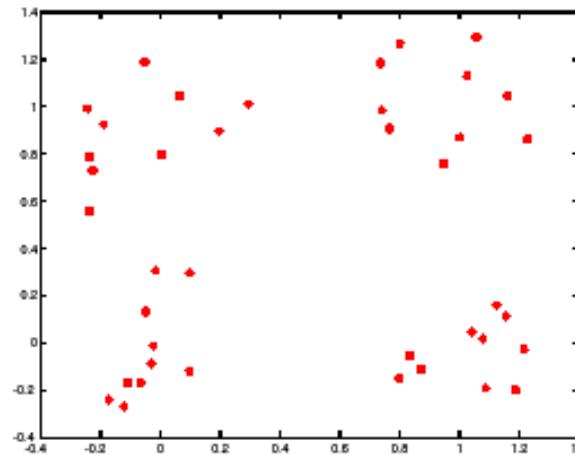
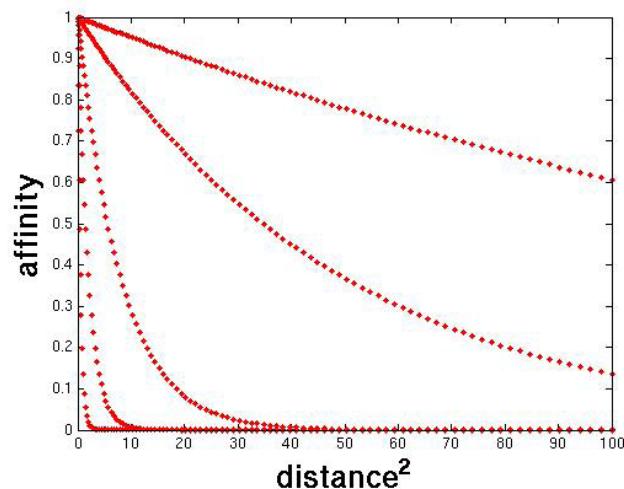
With P_b = probability of boundary



Slide credit: B. Freeman and A. Torralba

Scale affects affinity

- Small σ : group only nearby points
- Large σ : group far-away points



Slide credit: S. Lazebnik

Feature grouping by “relocalisation” of eigenvectors of the proximity matrix

British Machine Vision Conference, pp. 103-108, 1990

Guy L. Scott

Robotics Research Group

Department of Engineering Science

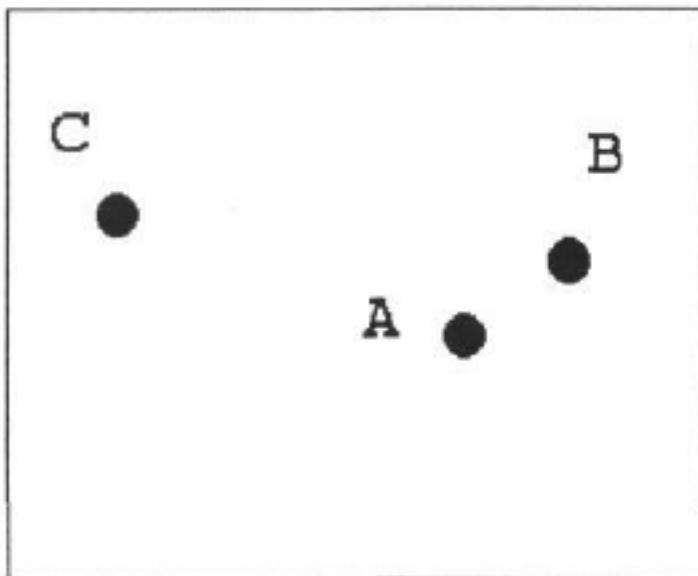
University of Oxford

H. Christopher Longuet-Higgins

University of Sussex

Falmer

Brighton



Three points in feature space

$$W_{ij} = \exp(-\| z_i - z_j \|^2 / s^2)$$

With an appropriate s

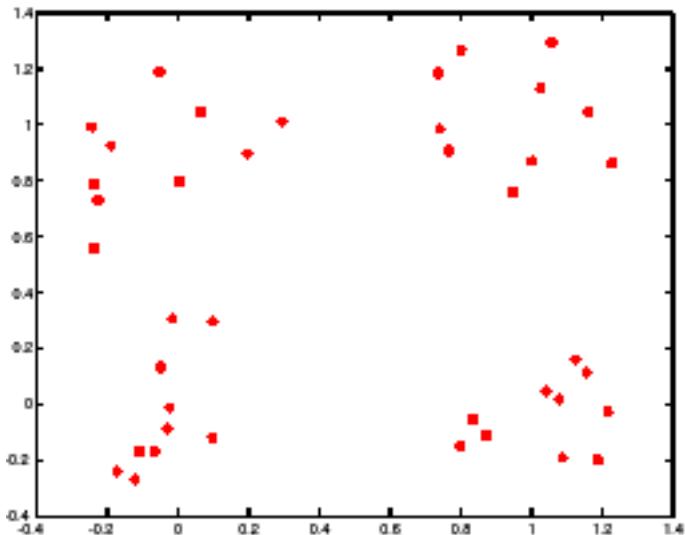
	A	B	C
A	1.00	0.63	0.03
B	0.63	1.00	0.0
C	0.03	0.0	1.00

The eigenvectors of W are:

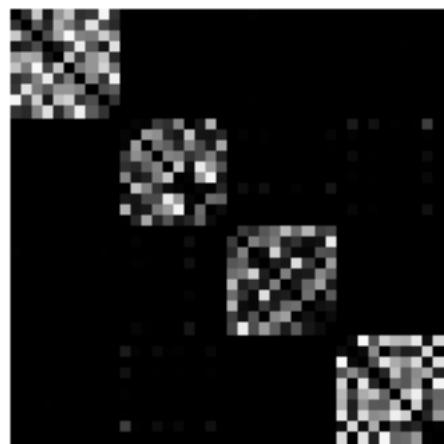
	E_1	E_2	E_3
Eigenvalues	1.63	1.00	0.37
A	-0.71	-0.01	0.71
B	-0.71	-0.05	-0.71
C	-0.04	1.00	-0.03

The first 2 eigenvectors group the points as desired...

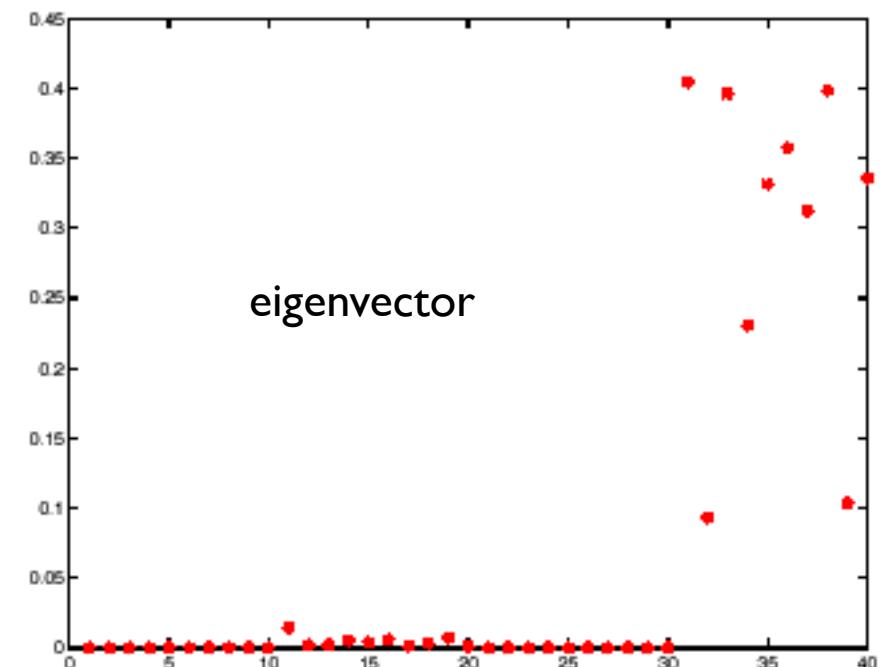
Example eigenvector



points

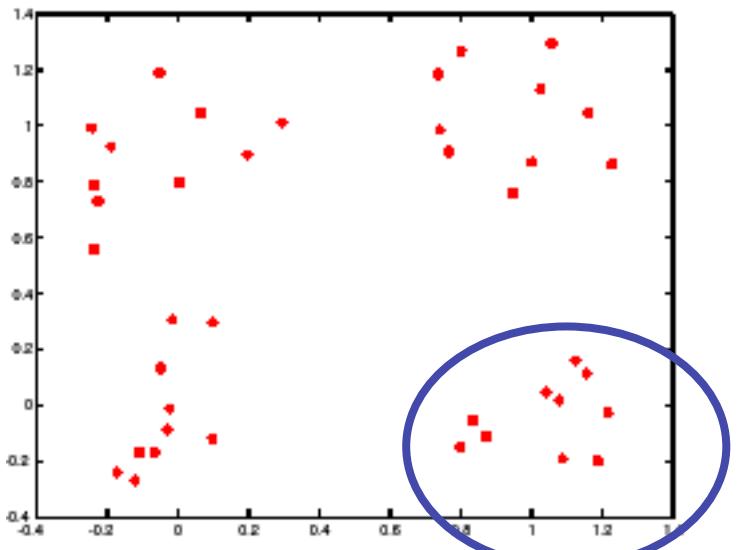


Affinity matrix

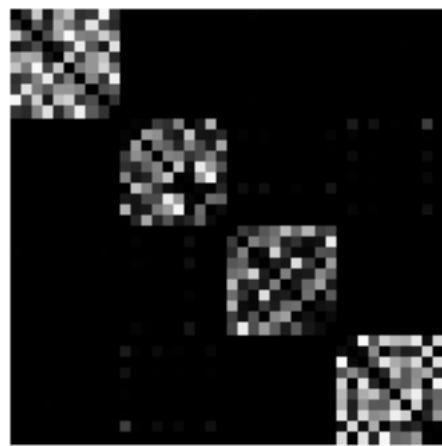


eigenvector

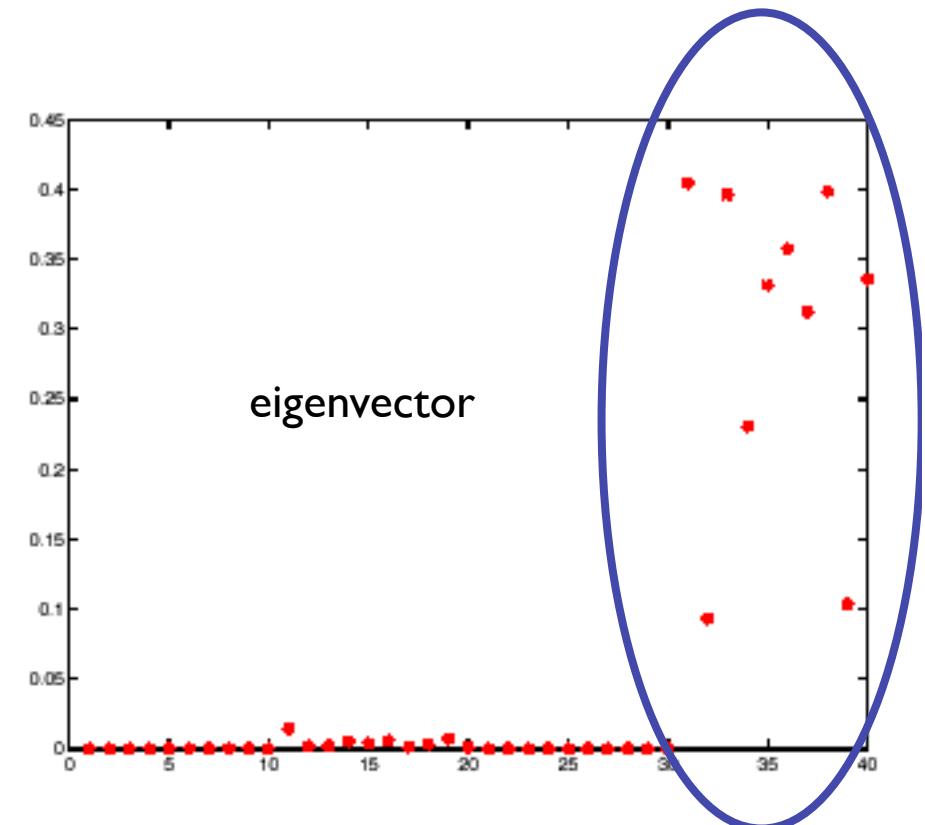
Example eigenvector



points

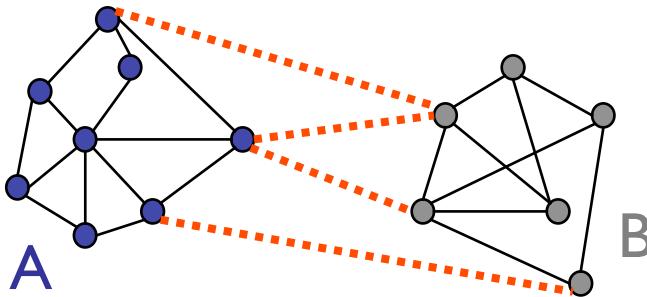


Affinity matrix



eigenvector

Graph cut



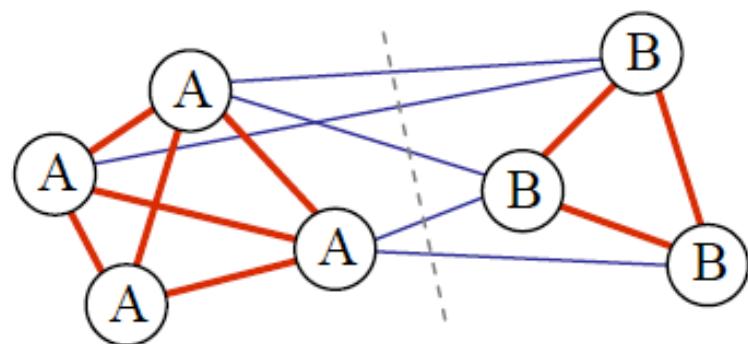
- Set of edges whose removal makes a graph disconnected
- Cost of a cut: sum of weights of cut edges
- A graph cut gives us a segmentation
 - What is a “good” graph cut and how do we find one?

Segmentation methods

- Segment foreground from background
- Histogram-based segmentation
- Segmentation as clustering
 - K-means clustering
 - Mean-shift segmentation
- **Graph-theoretic segmentation**
 - Min cut
 - Normalized cuts
- Interactive segmentation

Minimum cut

A cut of a graph G is the set of edges S such that removal of S from G disconnects G .



Cut: sum of the weight of the cut edges:

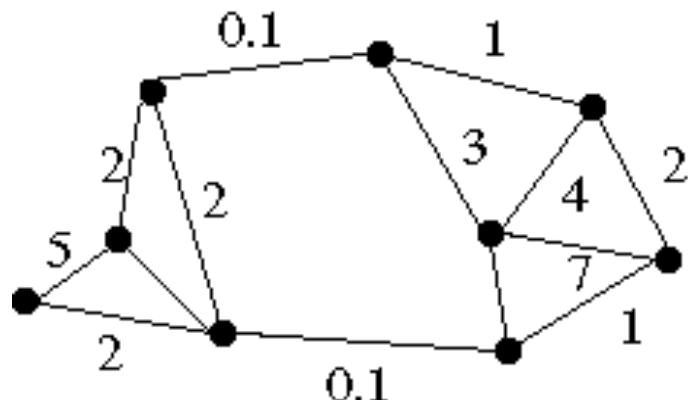
$$cut(A, B) = \sum_{u \in A, v \in B} W(u, v),$$

with $A \cap B = \emptyset$

Minimum cut

- We can do segmentation by finding the *minimum cut* in a graph
 - Efficient algorithms exist for doing this

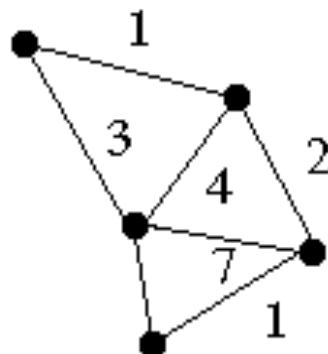
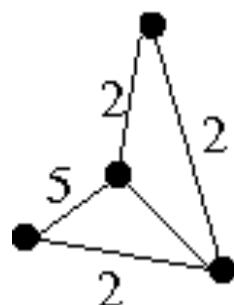
Minimum cut example



Minimum cut

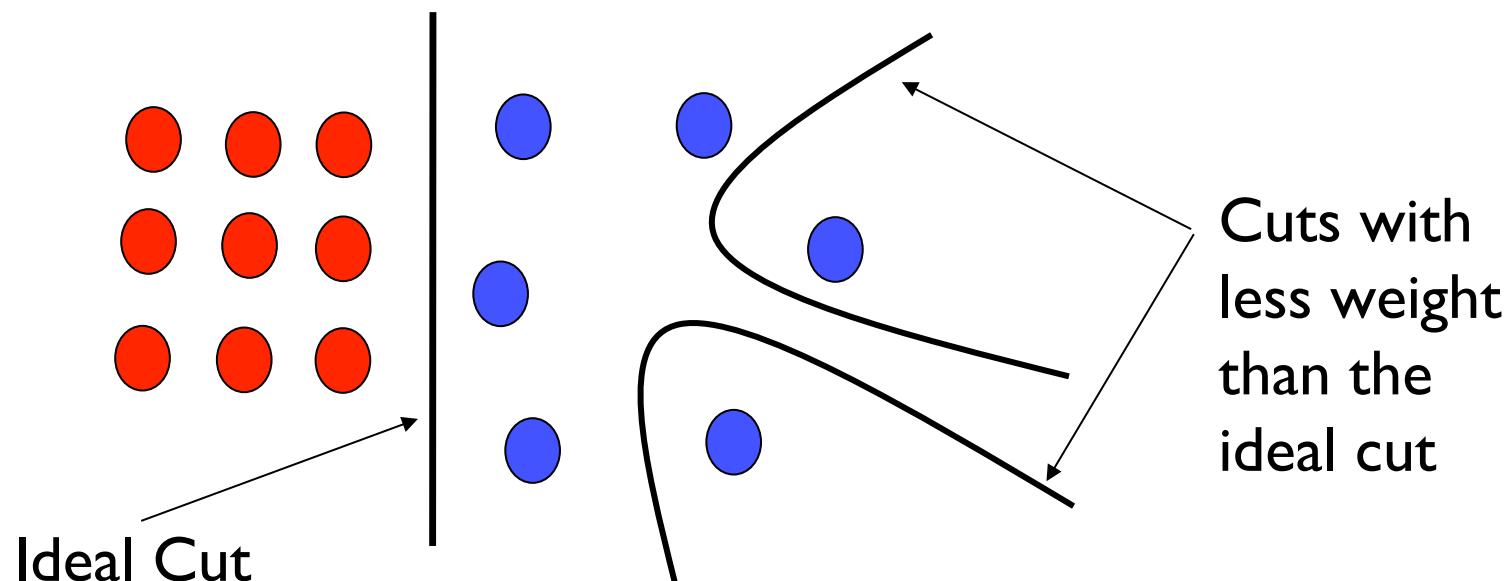
- We can do segmentation by finding the *minimum cut* in a graph
 - Efficient algorithms exist for doing this

Minimum cut example



Drawbacks of Minimum cut

- Weight of cut is directly proportional to the number of edges in the cut.

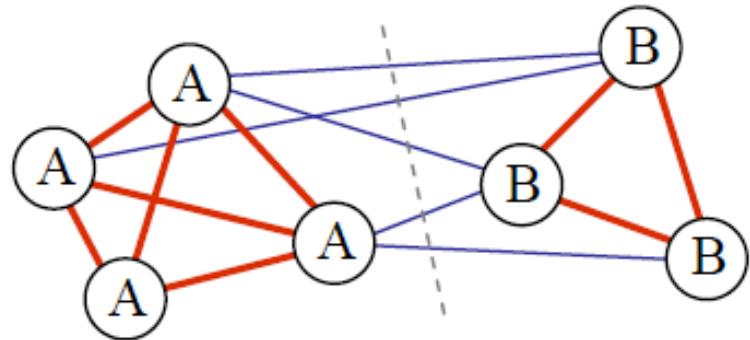


Segmentation methods

- Segment foreground from background
- Histogram-based segmentation
- Segmentation as clustering
 - K-means clustering
 - Mean-shift segmentation
- **Graph-theoretic segmentation**
 - Min cut
 - Normalized cuts
- Interactive segmentation

Normalized cuts

Write graph as V , one cluster as A and the other as B



$$Ncut(A,B) = \frac{cut(A,B)}{assoc(A,V)} + \frac{cut(A,B)}{assoc(B,V)}$$

$cut(A,B)$ is sum of weights with one end in A and one end in B

$$cut(A,B) = \sum_{u \in A, v \in B} W(u,v),$$

with $A \cap B = \emptyset$

$assoc(A,V)$ is sum of all edges with one end in A .

$$assoc(A,V) = \sum_{u \in A, v \in V} W(u,v)$$

A and B not necessarily disjoint

J. Shi and J. Malik. [Normalized cuts and image segmentation](#). PAMI 2000

Slide credit: B. Freeman and A. Torralba

Normalized cut

- Let W be the adjacency matrix of the graph
- Let D be the diagonal matrix with diagonal entries
$$D(i, i) = \sum_j W(i, j)$$
- Then the normalized cut cost can be written as

$$\frac{y^T (D - W) y}{y^T D y}$$

where y is an indicator vector whose value should be 1 in the i th position if the i th feature point belongs to A and a negative constant otherwise

注：根据 y 来划分 A and B - 对应 $y(i)=1, i \in A$,
其他点属于B

Normalized cut

- Finding the exact minimum of the normalized cut cost is NP-complete, but if we relax y to take on arbitrary values, then we can minimize the relaxed cost by solving the *generalized eigenvalue problem* $(D - W)y = \lambda Dy$
- The solution y is given by the generalized eigenvector corresponding to the second smallest eigenvalue
- Intuitively, the i th entry of y can be viewed as a “soft” indication of the component membership of the i th feature
 - Can use 0 or median value of the entries as the splitting point (threshold), or find threshold that minimizes the Ncut cost

例如，设阈值为0，则有 $y(i) > 0$, i点 A, 否则点 B

Normalized cut algorithm

1. Given an image or image sequence, set up a weighted graph $\mathbf{G} = (\mathbf{V}, \mathbf{E})$, and set the weight on the edge connecting two nodes being a measure of the similarity between the two nodes.
2. Solve $(\mathbf{D} - \mathbf{W})\mathbf{x} = \lambda \mathbf{D}\mathbf{x}$ for eigenvectors with the smallest eigenvalues.
3. Use the eigenvector with second smallest eigenvalue to bipartition the graph.
4. Decide if the current partition should be sub-divided, and recursively repartition the segmented parts if necessary.

对每个分割，重复以上Ncut，得到更多分割

Global optimization

- In this formulation, the segmentation becomes a global process.
- Decisions about what is a boundary are not local (as in Canny edge detector)

Boundaries of image regions defined by a number of attributes

- Brightness/color
- Texture
- Motion
- Stereoscopic depth
- Familiar configuration



[Malik]

Slide credit: B. Freeman and A. Torralba

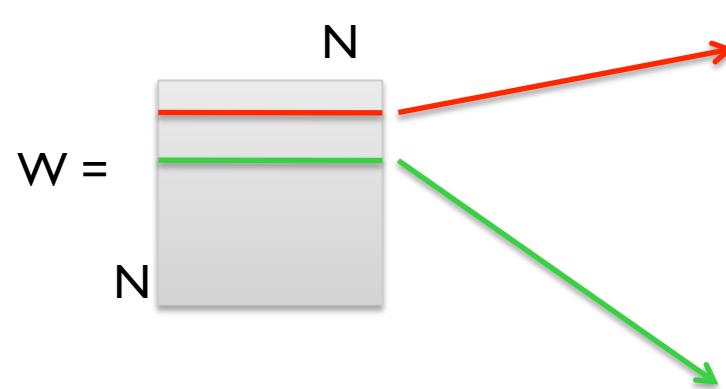
Example

Affinity:

$$w_{ij} = e^{\frac{-\|F(i) - F(j)\|_2^2}{\sigma_I}} * \begin{cases} e^{\frac{-\|X(i) - X(j)\|_2^2}{\sigma_X}} & \text{if } \|X(i) - X(j)\|_2 < r \\ 0 & \text{otherwise} \end{cases}$$



$$N_{\text{pixels}} = \text{ncols} * \text{nrows}$$



Slide credit: B. Freeman and A. Torralba

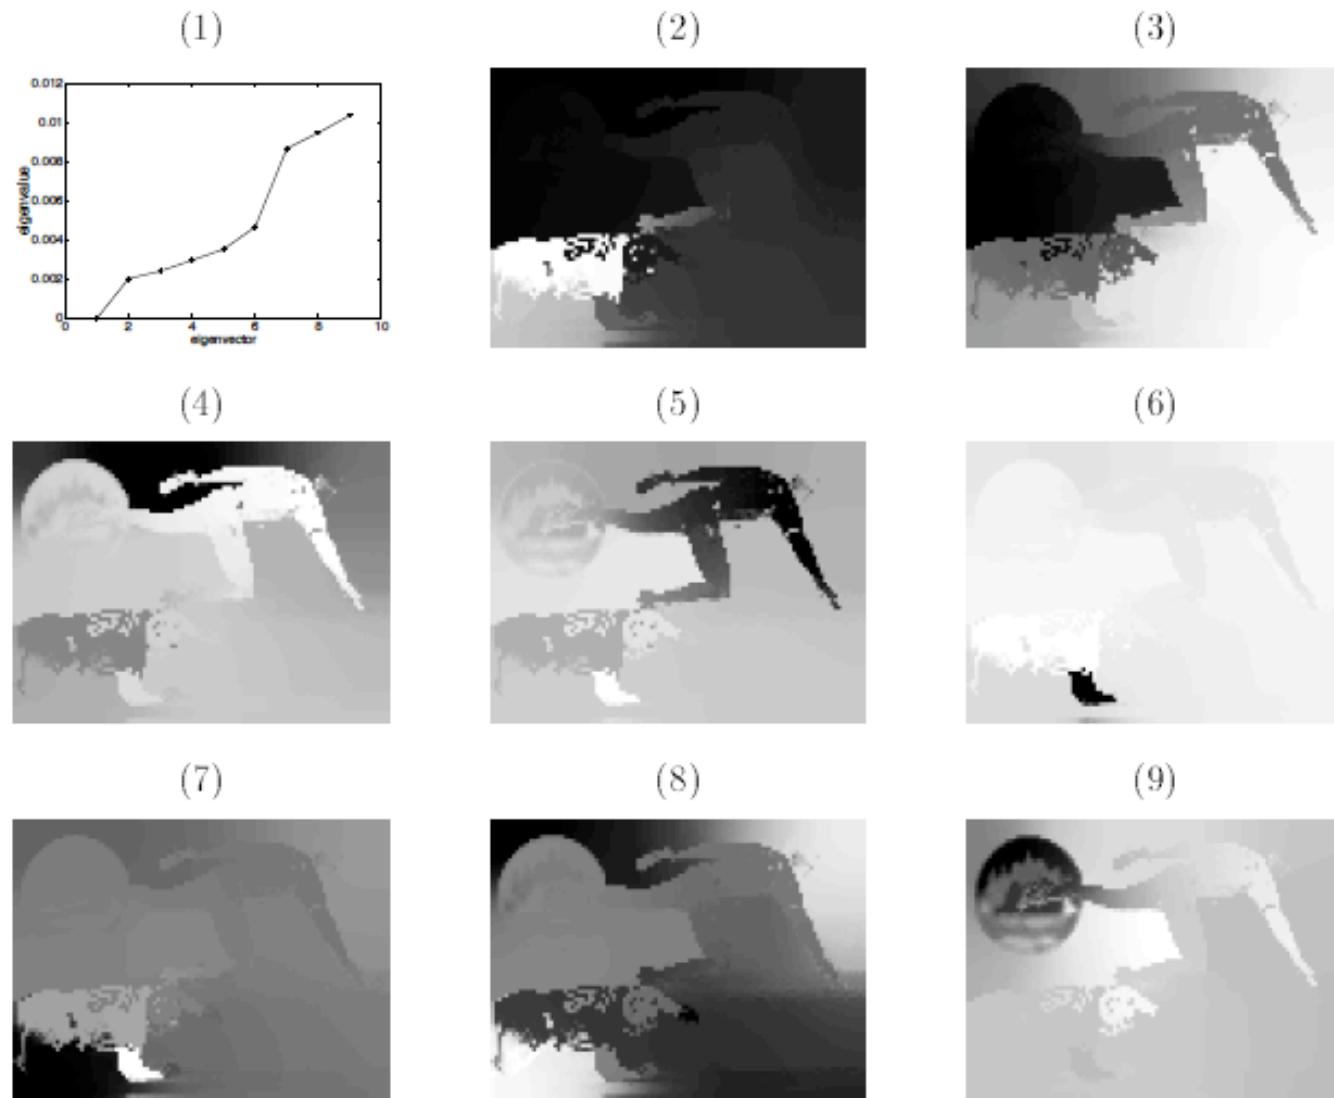


Figure 12: Subplot (1) plots the smallest eigenvectors of the generalized eigenvalue system (11). Subplot (2) - (9) shows the eigenvectors corresponding to the 2nd smallest to the 9th smallest eigenvalues of the system. The eigenvectors are reshaped to be the size of the image.

Slide credit: B. Freeman and A. Torralba

注：对于 $N \times M$ 图像，特征向量 y 的维数是 $M \times N$

Brightness Image Segmentation



converge. On the 100×120 test images shown here, the normalized cut algorithm takes about 2 minutes on Intel Pentium 200MHz machines.

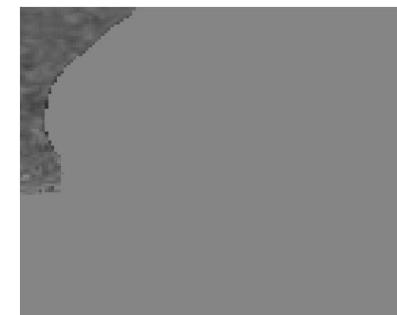
A multiresolution implementation can be used to reduce this running time further on larger images. In our current experiments, with this implementation, the running time on a 300×400 image can be reduced to about 20 seconds on Intel Pentium 300MHz machines. Furthermore, the bottleneck of the computation, a sparse matrix-vector

Brightness Image Segmentation



<http://www.cs.berkeley.edu/~malik/papers/SM-ncut.pdf>

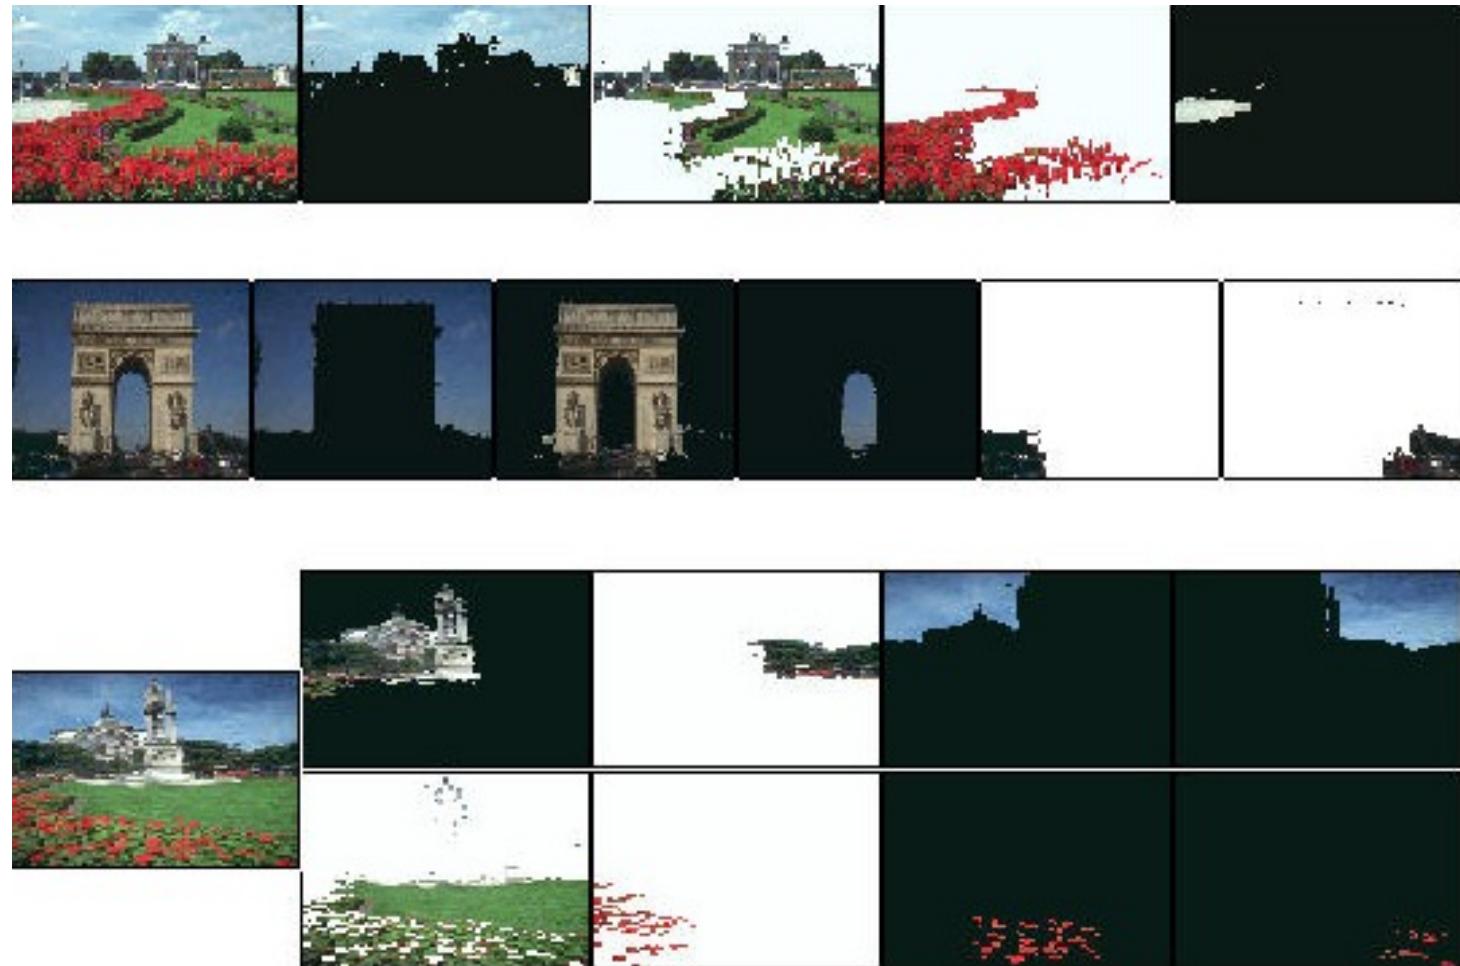
Slide credit: B. Freeman and A. Torralba



<http://www.cs.berkeley.edu/~malik/papers/SM-ncut.pdf>

Slide credit: B. Freeman and A. Torralba

Results on color segmentation

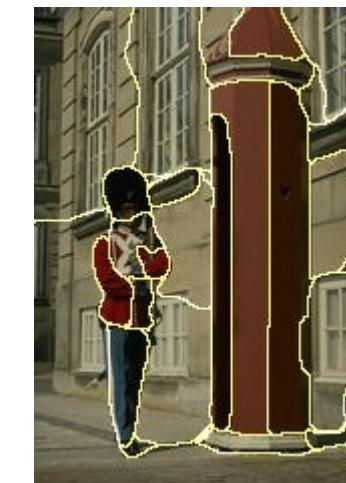
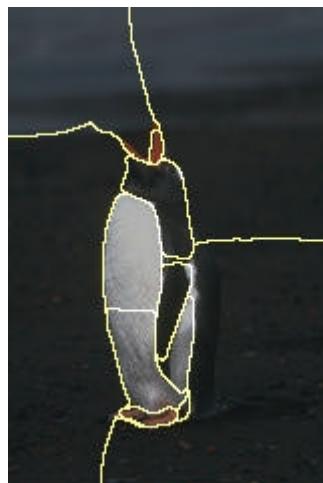
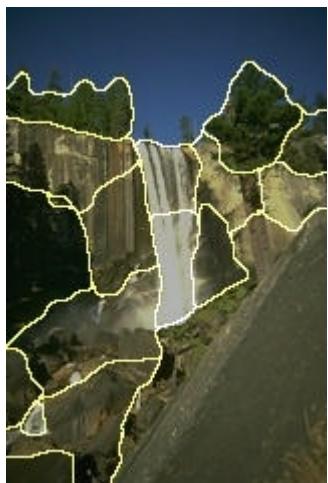
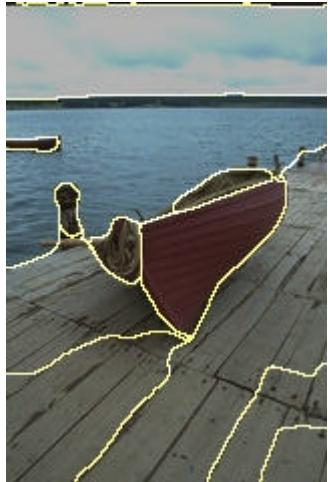


Example results



Slide credit: S. Lazebnik

Results: Berkeley Segmentation Engine



<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Slide credit: S. Lazebnik

Normalized cuts: Pro and con

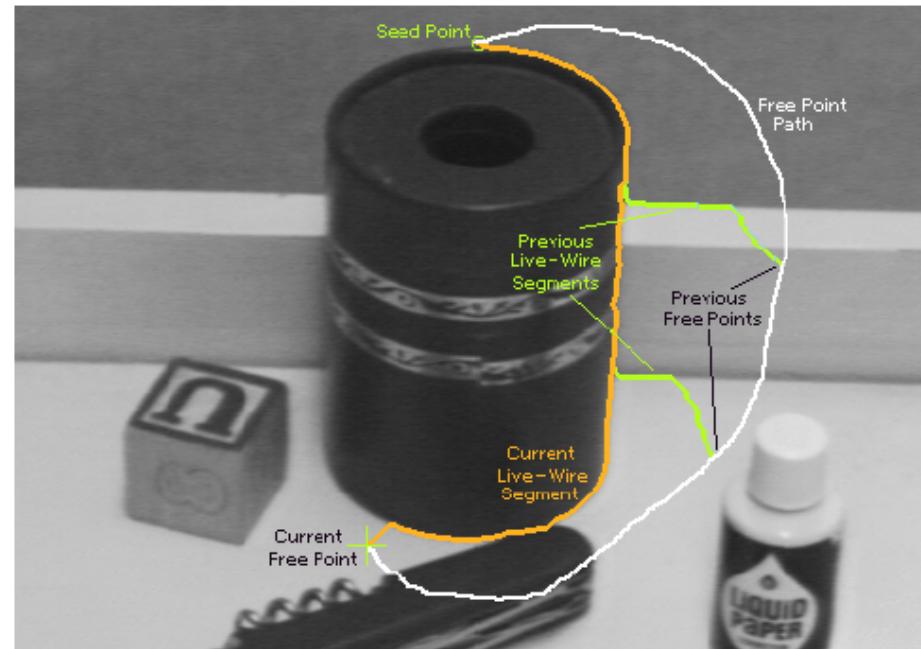
- Pros
 - Generic framework, can be used with many different features and affinity formulations
- Cons
 - High storage requirement and time complexity
 - Bias towards partitioning into equal segments

Segmentation methods

- Segment foreground from background
- Histogram-based segmentation
- Segmentation as clustering
 - K-means clustering
 - Mean-shift segmentation
- Graph-theoretic segmentation
 - Min cut
 - Normalized cuts
- Interactive segmentation

Intelligent Scissors [Mortensen 95]

- Approach answers a basic question
 - Q: how to find a path from seed to mouse that follows object boundary as closely as possible?

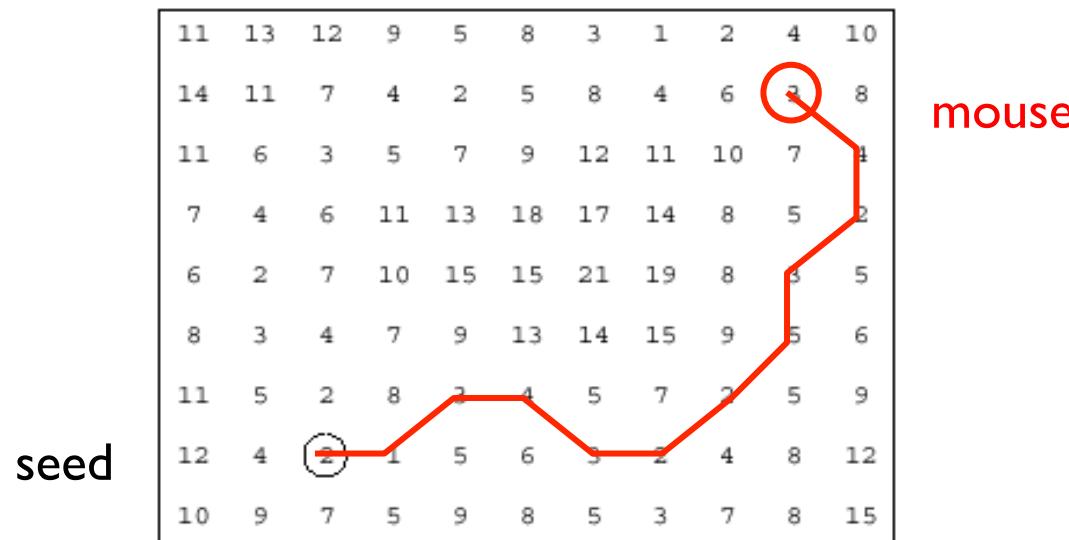


Mortensen and Barrett, Intelligent Scissors for Image Composition, Proc. 22nd annual conference on Computer graphics and interactive techniques, 1995

Figure 2: Image demonstrating how the live-wire segment adapts and snaps to an object boundary as the free point moves (via cursor movement). The path of the free point is shown in white. Live-wire segments from previous free point positions (t_0 , t_1 , and t_2) are shown in green.

Intelligent Scissors

- Basic Idea
 - Define edge score for each pixel
 - edge pixels have low cost
 - Find lowest cost path from seed to mouse



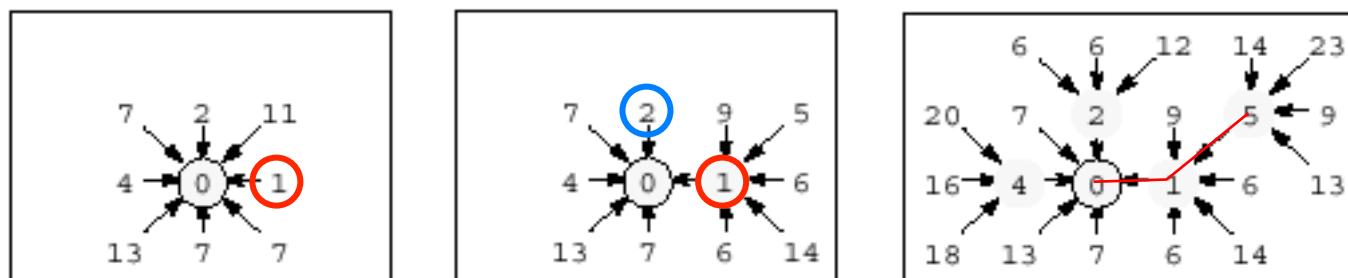
Questions

- How to define costs?
- How to find the path?

Path Search (basic idea)

- Graph Search Algorithm
 - Computes minimum cost path from seed to *all other pixels*

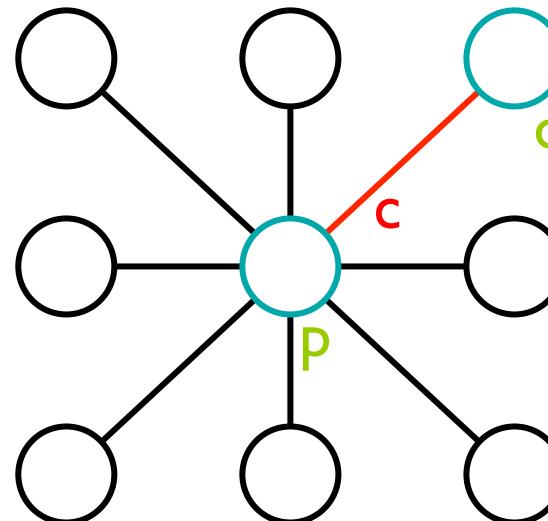
11	13	12	9	5	8	3	1	2	4	10
14	11	7	4	2	5	8	4	6	3	8
11	6	3	5	7	9	12	11	10	7	4
7	4	6	11	13	18	17	14	8	5	2
6	2	7	10	15	15	21	19	8	3	5
8	3	4	7	9	13	14	15	9	5	6
11	5	2	8	3	4	5	7	2	5	9
12	4	2	1	5	6	3	2	4	8	12
10	9	7	5	9	8	5	3	7	8	15



Slide credit: S. Seitz

How does this really work?

- Treat the image as a graph



Graph

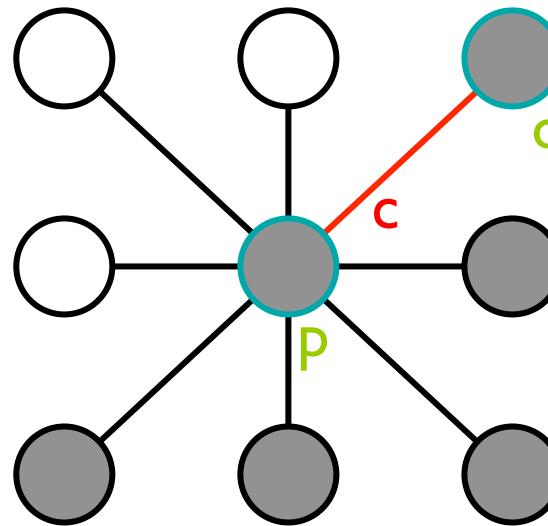
- node for every pixel **p**
- link between every adjacent pair of pixels, **p,q**
- cost **c** for each link

Note: each *link* has a cost

- this is a little different than the figure before where each pixel had a cost

Defining the costs

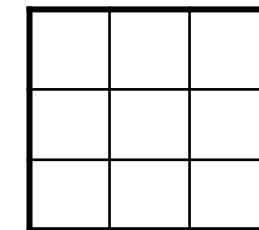
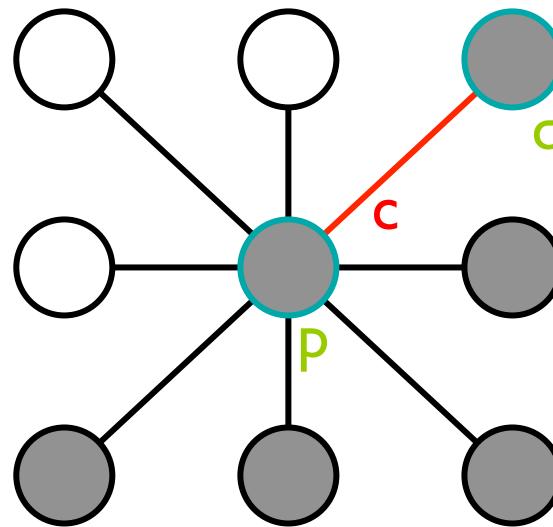
- Treat the image as a graph



Want to hug image edges: how to define cost of a link?

- the link should follow the intensity edge
 - want intensity to change rapidly \perp to the link 灰度值的差值
- $c \approx -|\text{difference of intensity } \perp \text{ to link}|$

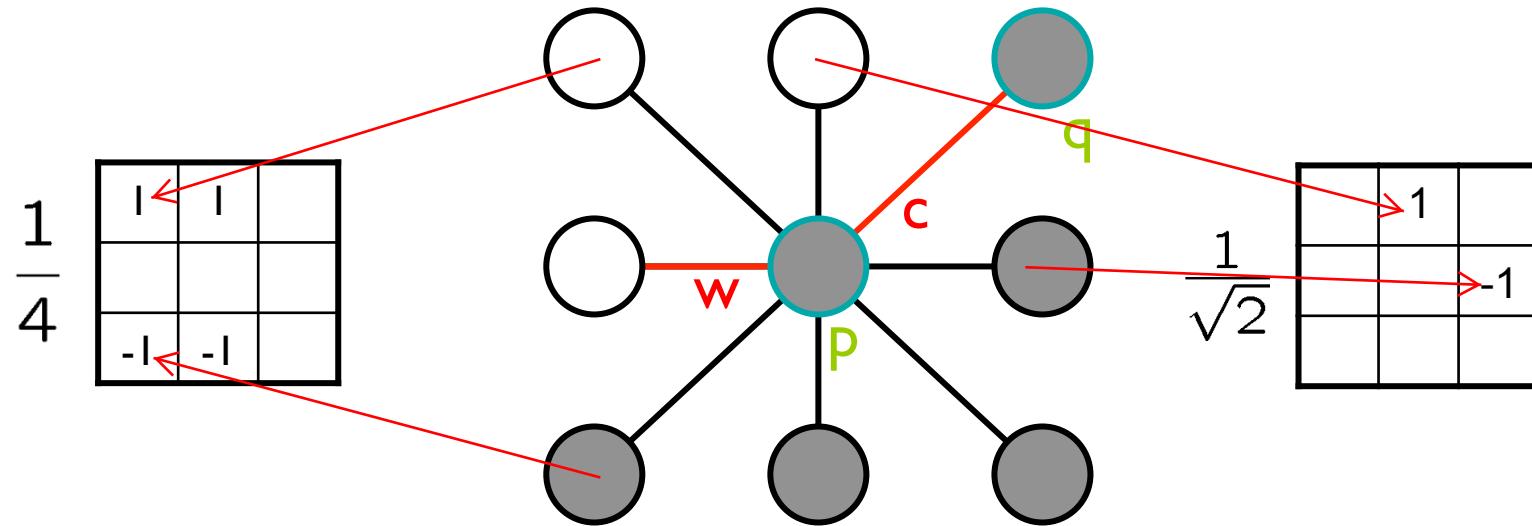
Defining the costs



- c can be computed using a cross-correlation filter
 - assume it is centered at P
- Also typically scale c by its length
 - set $c = (\max - |filter\ response|)$
 - where $\max = \text{maximum } |filter\ response| \text{ over all pixels in the image}$

Slide credit: S. Seitz

Defining the costs



- c can be computed using a cross-correlation filter
 - assume it is centered at P

注：如果响应很大，则是边缘的代价（风险）就很小

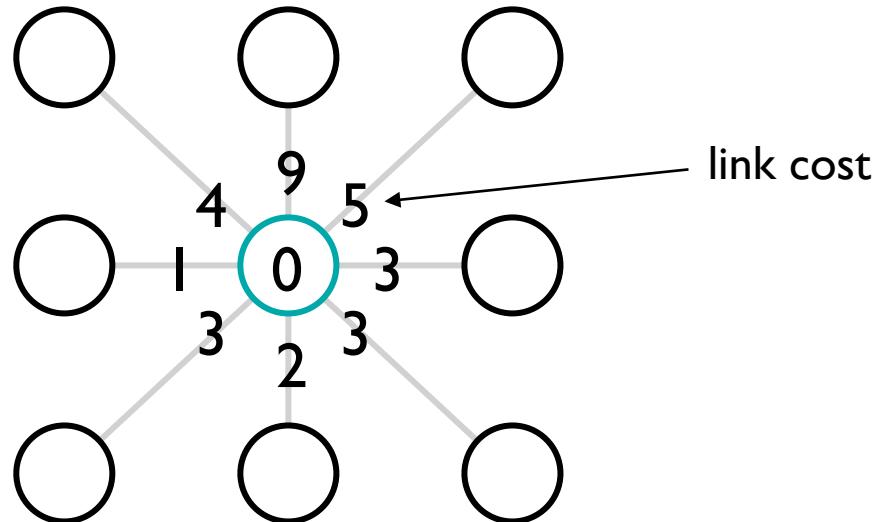
- Also typically scale c by its length

- set $c = (\max - |\text{filter response}|)$

- where $\max = \text{maximum } |\text{filter response}|$ over all pixels in the image

迪杰斯特拉（1972年图灵奖获得者）最短路径算法

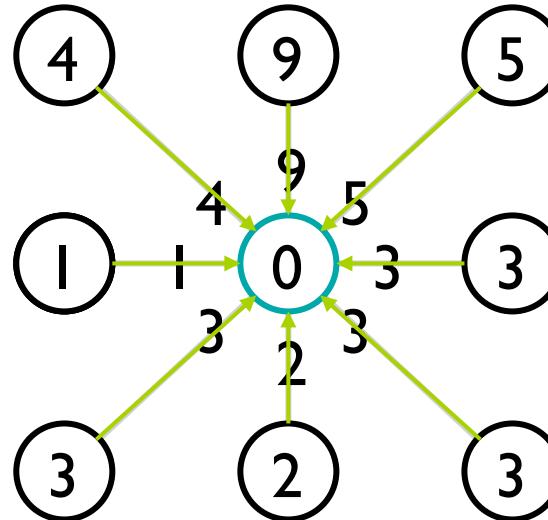
Dijkstra's shortest path algorithm



Algorithm

1. init node costs to ∞ , set p = seed point, $\text{cost}(p) = 0$
2. expand p as follows:
 - for each of p 's neighbors q that are not expanded
 - » set $\text{cost}(q) = \min(\text{cost}(p) + c_{pq}, \text{cost}(q))$

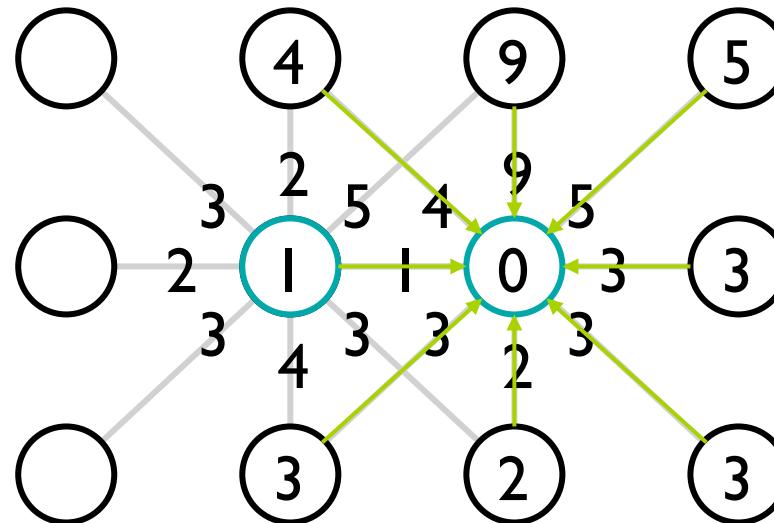
Dijkstra's shortest path algorithm



Algorithm

1. init node costs to ∞ , set p = seed point, $\text{cost}(p) = 0$
2. expand p as follows:
 - for each of p 's neighbors q that are not expanded
 - » set $\text{cost}(q) = \min(\text{cost}(p) + c_{pq}, \text{cost}(q))$
 - » if q 's cost changed, make q point back to p
 - » put q on the ACTIVE list (if not already there)

Dijkstra's shortest path algorithm

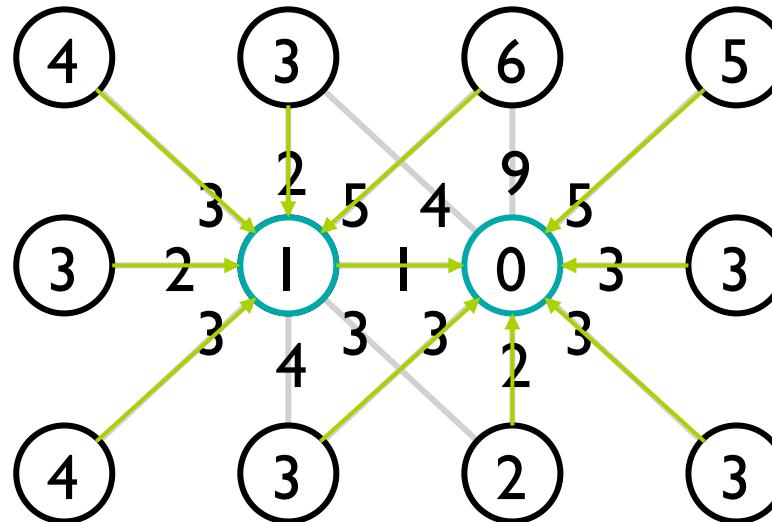


Algorithm

1. init node costs to ∞ , set p = seed point, $\text{cost}(p) = 0$
2. expand p as follows:
 - for each of p 's neighbors q that are not expanded
 - » set $\text{cost}(q) = \min(\text{cost}(p) + c_{pq}, \text{cost}(q))$
 - » if q 's cost changed, make q point back to p
 - » put q on the ACTIVE list (if not already there)
3. set r = node with minimum cost on the ACTIVE list
4. repeat Step 2 for $p = r$

ACTIVE list 包含 p 的所有邻点及其代价

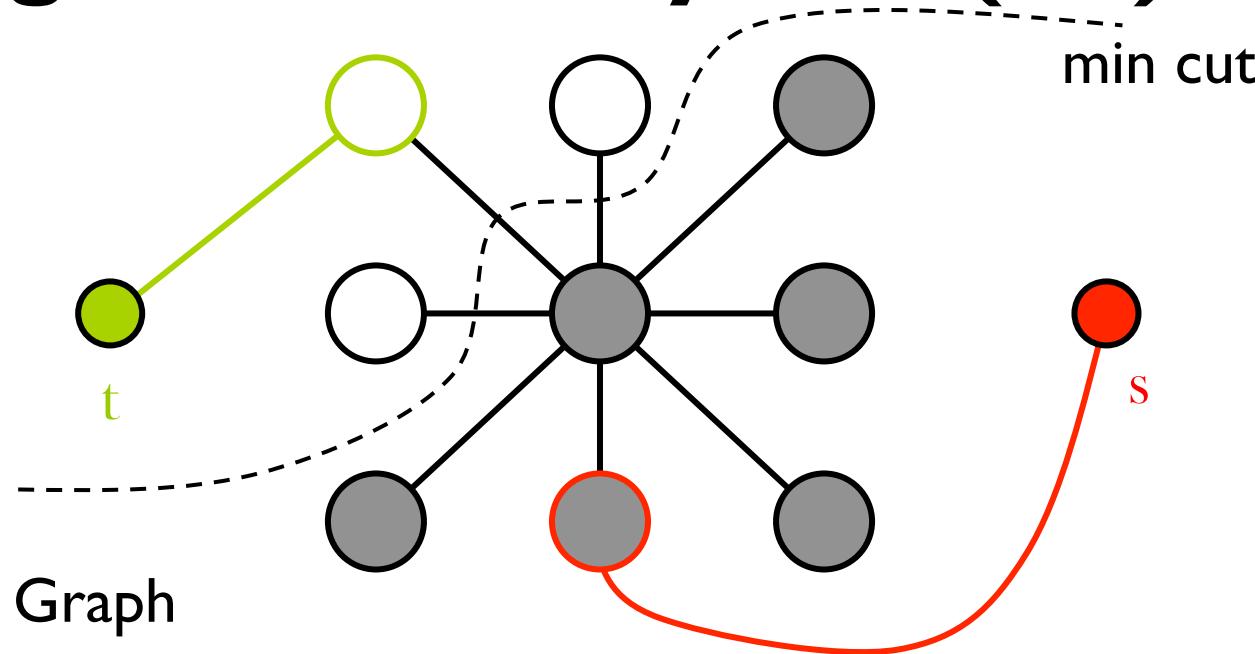
Dijkstra's shortest path algorithm



Algorithm

1. init node costs to ∞ , set $p = \text{seed point}$, $\text{cost}(p) = 0$
2. expand p as follows:
 - for each of p 's neighbors q that are not expanded
 - » set $\text{cost}(q) = \min(\text{cost}(p) + c_{pq}, \text{cost}(q))$
 - » if q 's cost changed, make q point back to p
 - » put q on the ACTIVE list (if not already there)
 - 3. set $r = \text{node with minimum cost on the ACTIVE list}$
 - 4. repeat Step 2 for $p = r$

Segmentation by min (s-t) cut

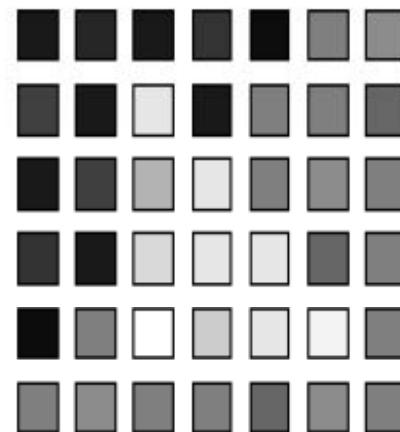


- Graph
 - node for each pixel, link between pixels
 - specify a few pixels as foreground and background
 - create an infinite cost link from each bg pixel to the “t” node
 - create an infinite cost link from each fg pixel to the “s” node
 - compute min cut that separates s from t
 - how to define link cost between neighboring pixels?

计算机视觉中的多标记问题

计算机视觉的很多问题可以看作一个最优标记 (labeling) 问题。如：

- 图像分割
- 图像恢复
- 立体视觉
- 三维重建



(a) An image

0	0	0	0	0	1	1
0	0	2	0	1	1	1
0	0	2	2	1	1	1
0	0	2	2	2	1	1
0	1	2	2	2	2	1
1	1	1	1	1	1	1

(b) A labeling

Key Issues

同样面临的两个问题：

- 什么是最优标记准则？
- 有没有有效求解算法

多标记问题的能量极小化模型

我们可以通过构造一个如下的能量函数来得到最优标记准则。

$$E(f) = E_{data}(f) + E_{smooth}(f) = \sum_{p \in P} D_p(f_p) + \sum_{(p,q) \in P} V_{p,q}(f_p, f_q)$$

$f: P \rightarrow L$ 的映射。P是像素点集，L是标记集。

Data项表示给每个像素点赋予标记（label）的代价

Smooth项表示每两个相邻的像素分别赋予标记 f_p 和 f_q 的代价

问题表达

我们希望最优标记应使得能量函数取最小值，即可建立下面最优化模型

$$\begin{aligned} & \min E(f) \\ & s.t. \quad f_p \in L \end{aligned}$$

下面具体建立各问题的能量极小化模型。

运用Graph Cut的问题求解

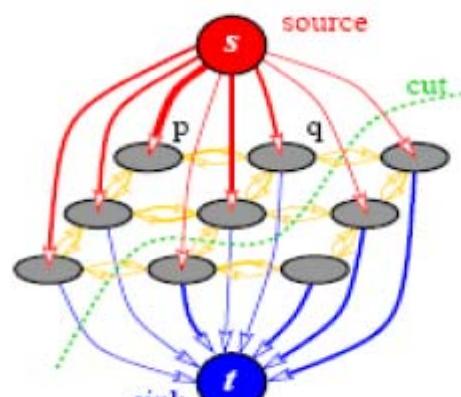
- graph cut与能量函数的对应关系
- 运用graph cut算法求解能量极小化问题
 - 两标记问题
 - 多标记问题

图割与能量函数的对应关系 (1)

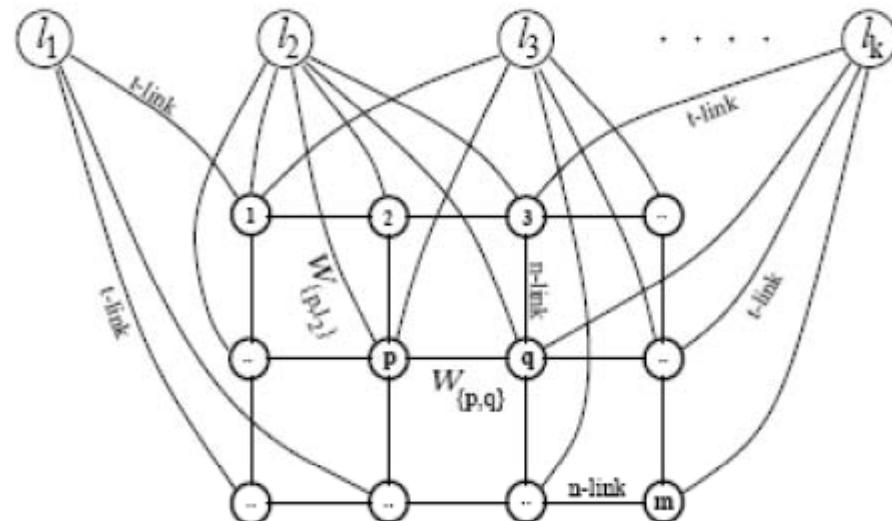
- 能量函数

$$E(f) = E_{data}(f) + E_{smooth}(f) = \sum_{p \in P} D_p(f_p) + \sum_{(p,q) \in N} V_{p,q}(f_p, f_q)$$

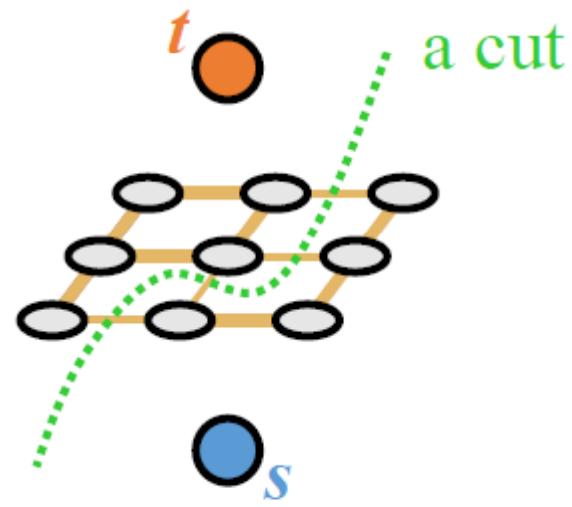
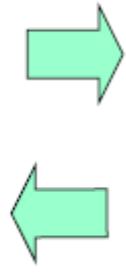
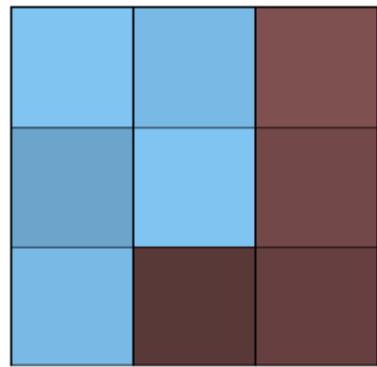
↔ ↔



(b) A cut on \mathcal{G}



图割与能量函数的对应关系 (2)



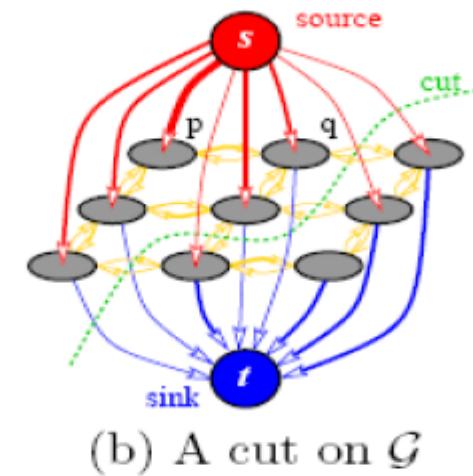
$$E(f) = E_{data}(f) + E_{smooth}(f) = \sum_{p \in P} D_p(f_p) + \sum_{(p,q) \in N} V_{p,q}(f_p, f_q)$$

图割与能量函数的对应关系 (3)

- 运用graph cut算法求解能量极小化问题

- 两标记问题

对于两标记问题，最小能量对应于图的最小割。图论中已有经典的算法，可以求得一个图的最小割，从而得到极小能量。具体算法见参考文献[1]、[3]。



图割与能量函数的对应关系 (4)

- 多标记问题

当标记数量大于2时，已经证明该问题是NP-hard问题。故很难求得该问题的全局极小值。

Boykov等^[4]构造了两个运用最小割求解该类能量函数的近似极小值的算法：

$\alpha - \beta$ swap

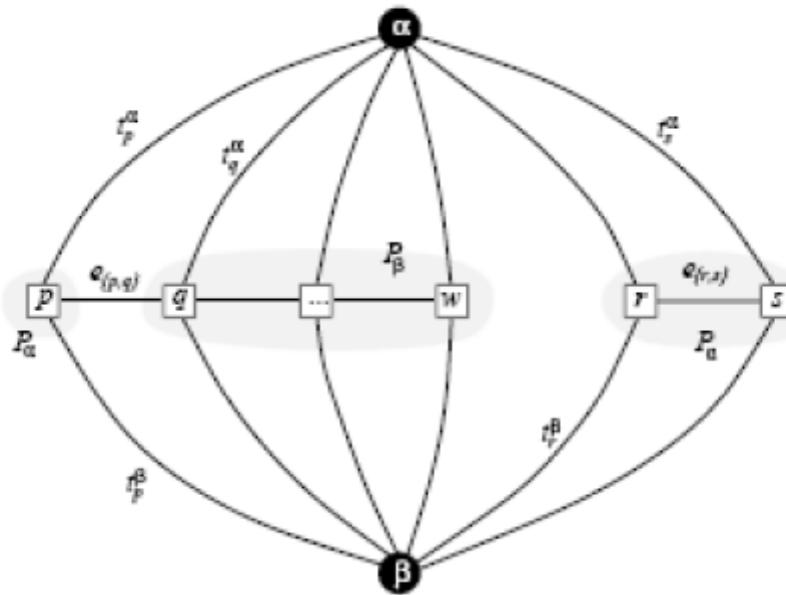
α expansion

这两个算法运算速度快，且能得到比较好的结果，从而得到了广泛应用，并使得用能量极小化模型和图割来处理计算机视觉中的一些问题成为目前的一个研究热点。

算法: $\alpha - \beta$ swap (1)

1. Start with an arbitrary labeling f
2. Set success := 0
3. For each pair of labels $\{\alpha, \beta\} \subset \mathcal{L}$
 - 3.1. Find $\hat{f} = \arg \min E(f')$ among f' within one α - β swap of f (Section 3)
 - 3.2. If $E(\hat{f}) < E(f)$, set $f := \hat{f}$ and success := 1
4. If success = 1 goto 2
5. Return f

算法： $\alpha - \beta$ swap (2)



edge	weight	for
t_p^α	$D_p(\alpha) + \sum_{\substack{q \in N_p \\ q \notin \mathcal{P}_{\alpha\beta}}} V_{\{p,q\}}(\alpha, f_q)$	$p \in \mathcal{P}_{\alpha\beta}$
t_p^β	$D_p(\beta) + \sum_{\substack{q \in N_p \\ q \notin \mathcal{P}_{\alpha\beta}}} V_{\{p,q\}}(\beta, f_q)$	$p \in \mathcal{P}_{\alpha\beta}$
$e_{\{p,q\}}$	$V_{\{p,q\}}(\alpha, \beta)$	$\begin{cases} \{p,q\} \in \mathcal{N} \\ p, q \in \mathcal{P}_{\alpha\beta} \end{cases}$

$$f_p^{\mathcal{C}} = \begin{cases} \alpha & \text{if } t_p^\alpha \in \mathcal{C} \text{ for } p \in \mathcal{P}_{\alpha\beta} \\ \beta & \text{if } t_p^\beta \in \mathcal{C} \text{ for } p \in \mathcal{P}_{\alpha\beta} \\ f_p & \text{for } p \in \mathcal{P}, p \notin \mathcal{P}_{\alpha\beta} . \end{cases}$$

算法: α expansion (1)

1. Start with an arbitrary labeling f
2. Set success := 0
3. For each label $\alpha \in \mathcal{L}$
 - 3.1. Find $\hat{f} = \arg \min E(f')$ among f' within one α -expansion of f (Section 4)
 - 3.2. If $E(\hat{f}) < E(f)$, set $f := \hat{f}$ and success := 1
4. If success = 1 goto 2
5. Return f

算法： a expansion (2)

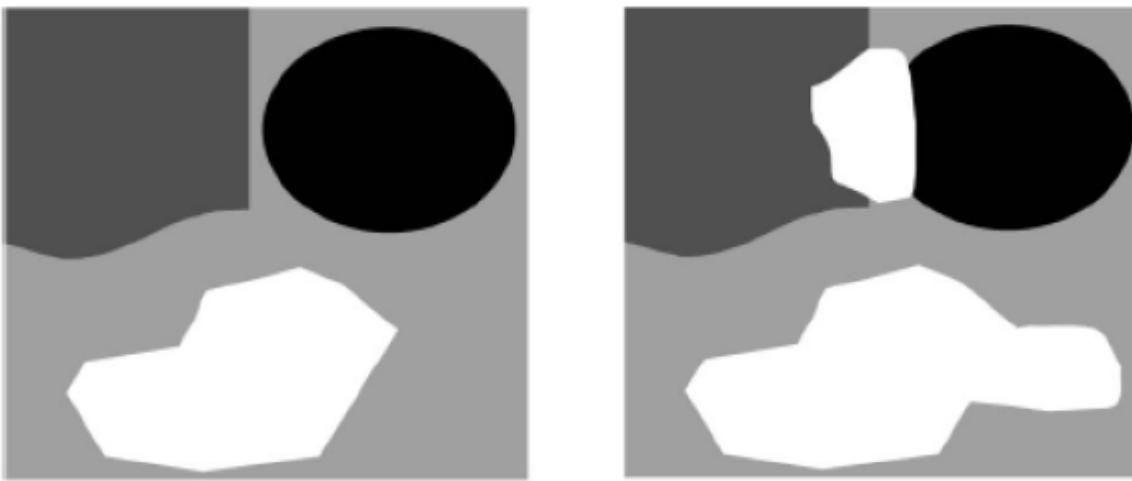


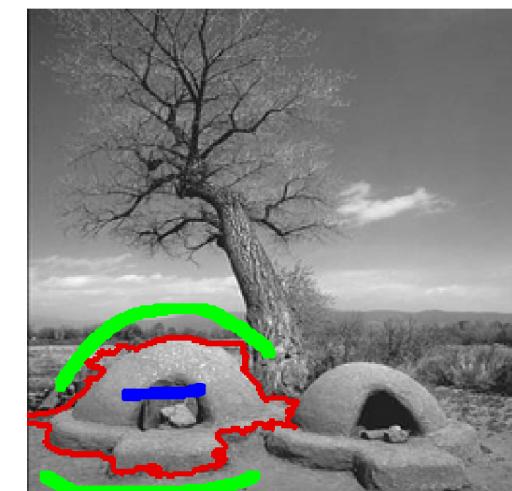
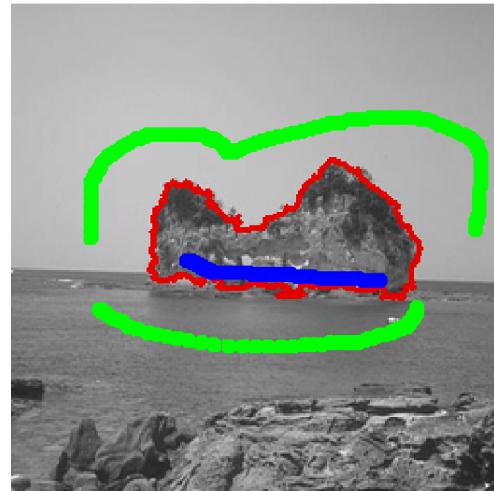
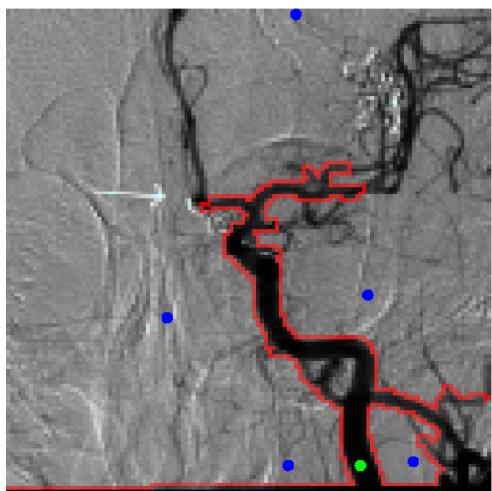
Fig. 1. An example of an expansion move. The labeling on the right is a white-expansion move from the labeling on the left.

参考文献

- [1] 楼世博等, 图论及其应用. 人民邮电出版社, 1982.
- [2] Jianbo Shi, Jitendra Malik, Normalized cuts and image segmentation[J]. IEEE Trans on PAMI, 2000; 22(8): 888-905.
- [3] Yuri Boykov and Vladimir Kolmogorov. An experimental comparison of min-cut/max-flow algorithms for energy minimization in vision. In International Workshop on Energy Minimization Methods in Computer Vision and Pattern Recognition (EMMCVPR), number 2134 in LNCS, pages 359–374, Sophia Antipolis, France, September 2001. SpringerVerlag.
- [4] Yuri Boykov, Olga Veksler, and Ramin Zabih. Fast approximate energy minimization via graph cuts. IEEE Transactions on Pattern Analysis and Machine Intelligence, 23(11):1222–1239, November 2001.

Random Walker

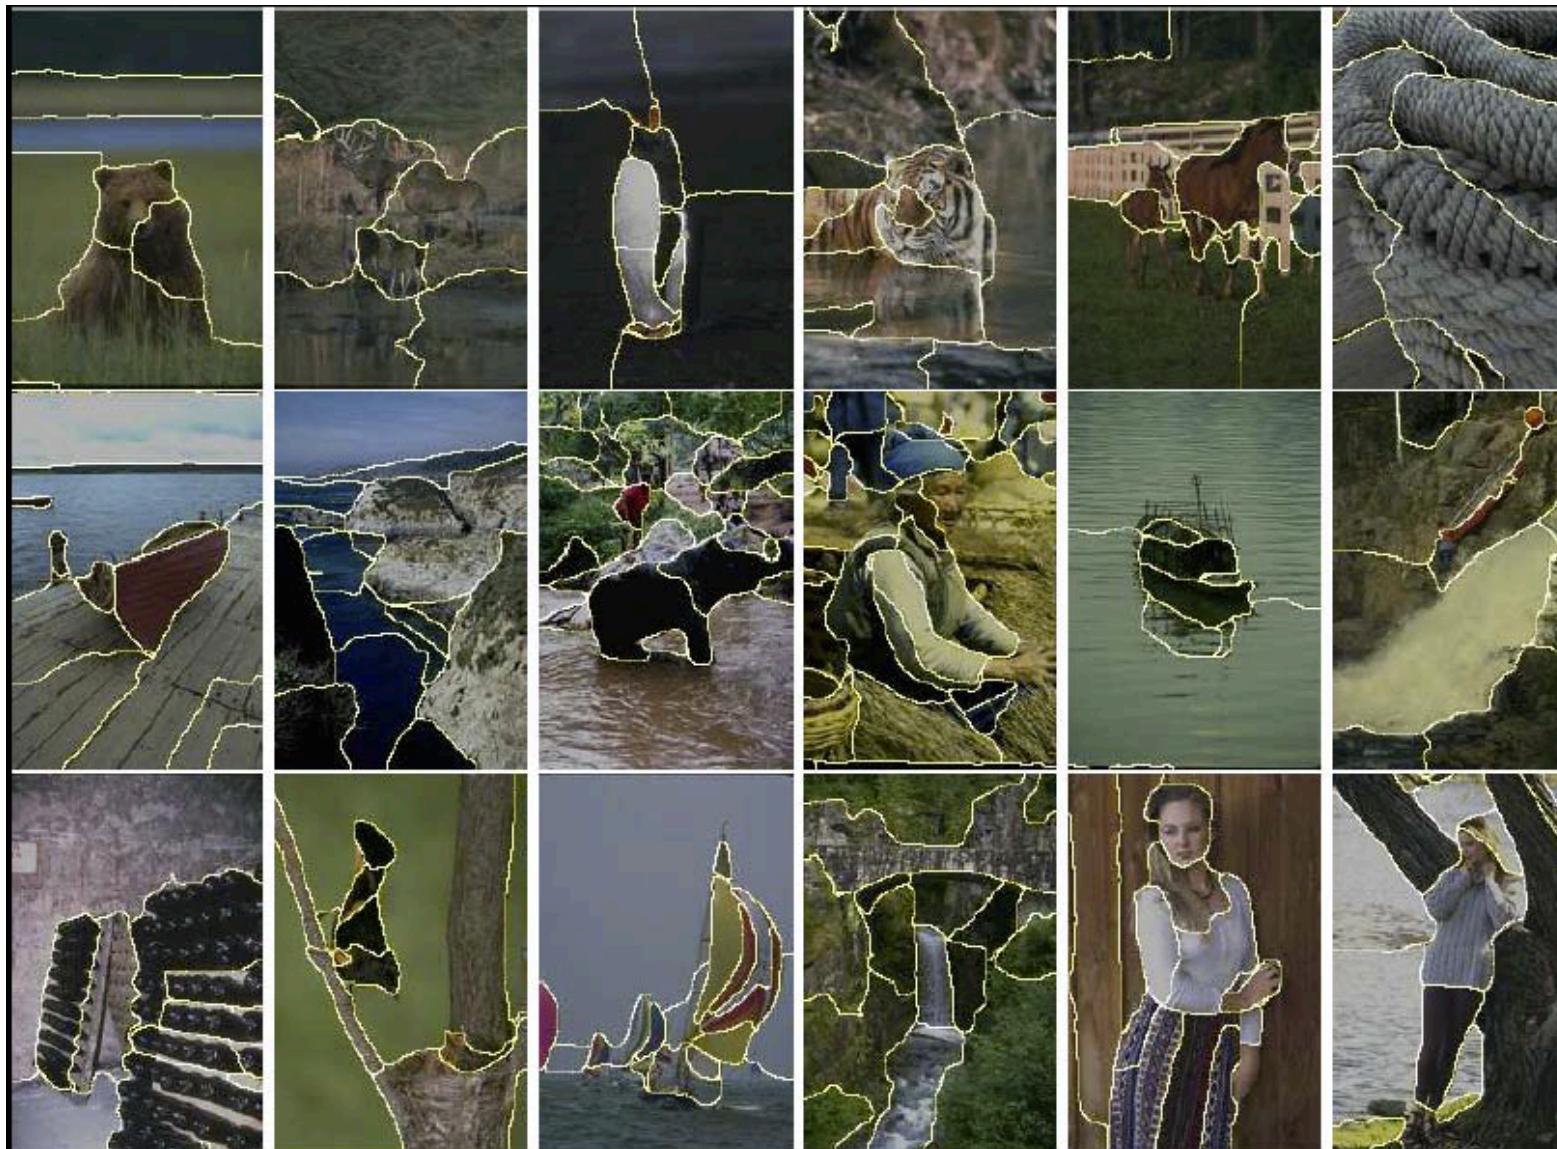
- Compute probability that a random walker arrives at seed



L. Grady, [Random Walks for Image Segmentation](#), IEEE T-PAMI, 2006

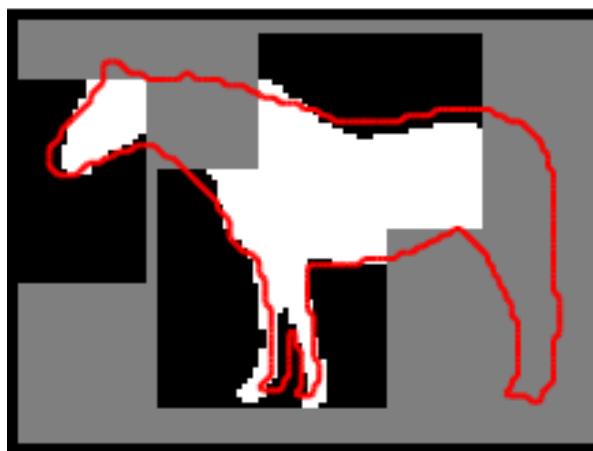
http://cns.bu.edu/~lgrady/Random_Walker_Image_Segmentation.html

Do we need recognition to take the next step in performance?



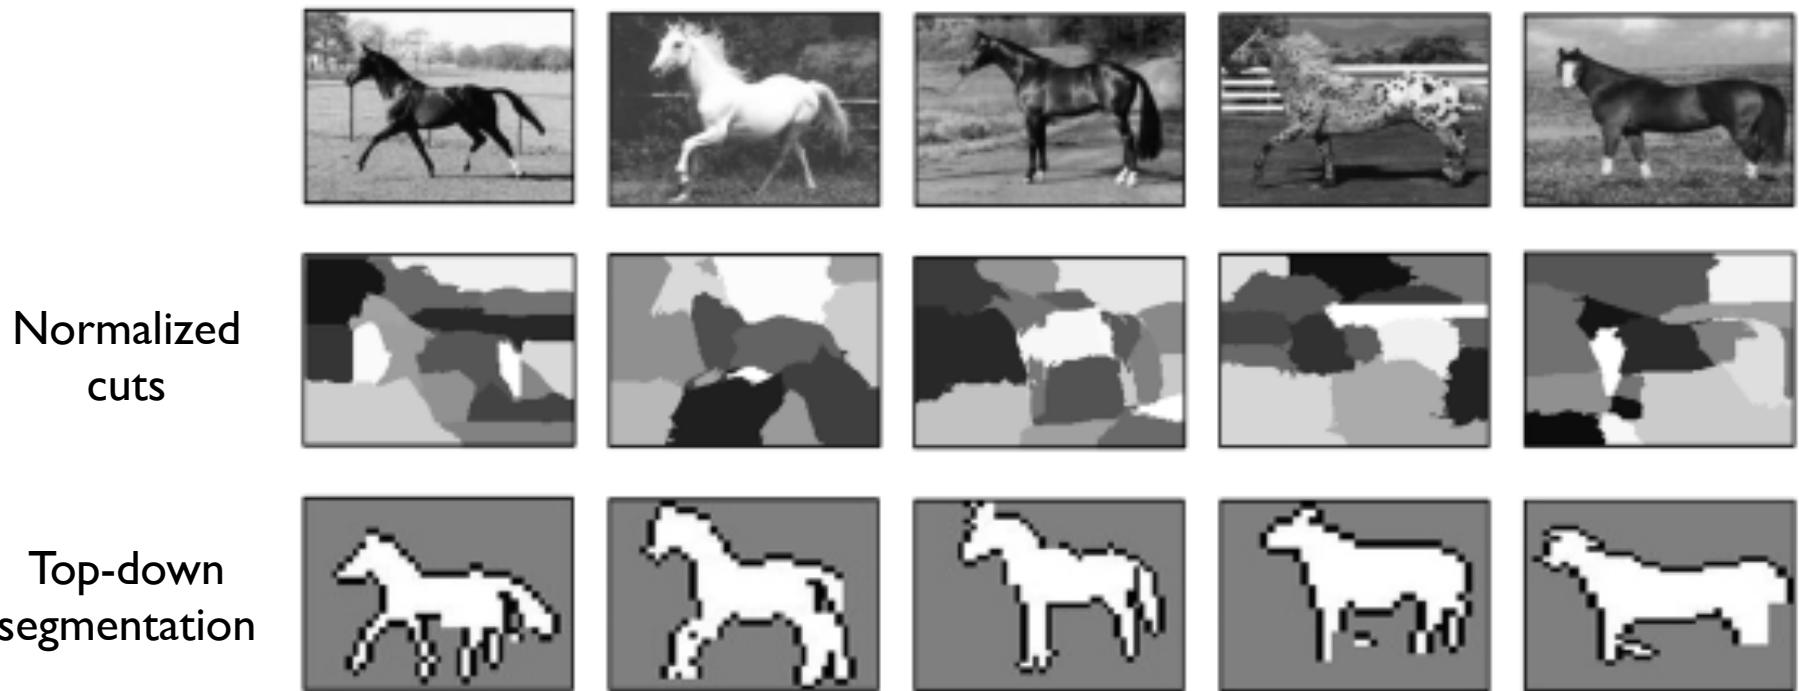
Slide credit: B. Freeman and A. Torralba

Top-down segmentation



- E. Borenstein and S. Ullman, [Class-specific, top-down segmentation](#), ECCV 2002
- A. Levin and Y. Weiss, [Learning to Combine Bottom-Up and Top-Down Segmentation](#), ECCV 2006.

Top-down segmentation



- E. Borenstein and S. Ullman, [Class-specific, top-down segmentation](#), ECCV 2002
- A. Levin and Y. Weiss, [Learning to Combine Bottom-Up and Top-Down Segmentation](#), ECCV 2006.

Motion segmentation



Input sequence

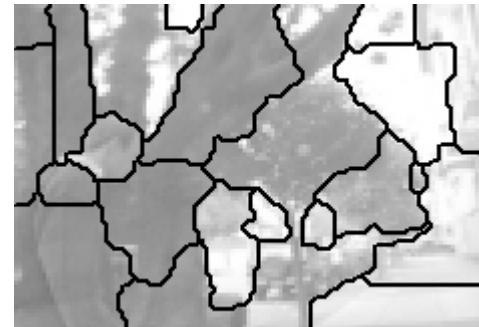
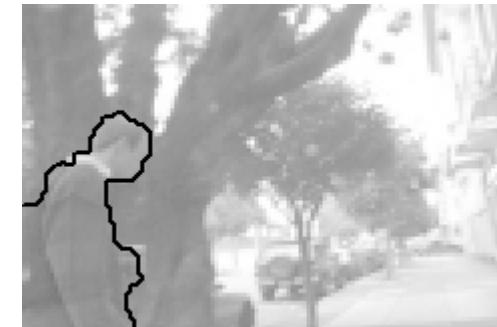


Image Segmentation



Motion Segmentation



Input sequence



Image Segmentation



Motion Segmentation

A. Barbu, S.C. Zhu. [Generalizing Swendsen-Wang to sampling arbitrary posterior probabilities](#), IEEE TPAMI, 2005.

Slide credit: K. Grauman